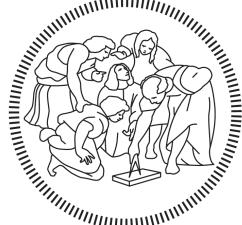


POLITECNICO DI MILANO
Scuola di Ingegneria dell'Informazione



POLITECNICO
MILANO 1863

POLO TERRITORIALE DI COMO
Master of Science in Computer Engineering

ECG-ira: An efficient mobile app for ECG analysis

Supervisor:

Prof. Giuseppe Pozzi

Co-Supervisors:

Diego Ulisse Pizzagalli, M.Sc.

Rodolfo Pizzuto, M.D.

Master Graduation Thesis by:

Antonello Fodde - 817371

Chai Botta - 817333

Academic Year 2015-2016

It's destiny

Abstract

In recent years, the healthcare market showed a continuously increasing trend of interest, and experienced a rapid development. This development has taken advantage of the potential made available by the latest mobile technologies: many companies created proprietary devices suitably designed for specific purposes, ranging from support to the world of fitness, such as analysis of body mass, or heartbeat, to those focused on the medical field, such as blood pressure sensors or glucose level meters. Many of these devices may easily interface with mobile apps, synchronizing and analyzing collected data, presenting them to the user in the most accessible way.

Our work fits in this area, more specifically in the medical one. Our work belongs to a wider project, which saw the work of some students who came first: it presents itself as an evolution in a more modern way. During our work, we used, as a starting point, a device and a PC software (previously developed) to interface with. The device is an ECG (electrocardiographic) signal acquisition module, namely ZEcg: its main features are its small size and its connection interface, which uses the Bluetooth technology, being so wireless. The software, originally deployed on PCs, aims at acquiring and visualizing in real time the ECG signals received from ZEcg, and at identifying possible arrhythmic events therein.

Leveraging the ZEcg device and the analysis algorithms, our work lead to the development of a mobile app running on smartphones and on tablets: the result is a real support for the physician on the field, and therefore provides all the necessary functionalities to capture, view and analyze an ECG recording.

The app was designed for maximum flexibility, extensibility and compatibility with more mobile devices. The responsive and speed features the application has as its first objectives, took us to the name “ECG-ira”, which stands for “ECG-instantaneous responsive analyzer”.

The following chapters will analyze the problems encountered and the implemen-

tation choices that have led to the creation of the system.

Sommario

Negli ultimi anni, il mercato della healthcare ha mostrato una continua tendenza di interesse, e subito un rapido sviluppo. Questo sviluppo ha sfruttato tutte le potenzialità rese disponibili dalle ultime tecnologie mobile: molte aziende hanno creato dispositivi proprietari opportunamente progettati per scopi specifici, andando dal supporto al mondo del fitness, come l'analisi della massa corporea o del battito cardiaco, a quelli focalizzati nel campo medicale, come i sensori di pressione sanguigna o i glucometri. Molti di questi dispositivi possono facilmente interfacciarsi con mobile app, sincronizzando e analizzando i dati raccolti, presentandoli all'utente nel modo più accessibile.

Il nostro lavoro fa parte di un progetto più ampio, che ha visto l'operato di alcuni studenti che sono venuti prima: esso si presenta come una sua moderna evoluzione. Durante il nostro lavoro, abbiamo utilizzato, come punto di partenza, un dispositivo e un software per PC (sviluppato in precedenza) con i quali ci siamo interfacciati. Il dispositivo è un modulo di acquisizione di segnale ECG (elettrocardiogramma), chiamato ZEcg: le sue principali funzionalità sono le ridotte dimensioni e la sua interfaccia di connessione, che sfrutta la tecnologia bluetooth, e quindi senza fili. Il software, originariamente sviluppato per PC, ha lo scopo di acquisire e visualizzare in tempo reale il segnale ECG ricevuto da ZEcg, e identificare possibili eventi aritmici all'interno di esso.

Sfruttando il dispositivo ZEcg ed i precedenti algoritmi di analisi, il nostro lavoro ci ha portato allo sviluppo di un'applicazione mobile per smartphone e tablet: il risultato è un vero e proprio supporto per il medico in ambito lavorativo, fornendo perciò tutte le funzionalità necessarie per acquisire, visualizzare e analizzare una registrazione ECG.

L'applicazione è stata progettata per la massima flessibilità, estensibilità e compatibilità con più dispositivi mobili. Le caratteristiche di adattamento e rapidità che

l'applicazione avrà come primi obiettivi, ci hanno portato al nome “ECG-ira”, il quale sta per “ECG-instantaneous responsive analyzer”.

Nei prossimi capitoli si analizzeranno i problemi riscontrati e le scelte implementative che hanno portato alla realizzazione del sistema finale.

Contents

Abstract	v
Sommario	vii
1 Introduction	1
2 Electrocardiography overview	5
2.1 The heart	5
2.1.1 Human heart structure	5
2.1.2 Human heart function	7
2.2 Heart electrical activity	9
2.3 Electrocardiogram	10
2.3.1 Lead I	12
2.3.2 Lead II	12
2.3.3 Lead III	13
2.3.4 Augmented leads	13
2.3.5 Precordials leads	13
2.3.6 How to read a ECG record	14
2.4 Noises and interferences	15
2.4.1 Artifact	15
2.4.2 Interference	16
2.4.3 Wandering baseline	16
2.4.4 Faulty equipment	17
3 State of Art	19
3.1 Device	19
3.1.1 Mortara ELI 10 Mobile	19

3.1.2	Philips DigiTrak XT Holter Recorder	20
3.1.3	AliveCor ECG	21
3.1.4	M-Trace (PC)Mobile	21
3.1.5	ECG Expert	22
3.2	Mobile application	22
3.2.1	Visualization only application	23
3.2.2	Visualization and Analysis	23
4	Objective	25
4.1	Preface	25
4.2	Fully functional medical mobile app as replacement to desktop app .	25
5	Requirement	27
5.1	Functional	27
5.1.1	Connection management with the acquisition device	27
5.1.2	Acquisition, storing and management of ECG records	27
5.1.3	Different ECG formats support	28
5.1.4	Dynamic display scaling	28
5.1.5	ECG record analysis integration on mobile platform	28
5.1.6	Analysis results displaying	29
5.1.7	Highly parameterizable	30
5.2	Nonfunctional	30
5.2.1	Reduced memory usage	30
5.2.2	Minimum performance rate and scalability on performance .	31
5.2.3	Wide platform compatibility and accessibility	31
5.2.4	Documentation	31
6	Problems	33
6.1	Mobile platform fragmentation	33
6.2	Native vs Cross-Platform	34
6.2.1	Native	35
6.2.2	Cross-Platform	36
6.3	Mobile hardware limitations	37
6.3.1	Memory limitations	37
6.3.2	Performance limitations	37

6.4	ECG baseline wander	38
6.4.1	State of Art	38
6.5	Signal visualization	40
7	Solution choices	43
7.1	Android Platform	43
7.2	Why native?	44
7.3	Android concurrency exploitation	45
7.3.1	Thread Overview	45
7.3.2	Threads in Android	46
7.3.3	Thread communication in Android	47
7.3.4	HandlerThread	49
7.3.5	Thread Pools	52
7.3.6	AsyncTask	54
7.4	Baseline wander solutions	57
7.4.1	Adaptive vertical displaying	57
7.4.2	Moving average filter	58
7.5	Custom View Drawing	59
7.5.1	Custom Libraries	59
7.5.2	Hardware Accelerated Drawing (GPU)	60
7.5.3	Not hardware Accelerated Drawing (CPU)	60
8	System architecture	63
8.1	Acquisition device	63
8.1.1	Mobile app	64
9	Implementation details	71
9.1	Main components	71
9.1.1	Data Sources	72
9.1.2	Display	79
9.1.3	Connection	87
9.1.4	Operations Flow	88
10	Final result	93
10.1	App Screens	93
10.1.1	Performance	98

10.1.2 Evaluation	100
10.1.3 Response Time	118
11 Conclusions	121
11.1 Future works	122
Acknowledgement	123

List of Figures

2.1	The human heart structure.[1]	6
2.2	The circulatory system with blood flow.	8
2.3	The SA node fires and electrical impulses travel through the right and left atrium.	9
2.4	The impulse then moves to the ventricular area.	10
2.5	An example of a normal ECG waveform.	11
2.6	The Einthoven's triangle shows the proper placement of leads over the chest.	12
2.7	Precordial leads and their position related to the heart and the chest horizontal plane.	14
2.8	A typical ECG paper.	14
2.9	ECG waveform interference due to artifact may cause monitoring to fail due to unreadable signals.	15
2.10	Electrical interference causes the baseline to be unstable and the signal is corrupted.	16
2.11	An example of baseline wandering due to artifacts.	16
3.1	Mortara ELI 10 Mobile, ECG acquisition device box.	20
3.2	DigiTrack, the ECG visualization.	20
3.3	AliveCor device real time acquisition on a tablet.	21
3.4	M-Trace PC device for ECG acquisition.	22
3.5	M-Trace PC device for ECG acquisition.	22
5.1	Istogram from the desktop application resulting from an analysis on a MIT/BIH record.	29
5.2	Tacogram from the desktop application resulting from an analysis on a MIT/BIH record.	29

5.3 ST+/ST- graph from the desktop application resulting from an analysis on a MIT/BIH record.	30
6.1 Mobile platform's share evolution (smartphone sales), 2009–13.	34
6.2 Native app development process.	35
6.3 Native vs Cross-Platform cost and time factors.[2]	36
6.4 Memory reserved to an app in Android.	38
6.5 The baseline wander artifact in the ECG signal.	39
6.6 Relation between distortion and moving window length for wavelet package translation (WPT), traditionally used MA, and our proposed high-pass filter based on a statistical weighted moving average (SMA). (a) Maximum error (ME) vs. moving window length; (b) Normalized root mean square error (NRMSE) vs. moving window length. M is the number of sub-bounds. SMA is the same as MA when M=1.	41
7.1 Top Mobile Operating Systems Per Country, Jan 2016. Statcounter.com	44
7.2 Overview of the message-passing mechanism between multiple producer threads and one consumer thread.	48
7.3 The execution lifecycle of AsyncTask.	56
7.4 The dynamic displaying result after a patient movement, causing the baseline wander artifact.	58
8.1 ZEcg device block diagram including all the modules components. . .	64
8.2 ADS1198 functional diagram showing the 8 channels representing the 8 leads used during the ECG record acquisition.	65
8.3 Android Build system process. How and which component are involved during an android application build and compilation.	66
8.4 ECG-ira project structure and packages inside Android Studio IDE. .	67
9.1 Class diagram of the SampleSource interface.	73
9.2 The movement of the file pointer in the DatReader, after the user change direction of scrolling the ECG paper.	75
9.3 The iteration between SampleSource and SampleDisplay after the user scroll the ECG paper.	76
9.4 The animation mechanism of DatReader using ScheduledThreadPoolExecutor.	77

9.5	Class diagram of the DatReader class.	78
9.6	Class diagram of the SampleDisplay interface.	79
9.7	The size of the ECG paper block after the computation of the method calculateMetrics().	81
9.8	Class diagram of the SignalScrollView class.	83
9.9	The different ECG visualization types. The one on top is the scrolling drawing, the one in the bottom is the oscilloscope drawing. The time elapsed after the movement of the ECG paper is one second.	84
9.10	Class diagram of the DrawingHelper class.	86
9.11	Class diagram of the Drawer class.	87
9.12	Activity diagram of a new ECG record acquisition operation.	89
9.13	Activity diagram of a new ECG record acquisition operation.	91
10.1	Two screens of the application, respectively the home screen and the form screen to be compiled by the patient	94
10.2	Two screens of the application representing the realtime acquisition and the list of records within the device	95
10.3	Two screens of the application representing the info dialog for a certain record and the visualization of that record	96
10.4	Two screens of the application representing the info dialog for a certain record and the visualization of that record	97
10.5	Screen of the three graphs generated after an analysis of the records. They are graphs summing up the complete record characteristics. Useful for a general overview of the patient record and health status.	98
10.6	The customization section of the application. A list of available settings parameters to tune the application behaviour.	99
10.7	An example of detailed view of an application memory usage.	102
10.8	Memory usage during application idle state. It is quite constant (due to user interaction it can slightly increase or decrease if no interaction at all for a while).	103
10.9	Details about memory usage and allocation over the different RAM section for the application EC-ira.	104
10.10	Two memory usage views, relatively the realtime view with the application running and the per process memory view and allocation.	105

10.11Memory usage during an analysis launched over a record. The step is due to the record allocation in a memory buffer.	106
10.12Memory usage at the exact time the algorithm performs the final steps by computing the ST+/St- areas (execution finish at 10m34s).	106
10.13Memory usage details by RAM types.	107
10.14Memory usage of ECG-ira in Idle state.	108
10.15Memory usage of ECG-ira in Idle state by RAM types.	108
10.16Two memory usage views during a normal app execution, on the second figure a GC is called before the memory reach the heap	109
10.17Memory usage details by RAM types.	111
10.18Two memory usage views during the execution of the analysis algorithms, on the first figure 10.18a the step sign the execution satrting point, in figure 10.18b the step is due to the end of the execution . .	112
10.19Memory usage when a record analysis is called. The memory usage is divide by RAM types.	113
10.20Cpu usage during idle state, the application is on foreground on a static page.	114
10.21Cpu usage when a record is plot on screen.	115
10.22Cpu usage when we close a plotting screen and reopen it.	115
10.23CPU usage during analysis state. An analysis command is launched over an ECG record strip.	116
10.24CPU usage on idle pages(app open on static screen page)	117
10.25CPU usage during a record plotting)	117
10.26usage when the analysis algorithm is launched over a record)	118

Chapter 1

Introduction

ECG-ira stands for ECG (electrocardiographic) instantaneous responsive analyzer, and it is an Android application developed for the acquisition, visualization and analysis of the electrocardiographic signals. The application acquires and stores the leads from the ZEcg acquisition device (using the ZEcg device format) property of the Politecnico of Milan: the app can also open and visualize other ECG formats from other devices.

Ecg-ira is part of a greater and long term project (ZEcg itself was part of a first step). ECG-ira aims at exploiting the reliability and performance of a complex software for electrocardiographic signal acquisition, visualization and processing running over a smartphone device through a mobile application. Apart of the analysis algorithm which was already implemented during a previous thesis (by Diego Ulisse Pizzagalli and Simone Battaglia [3]), the application of the current thesis was designed and implemented from scratch. During the process we dealt with the device limitations, in terms of performance power and limited amount of available memory. For instance, we had to deal with the exclusive memory allocated for an app by the OS. This memory allocation is called heap and can vary a lot, depending on the devices. In some low-end devices, it can be limited to 16MB.

We overcame many issues, related also to small and different screen sizes and density of pixel, due to the great number of different devices currently on the market. We had to exploit the multithreading and efficient memory usage strategies in order to achieve responsiveness and fulfill the medical requirement for an ECG compliant application.

The result is an application which is easy to use, following all the best design

principles as specified by the official Google guidelines for responsive UI and UX [4]. By taking advantage of the multithreading capabilities and the usage of all the available cores into a device, we achieved an application which performs fast and well.

This thesis is structured as follows:

- the Introduction chapter gives a brief but detailed overview about the heart, its functionalities and how an electrocardiogram is related to it. We describe the ECG and the electrical signals it detects.
- the State of Art chapter depicts the scenario of the actual devices for ECG acquisition and some available applications on the market which allow one to open and read an ECG record.
- the Objective chapter describes and explains the goal of this thesis work.
- the Requirements chapter lists all the functional and the nonfunctional sets of features which came up during the planning phase. The final result must comply all of them.
- the Problem chapter deals with all the issues related to the project development. We discuss the issues about the development platform to be adopted, the hardware limitations on a smartphone device and the problems strictly related to the ECG signals.
- the Solution Choices chapter describes our implementation choices, from the platform choice to the programming language choice to the proprietary implementation of a drawing library.
- the System Architecture chapter provides the reader with a general overview of the project structure, considering the acquisition device, the architecture of the mobile application, and its main functionalities.
- the Implementation Details chapter describes all the major components and the Java classes. This chapter also provides the reader with hints for future implementations and enhancements.
- the Results chapter proposes some screenshots from the application; the chapter also describes some preliminary performance analysis considering the required amount of memory, the CPU workload, and the response time of some portions

of the code. Tests were performed on many different devices, but data from the older devices (2011) and the newest (November 2015) one are compared, only.

- the Conclusion chapter sums up the overall results and an overview of the work done. The chapter also provides the reader with hints for further implementations, to make ECG-ira really an all in one solution as an ECG mobile application. We also discuss the future of mobile software with respect to the desktop one, according to the trends and increasing wide-spread diffusion of mobile environments, especially in the health care scenario.

Chapter 2

Electrocardiography overview

This chapter will introduce some basic but fundamental concepts about electrocardiography starting from the heart to the ECG and all issues related to the topic. We will start introducing the heart, its functionality and the entire circulatory system. After that we will describe in details the electrical activity inside the heart and how heart beats are generated. Following there will be a description of the electrocardiogram and the ECG signals. In the last section of this chapter we will discuss about all the noises and interference related to the ECG signal during its acquisition.

2.1 The heart

This chapter's focus is to describe in details the human heart. We will start from the structure to end up describing the heart functionality.

2.1.1 Human heart structure

The human heart is an organ that pumps blood throughout the body via the circulatory system, supplying oxygen and nutrients to the tissues and removing carbon dioxide and other wastes.

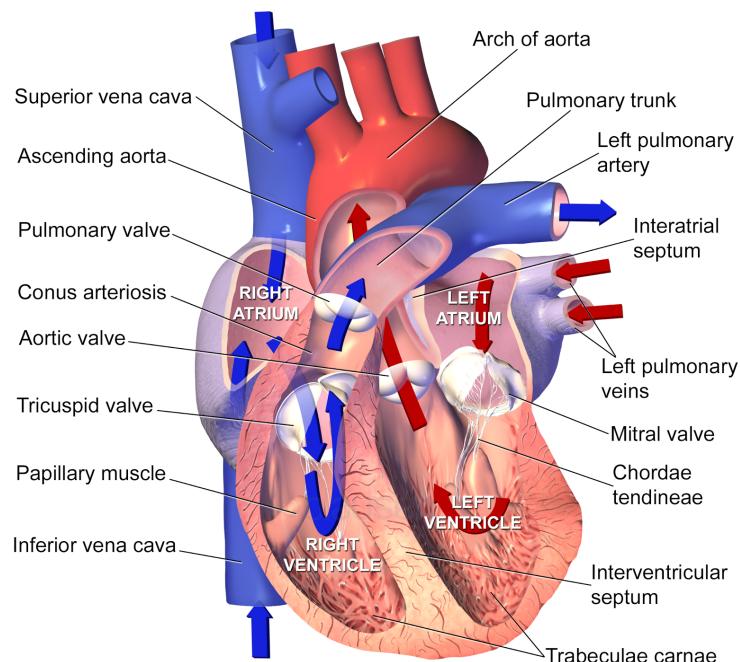
This fundamental organ has four chambers: two upper chambers(the atrial) and two lower ones(the ventricles). The right atrium and the right ventricle together make up the "right heart", and the left atrium and left ventricle make up the"left heart". The two sides of the heart are separated by a muscle called the septum.

A double-walled sac called the pericardium, encases the heart, which serves to protect

the heart and anchors it inside the chest. Between the outer layer, the parietal pericardium, and the inner layer, the serous pericardium, runs pericardial fluid, which lubricates the heart during contractions and movements of the lungs and diaphragm. The heart outer wall consists of three layers. The outermost wall layer, or epicardium, is the inner wall of the pericardium. The middle layer, or myocardium, contains the muscle that contracts. The inner layer, or endocardium, is the lining that contacts the blood.

The tricuspid valve and the mitral valve make up the atrioventricular (AV) valves, which connect the atria and the ventricles. The pulmonary semilunar valve separates the right ventricle from the pulmonary artery, and the aortic valve separates the left ventricle from the aorta. The heartstrings, or chordae tendineae, anchor the valves to heart muscles.

The sinoatrial node produces the electrical pulses that drive heart contractions.



Sectional Anatomy of the Heart

Figure 2.1: The human heart structure.[1]

2.1.2 Human heart function

The heart circulates blood through two pathways: the pulmonary circuit and the systemic circuit.

In the pulmonary circuit, deoxygenated blood leaves the right ventricle of the heart via the pulmonary artery and travels to the lungs, then returns as oxygenated blood to the left atrium of the heart via the pulmonary vein.

In the systemic circuit, oxygenated blood leaves the body via the left ventricle to the aorta, and from there enters the arteries and capillaries where it supplies the body's tissues with oxygen. Deoxygenated blood returns via veins to the venae cavae, re-entering the heart's right atrium.

Of course, the heart is also a muscle, so it needs a fresh supply of oxygen and nutrients too. After the blood leaves the heart through the aortic valve, two sets of arteries bring oxygenated blood to feed the heart muscle. The left main coronary artery, on one side of the aorta, branches into the left anterior descending artery and the left circumflex artery. The right coronary artery branches out on the right side of the aorta.

Blockage of any of these arteries can cause a heart attack, or damage to the muscle of the heart. A heart attack is distinct from cardiac arrest, which is a sudden loss of heart function that usually occurs as a result of electrical disturbances of the heart rhythm. A heart attack can lead to cardiac arrest, but the latter can also be caused by other problems.

The heart contains electrical "pacemaker" cells, which cause it to contract — producing a heartbeat.

Each cell has the ability to be the 'band leader' and have everyone follow. In people with an irregular heartbeat, or atrial fibrillation, every cell tries to be the band leader, which causes them to beat out of sync with one another.

A healthy heart contraction happens in five stages:

1. In the first stage (early diastole), the heart is relaxed.
2. Then the atrium contracts (atrial systole) to push blood into the ventricle.
3. Next, the ventricles start contracting without changing volume.
4. Then the ventricles continue contracting while empty.
5. Finally, the ventricles stop contracting and relax.

Then the cycle repeats.

Valves prevent backflow, keeping the blood flowing in one direction through the heart.

Some interesting data about the human heart are:

- A human heart is roughly the size of a large fist.
- The heart weighs between about 280 to 340 grams in men and 230 to 280 grams in women.
- The heart beats about 100,000 times per day (about 3 billion beats in a lifetime).
- An adult heart beats about 60 to 80 times per minute.
- Newborns' hearts beat faster than adult hearts, about 70 to 190 beats per minute.
- The heart pumps about 6 quarts (5.7 liters) of blood throughout the body.
- The heart is located in the center of the chest, usually pointing slightly left

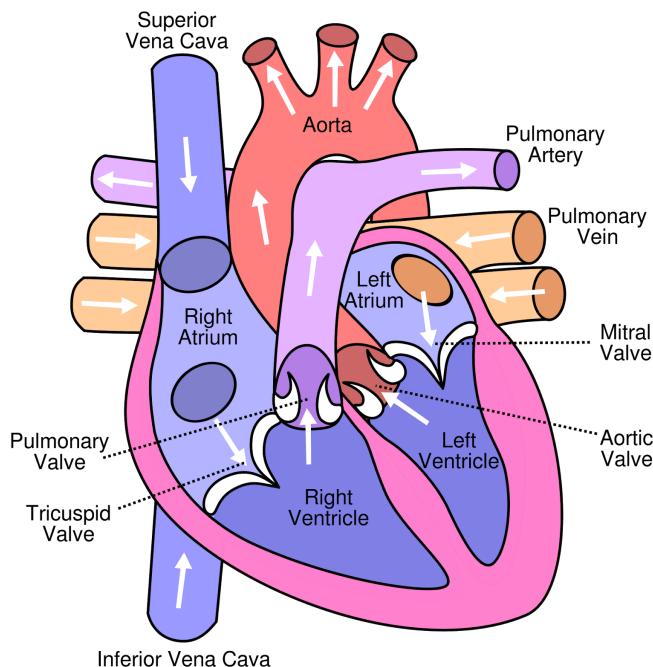


Figure 2.2: The circulatory system with blood flow.

2.2 Heart electrical activity

The heart has a natural pacemaker that regulates the pace or rate of the heart. It sits in the upper portion of the right atrium (RA) and is a collection of specialized electrical cells known as the SINUS or SINOATRIAL (SA) node.

Like the spark-plug of an automobile it generates a number of "sparks" per minute. Each "spark" travels across a specialized electrical pathway and stimulates the muscle wall of the four chambers of the heart to contract (and thus empty) in a certain sequence or pattern. The upper chambers or atria are first stimulated. This is followed by a slight delay to allow the two atria to empty. Finally, the two ventricles are electrically stimulated. In an car, the number of sparks per minute generated by a spark plug is increased when you press the gas pedal or accelerator. This revs up the motor. In case of the heart, adrenaline acts as a gas pedal and causes the sinus node to increase the number of sparks per minute, which in turn increases the heart rate. The release of adrenaline is controlled by the nervous system. The heart normally beats at around 72 times per minute and the sinus node speeds up during exertion, emotional stress, fever, etc., or whenever our body needs an extra boost of blood supply. In contrast, it slows down during rest or under the influence of certain medications. Well trained athletes also tend to have a slower heart beat.

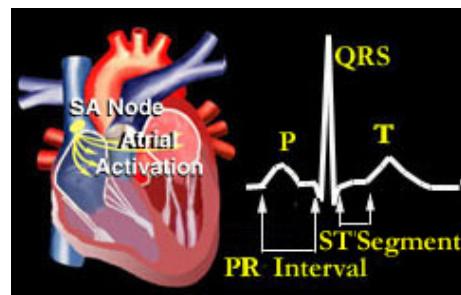


Figure 2.3: The SA node fires and electrical impulses travel through the right and left atrium.

The sequence of electrical activity within the heart is displayed in the diagrams above and occurs as follows:

1. As the SA node fires, each electrical impulse travels through the right and left atrium. This electrical activity causes the two upper chambers of the heart to

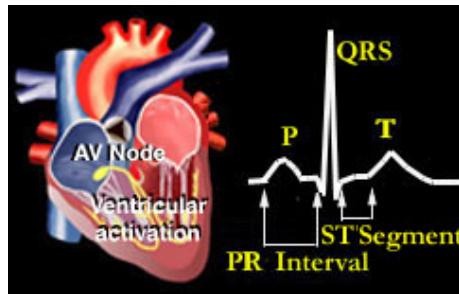


Figure 2.4: The impulse then moves to the ventricular area.

contract. This electrical activity can be recorded from the surface of the body as a "P" wave" on the patient's EKG or ECG (electrocardiogram).

2. The electrical impulse then moves to an area known as the AV (atrio-ventricular) node. This node sits just above the ventricles. Here, the electrical impulse is held up for a brief period. This delay allows the right and left atrium to continue emptying its blood contents into the two ventricles. This delay is recorded as a "PR interval." The AV node thus acts as a "relay station" delaying stimulation of the ventricles long enough to allow the two atria to finish emptying.
3. Following the delay, the electrical impulse travels through both ventricles (via special electrical pathways known as the right and left bundle branches). The electrically stimulated ventricles contract and blood is pumped into the pulmonary artery and aorta. This electrical activity is recorded from the surface of the body as a "QRS complex". The ventricles then recover from this electrical stimulation and generates an "ST segment" and T wave on the ECG.

2.3 Electrocardiogram

An electrocardiogram(abbreviated as ECG or EKG) is a test that measures the electrical activity of the heartbeat. With each beat, an electrical impulse (or wave) travels through the heart. This wave causes the muscle to squeeze and pump blood from the heart. A normal heartbeat on ECG will show the timing of the top and lower chambers.

The right and left atria or upper chambers make the first wave called a "P wave" following a flat line when the electrical impulse goes to the bottom chambers. The right and left bottom chambers or ventricles make the next wave called a "QRS

complex.” The final wave or “T wave” represents electrical recovery or return to a resting state for the ventricles.

Each waves in the figure 2.5 is no other than the result of different views or

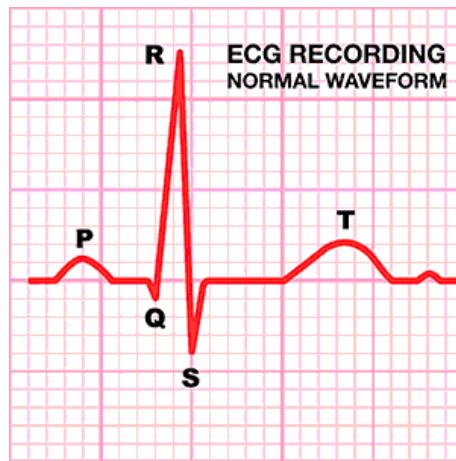


Figure 2.5: An example of a normal ECG waveform.

perspectives of the waveforms generated from the current in the heart.

There are two type of ECGs recordings: the 12-lead ECG and the rhythm strip. Both give valuable information about heart function.

We will focus our attention on the 12-lead ECG. It records information from 12 different views of the heart and provides a complete picture of electrical activity. The limb leads and the chest, or precordial, leads reflect information from the different planes of the heart. Different leads provide different information. The six limb leads I, II, III, augmented vector right (aVR), augmented vector left (aVL), and augmented vector foot (aVF) provide information about the heart’s frontal (vertical) plane. Leads I, II, and III require a negative and positive electrode for monitoring, which makes those leads bipolar. The augmented leads record information from one lead and are called unipolar.

The six precordials or V leads V1, V2, V3, V4, V5, and V6 provide information about the heart’s horizontal plane. Like the augmented leads, the precordial leads are also unipolar, requiring only a single electrode. The opposing pole of those leads is the center of the heart as calculated by the ECG.

The position of the leads are crucial for a right ECG recordings. It is common to use the so called Einthoven’s triangle, a set of positions to set up standard limb leads. The electrodes for leads I, II, III are about equidistant from the heart and form an

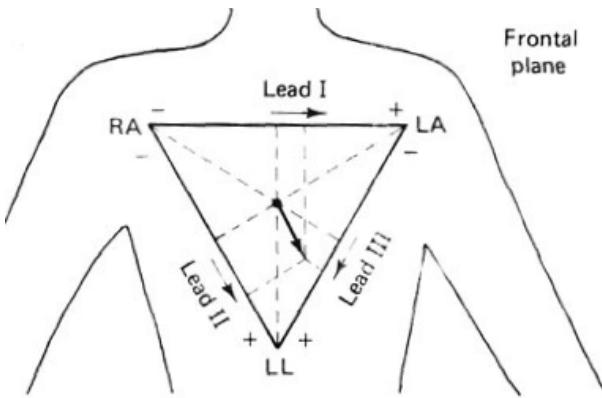


Figure 2.6: The Einthoven's triangle shows the proper placement of leads over the chest.

equilateral triangle.

2.3.1 Lead I

It provides a view of the heart that shows current moving from right to left. Because the current flows from negative to positive, the positive electrode for this lead is placed on the left arm or on the left side of the chest; the negative electrode is placed on the right arm. Lead I produces a positive deflection on ECG tracings and is helpful in monitoring atrial and hemiblock.

2.3.2 Lead II

Lead II produces a positive deflection. Place the positive electrode on the patient's left leg and the negative electrode on the right arm. For continuous monitoring, place the electrodes on the torso for convenience, with the positive electrode below the lowest palpable rib at the left midclavicular line and the negative electrode below the right clavicle. The current travels down and to the left in this lead. Lead II tends to produce a positive, high voltage deflection, resulting in tall P, R, and T waves. This lead is commonly used for routine monitoring and is useful for detecting sinus node and atrial arrhythmias.

2.3.3 Lead III

Lead III produces a positive deflection. The positive electrode is placed on the left leg; the negative electrode, on the left arm. Along with lead II, this lead is useful for detecting changes associated with an inferior wall myocardial infarction. The axes of the three bipolar limb leads I, II, and III form a triangle around the heart and provide a frontal plane view of the heart.

2.3.4 Augmented leads

Leads aVR, aVL, and aVF are called augmented leads. They measure electrical activity between one limb and a single electrode. Lead aVR provides no specific view of the heart. Lead aVL shows electrical activity coming from the heart's lateral wall. Lead aVF shows electrical activity coming from the heart's inferior wall.

2.3.5 Precordials leads

The six unipolar precordial leads (V1, V2, V3, V4, V5 and V6) are placed in sequence across the chest and provide a view of the heart's horizontal plane.

- Lead V1—The precordial lead V1 electrode is placed on the right side of the sternum at the fourth intercostal rib space. This lead corresponds to the modified chest lead MCL1 and shows the P wave, QRS complex, and ST segment particularly well. It helps to distinguish between right and left ventricular ectopic beats that result from myocardial irritation or other cardiac stimulation outside the normal conduction system. Lead V1 is also useful in monitoring ventricular arrhythmias, ST-segment changes, and bundle-branch blocks.
- Lead V2—Lead V2 is placed at the left of the sternum at the fourth intercostal rib space.
- Lead V3—Lead V3 goes between V2 and V4. Leads V1, V2, and V3 are biphasic, with both positive and negative deflections. Leads V2 and V3 can be used to detect ST-segment elevation.
- Lead V4—Lead V4 is placed at the fifth intercostal space at the midclavicular line and produces a biphasic waveform.

- Lead V5—Lead V5 is placed at the fifth intercostal space at the anterior axillary line. It produces a positive deflection on the ECG and, along with V4, can show changes in the ST segment or T wave.
- Lead V6—Lead V6, the last of the precordial leads, is placed level with V4 at the midaxillary line. This lead produces a positive deflection on the ECG.

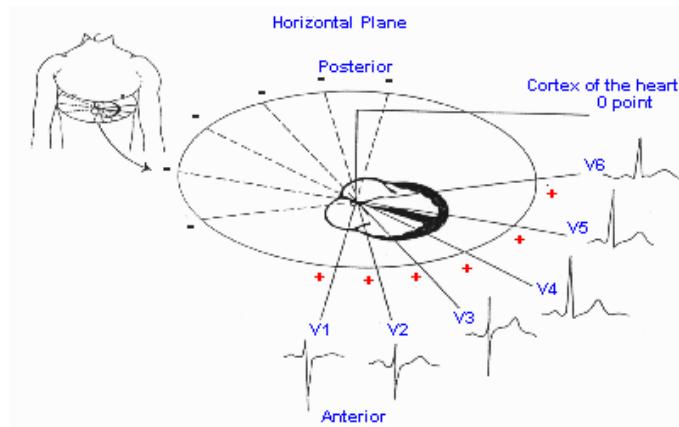


Figure 2.7: Precordial leads and their position related to the heart and the chest horizontal plane.

2.3.6 How to read a ECG record

Waveforms produced by the heart's electrical current are recorded on graphed ECG paper by a stylus. An ECG paper consists of horizontal and vertical lines forming a grid. A piece of ECG paper is called an ECG strip or tracing. The horizontal axis of

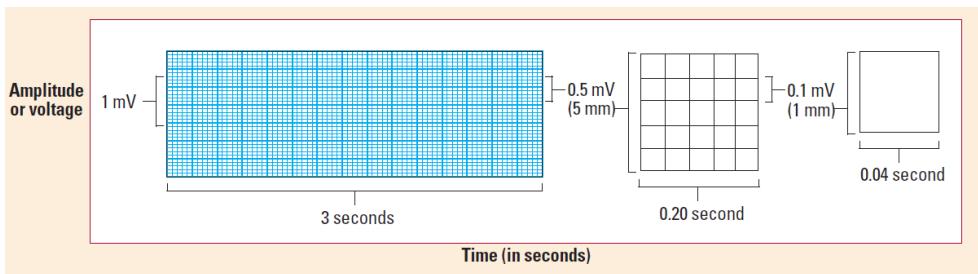


Figure 2.8: A typical ECG paper.

the ECG strip represents time. Each small block equals 0.04 second, and five small blocks form a large block, which equals 0.2 second. This time increment is determined

by multiplying 0.04 second (for one small block) by 5, the number of small blocks that compose a large block. Five large blocks equal 1 second (5×0.2). When measuring or calculating a patient's heart rate, a 6-second strip consisting of 30 large blocks is usually used. The ECG strip's vertical axis measures amplitude in millimeters (mm) or electrical voltage in millivolts (mV). Each small block represents 1 mm or 0.1 mV; each large block, 5 mm or 0.5 mV. To determine the amplitude of a wave, segment, or interval, count the number of small blocks from the baseline to the highest or lowest point of the wave, segment, or interval.

2.4 Noises and interferences

Obtaining a reliable ECG recording is still an issue. In fact there may occur many problems interfering with the signals. Some of these problems include artifacts from patient movement and poorly placed or poorly functioning equipment.

2.4.1 Artifact

Artifact , also called waveform interference, may be seen with excessive movement (somatic tremor). The baseline of the ECG appears wavy, bumpy, or tremulous. Dry electrodes may also cause this problem to occur due to poor contact.

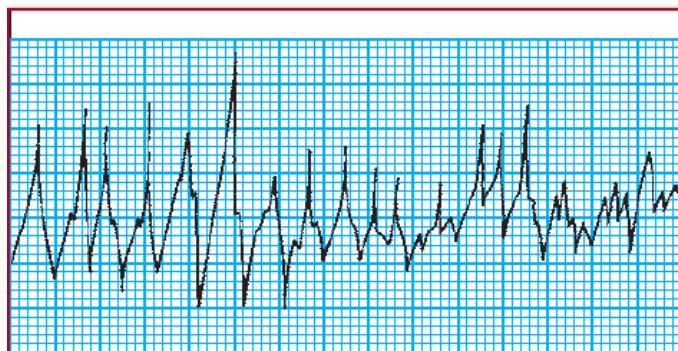


Figure 2.9: ECG waveform interference due to artifact may cause monitoring to fail due to unreadable signals.

2.4.2 Interference

Electrical interference, also called 60-cycle interference, is caused by electrical power leakage. It may occur due to interference from other room equipment or improperly grounded equipment. As a result, the lost current pulses at a rate of 60 cycles per second. This interference appears on the ECGs as a baseline that is thick and unreadable.

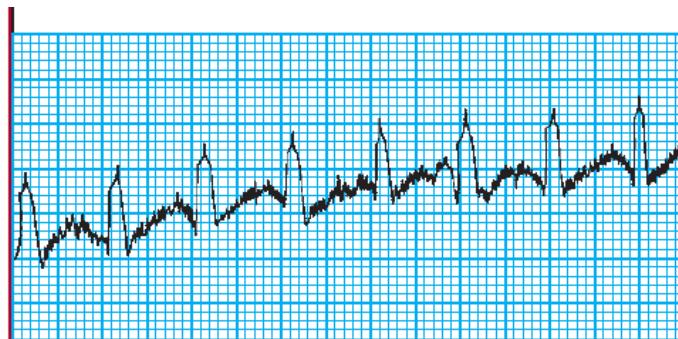


Figure 2.10: Electrical interference causes the baseline to be unstable and the signal is corrupted.

2.4.3 Wandering baseline

A wandering baseline undulates, meaning that all waveforms are present but the baseline is not stationary. It can be caused by movement if the chest wall during respiration, poor electrode placement, or poor electrode contact.

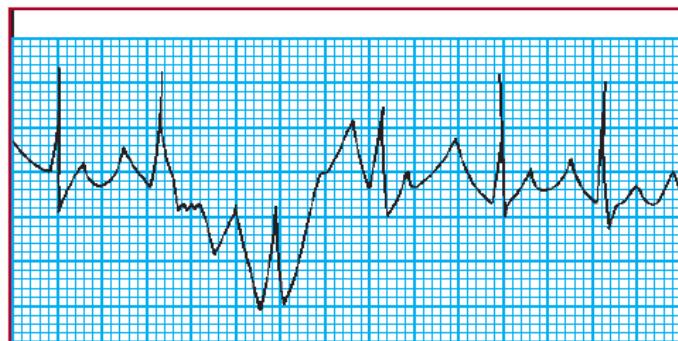


Figure 2.11: An example of baseline wandering due to artifacts.

2.4.4 Faulty equipment

Faulty equipment, such as broken lead wires and cables, can also cause monitoring problems.

Chapter 3

State of Art

3.1 Device

The personal health care market has changed a lots and recently new products and devices are showing up on the market. We will describe briefly the most relevant and similar products as mobile ECG acquisition devices. We evaluate the following solutions:

- Mortara ELI 10 Mobile
- Philips DigiTrak XT Holter Recorder
- M-Trace (PC) Mobile
- ECG Expert

3.1.1 Mortara ELI 10 Mobile

This device offers an all in one solution for 12 leads ECG acquisition. It is compact and complete as it provides an alphanumeric keyboard and a screen for real time visualization and the possibility to send the record via GPRS/3G channels. For each devices a SIM card is required . The device can also read and interpret the ECG supporting the doctor. Interesting feature is its great interoperability with the main ECG data management systems.



Figure 3.1: Mortara ELI 10 Mobile, ECG acquisition device box.

3.1.2 Philips DigiTrak XT Holter Recorder

This is the smaller acquisition device on the market. Thanks to a proprietary algorithm from Philips it can derive all the 12 ECG leads using only 5 leads. It weighs 62g and the internal battery lasts till 7 days. It also has a small screen showing 1 real time signal at a time.



Figure 3.2: DigiTrack, the ECG visualization.

3.1.3 AliveCor ECG

An innovative solution even though it doesn't offer a complete solution for ECG acquisition and analysis. This small sensor can be attached on the back of your smartphone making it an ECG acquisition device. It can record only one ECG signal (D1), so also the analysis is limited to a few types of arrhythmias . The record length is also limited to 5 minutes.



Figure 3.3: AliveCor device real time acquisition on a tablet.

3.1.4 M-Trace (PC)Mobile

M-Trace PC is an completed 12 leads ECG acquisition device. With the device it comes a mobile application and a desktop pc application used to visualize and analyze the ecg signals. The device is really portable with dimensions 95x64x28mm. The company offers also a more portable device (M-Trace Mobile) to be used by privates at their home. The mobile version cannot acquire a full record but only test records with 6 leads. Its main purpose it to send the test records via GSM/GPRS to the doctor for a faster review.



Figure 3.4: M-Trace PC device for ECG acquisition.

3.1.5 ECG Expert

ECG Expert produced by CSE Medical is a completed solution for ECG acquisition. The device comes with fully supported software for both PC desktop (Windows and Mac) both smartphones (Android, iOS). The device is rechargeable and makes use of a wireless connection via Bluetooth as exchanges data communication with the handheld smartphone or PC software.

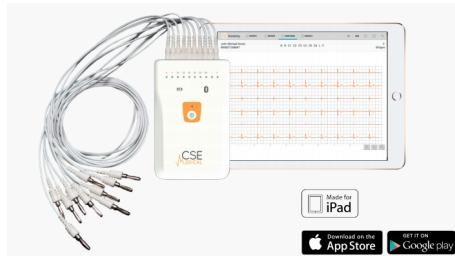


Figure 3.5: M-Trace PC device for ECG acquisition.

3.2 Mobile application

There are already mobile applications on the market store for ECG visualization and analysis supporting different formats. We can distinguish applications that only visualize the signal and the ones that also apply some analysis on the ECG signals. We listed only applications on the Google Play Store, so only Android applications because they are the only comparable with the solution we propose.

3.2.1 Visualization only application

The application on the market able to visualize the ECG signal are:

- *StribogECG*: an Android application based on an open source project under GLP v3 licence. It uses Biosig library to read ECG formats such as scp, xml (hl7), ecg and dgf. The software is only provided as it is and it requires to the user to already have the ecg files stored in those supported formats.
- *AndroidECG*: application on beta release, it was developed by Paco Gonzàles as thesis project during the Master course in Computer Science at the University of Murcia. The application is able to show ECG signals of the following formats: binary, scp, 212. As additional feature it is a basic analysis over the signal to detect QRS complex, P waves, ST segments and T waves. It is also possible to send the ECG record via email.

3.2.2 Visualization and Analysis

The applications on the market that also provide a more detailed analysis over the ECG signal are all bind to a specific proprietary acquisition device. By this way they lack the compatibility and interoperability requirement with other software and ECG formats.

- *M-Trace PC*: the application was developed by *M4Medical*, a Poland company providing medical devices for professionals and private customers. The application only works with the company 12-channel ECG M-Trace PC register device. The main features are the real-time monitoring interface, a patients' database management system and the possibility to share the record.
- *ECG Expert*: developed by *CSE Medical* the application works only connected to an ECG-Expert acquisition device. The main features are the real-time view of the acquisition, the analysis of the record providing information about QRS complex and heart rate, the possibility to manage patient information bind to the record and a heartbeat Normal/Abnormal classifier.

One last mobile application, which is not strictly related to ECG signal visualization and analysis but it worth to be mentioned, is the *ECG Interpretation*. This application instead provides enough detailed information about how to read and interpret the

ECG signals through 32 small lectures. All the lectures provide a picture and a short description and explanation.

Chapter 4

Objective

4.1 Preface

For a clear understanding of the next chapters we will make use of some terms listed below with the proper meanings:

1. Mobile application: it is a software running on smartphones and tablets
2. Desktop application: it is a software running on desktop pcs or notebooks
3. Acquisition device: named ZEcg, it is the device (hardware) used to acquire the ECG signal from the electrodes connected to a patient body.

4.2 Fully functional medical mobile app as replacement to desktop app

The main purpose for this thesis is to develop a medical mobile application as replacement to an original desktop application. The application needs to be standalone and independent from other software, still it can share its content and integrate other software content.

As a starting point we planned to reproduce all the desktop features such as the connection between the application and the remote device ZEcg for the ECG signal acquisition. It should also save the ECG records inside the mobile device, plot the signals and run arrhythmia recognition algorithms on them. We are aware that the user experience is different from a desktop one due to the differences in capabilities

and functionalities. Having in mind these differences, we did not try to reproduce the desktop experience. We developed instead the application having a mobile experience at first position, following the standards of mobile application designs and principles. We took advantage of the new and latest technologies mounted on the new smartphones, trying to provide to the end user the best in term of user experience, performance and application design. The main difficulty is probably to redesign and re-imagine the desktop feature from a mobile point of view. For example, if a desktop application usually makes use of keyboard and mouse, inside a modern mobile application there is only the touch input as user interaction. The differences in term of screen size, memory and CPU performance matters and should always be kept in mind during the initial planning phase. We will deal with these and others limitations, trying to achieve the best results and performances.

We believe this application can be really a replacement to a desktop application as the technology trend points to future devices with better performances in term of lower power consumption and higher operational capabilities.[5]

Chapter 5

Requirement

In the project, there was the need for a deep analysis of all the tied requirements. The result of this analysis was essential to identify the subsequent problems. We will describe all of them, distinguishing between functional and nonfunctional ones.

5.1 Functional

5.1.1 Connection management with the acquisition device

Fundamental feature to be included inside the mobile application is the capability to directly connect the smartphone device to the acquisition device ZEcg. Since this last one was designed to transmit the signals through a bluetooth channel, we have to implement and manage a bluetooth socket connection inside the application in order to receive the data.

5.1.2 Acquisition, storing and management of ECG records

For a matter of medical feature as it is a fact that there are many “standards” on saving an ECG signal, the application has to be able to manage different formats. Even though this application is designed to be used mostly for acquisition from the ZEcg device, it is also able to open and read other standard format such as the MIT-BIH, one of the most common standard in the literature of ECG. The code software behind is designed in a such way that the integration of other format is made extremely easy to add just by implementing few interfaces and classes.

5.1.3 Different ECG formats support

For a matter of medical feature as it is a fact that there are many “standards” on saving an ECG signal, the application has to be able to manage different formats. Even though this application is designed to be used mostly for acquisition from the ZEcg device, it is also able to open and read other standard format such as the MIT-BIH, one of the most common standard in the literature of ECG. The code software behind is designed in a way that the integration of other format is made extremely easy to add just by implementing few interfaces and classes.

5.1.4 Dynamic display scaling

The mobile device market is huge and there are a very large number of devices with completely different hardware and screens. As first classification we can distinguish mobile devices into smartphones and tablets. The most obvious difference is based on the screen size and the pixel density. Building a mobile application means also to deal with this number of different devices. To achieve the same experience and look and feel the application should be able to scale its view according to the device screen and the pixel density. A typical ECG signal is plot on a paper with squares of well defined size in millimeters. The mobile application has to respect such a standard independently on the screens capabilities and pixel density, so it should be able to properly scale the view and the plotting based on the hardware provided by the device.

5.1.5 ECG record analysis integration on mobile platform

To complete the set of features for the application we plan to integrate the algorithms of ECG signal processing. To have a mobile device able not only to acquire and visualize in real time the ECG but also to analyze it at runtime, it can be of vital importance, especially if the user has little knowledge about reading and interpreting an ECG graph. The integrated algorithms for arrhythmia analysis are based on a Neural Network trained to recognize the nature of the signals for the given record with high accuracy. The algorithms come from a previous thesis work[3] which belongs to Ulisse Pizzagalli, student at Politecnico of Milan.

5.1.6 Analysis results displaying

After analysis there are results that need to be shown to the user in the most friendly and understandable way. The most important results from an ECG analysis are called Istogram, Tacogram and the ST+/ST-. They are respectively graphs showing the number of heart rates of a certain value, the average heart rate at each heartbeat and the difference between the area ST+ and ST-, the area above the segment ST and the one below. This last graph is useful for ischemia detection.[6]

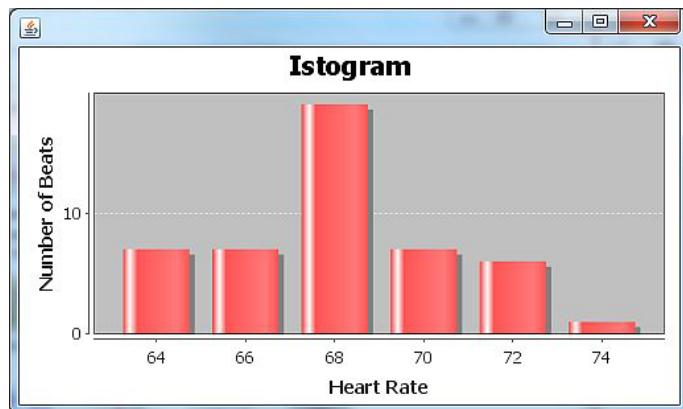


Figure 5.1: Istogram from the desktop application resulting from an analysis on a MIT/BIH record.

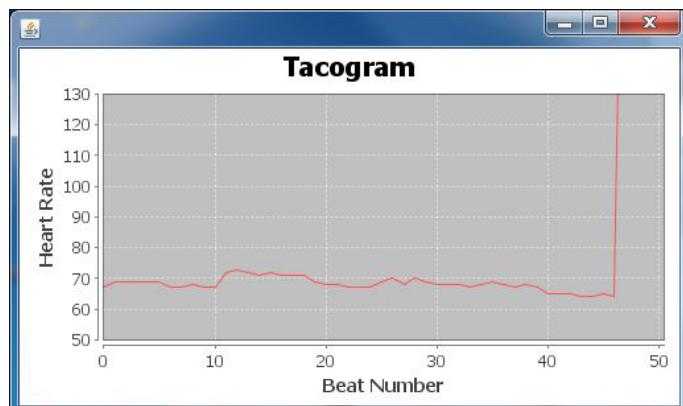


Figure 5.2: Tacogram from the desktop application resulting from an analysis on a MIT/BIH record.

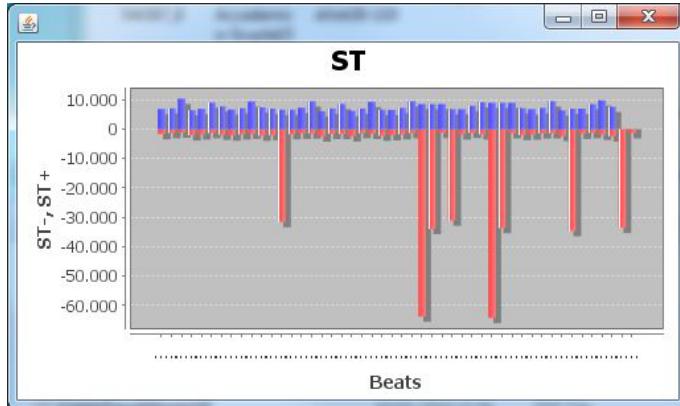


Figure 5.3: ST+/ST- graph from the desktop application resulting from an analysis on a MIT/BIH record.

5.1.7 Highly parameterizable

We believe in dynamic software, that is why we planned from the beginning on making this application dynamic. Even if the application is build on top of the ZEcg standard, we plan to make the software responsive also to other standards. To achieve this, we planned to abstract all the acquisition device independent features and functionalities. In order to support as much as possible any variants of the original acquisition device, we plan to setup customizable parameters, the only related to the hardware implementation. With a little of changes the application will be able to interface with other devices as well for acquisition.

5.2 Nonfunctional

5.2.1 Reduced memory usage

This requirement is fundamental for any project related to mobile application development. In fact, if a desktop pc in general doesn't have any problem related to memory usage (even if it is a good practice not to waste memory), on mobile devices this over-usage can bring the application to crash and get killed by the OS. The memory available is higher on new devices with respect to older ones, but it is still small so it is always a good practice to use it carefully.

5.2.2 Minimum performance rate and scalability on performance

Nowadays the new high-level mobile devices has quad-core or even octa-core CPU processors. Any application should take advantage of a such configuration, but on the other hand mobile application developers should always consider the fact that the market is still full of older and low-end devices. In order to cover at least most of the market devices their application have to run fine (with a minimum acceptable performance rate) starting from the low-end devices and, at the same time, taking advantage of last devices capabilities.

We believe modern applications should seriously take this aspect in consideration, because it will make their application scalable also from a performance point of view.

5.2.3 Wide platform compatibility and accessibility

Developing a mobile application implies building a software that has to be executed on many different platforms. The smartphone and tablet market is huge with many different devices mounting different hardware and running of the three major mobile OS (iOS, Android and Windows Phone). In the next chapters we will deal with this issue.

5.2.4 Documentation

This thesis includes also a more technical documentation about the development phase and the choices we starting from the planning phase to the development phase. The software is fully documented and with annotations and comments to increase code readability and future development on top of it. The technical documentation is included in the next chapters where we are going to discuss and motivate the implementation and the results.

Chapter 6

Problems

By identifying the requirements, we could then be able to highlight the related problems that we had to face in order to fulfill all of them. We will describe the related requirement as source of each faced problem.

6.1 Mobile platform fragmentation

Considering the requirement of a wide platform compatibility, we obviously need to face a really big problem in the mobile application world: the platform fragmentation. Starting from the first smartphone release on 2007, the iPhone from Apple, the sale of such devices keeps increasing each year. Between 2007 and 2008, sales proceeded upwards reaching the same sale rate of their computing parent, the PC. On 2009 the market signed an important inflection point, representing the beginning of an inexorable vertical rise. Although PCs were still the only ones to offer some types of functionality due to their longer replacement cycle, they were sold with a ratio of 1 : 2, compared to smartphones, over a 5-year period. This new market' growth leads to an obvious seeking of various participants, some by choice and some by necessity, in order to extract value. Android has been the prime beneficiary, having been announced on 2007 and having gone on to account for well over half of smartphone sales worldwide. Apple, meanwhile, has maintained a steady ship, leveraging its skill in product design and user experience with a finely honed marketing machine, attracting a customer that rarely defects.^[7] Fast-forwarding to 2016, if we consider the global market share held by the leading smartphone operating systems, at first position we have Android with a share of 80.7%, followed by iOS with a 17.7%. Windows Phone (1.1%), RIM

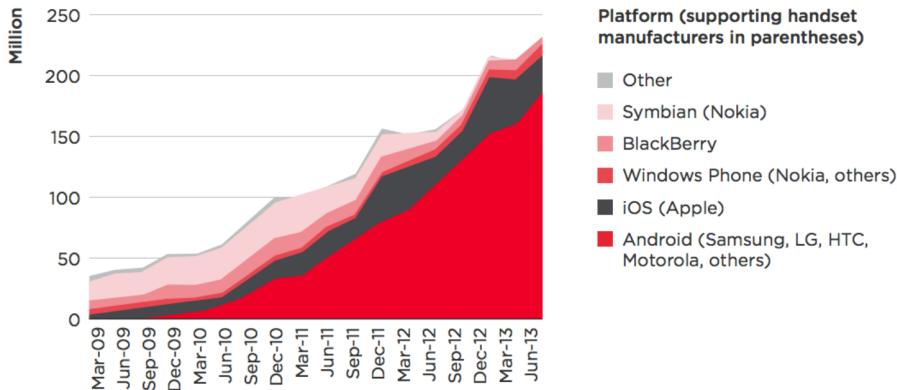


Figure 6.1: Mobile platform's share evolution (smartphone sales), 2009–13.

(0.2%) and others (0.2%) represent minor percentages.[8]

6.2 Native vs Cross-Platform

As mentioned earlier, one of the main challenges when moving to a mobile solution is the software development within a technology landscape that is highly fragmented and rapidly evolving. Mobile apps require a fair amount of customization to run on diverse platforms and a constant update due to the steady stream of new hardware, OS versions and browsers. Even a single platform (Android, Windows, Blackberry, and to a lesser extent Apple) has numerous flavors that require some degree of customization. There are also other factors such as the overlay software from different manufacturers which can affect the behaviour of an app on a particular device.

In response, the mobile industry has spawned a rapidly growing ecosystem of cross-platform and cross-device frameworks, source code analyzers, libraries of reusable components, and other tools designed to accelerate and simplify multi-platform development. New tools are constantly emerging, with new functionality, different capabilities, strengths and weaknesses.

Developer's preferences change and evolve, particularly as new tools and capabilities become available. However, the basic goals are the same: to code less and accomplish more, to reuse and recycle across multiple platforms as much as possible, and consider developing from scratch as the last resort. In addition, any tool or framework should be able to work with current and future evolving offers, and not to be locked because of a particular platform or technology.[9]

6.2.1 Native

Native app development involves developing software specifically designed for a specific platform, hardware, processor and its set of instructions. Typical programming languages are Java, Object C, Swift, C# and many other.

Native apps' major advantage over cross-platform apps is the ability to leverage device-specific hardware and software. This means that native apps can take advantage of the latest technology and API available on the mobile devices and it can well interface with other platform apps. Other advantages are a predictable increase on performance, streamlined support, native UI, native API and coherent library updates. However, the mobile platform's fragmentation makes the task of keeping up with the pace of emerging technology onerous and costly, having to develop different software for each of the different platforms (Android, iOS, Windows Phone, Symbian).

Going further on a more technical analysis, native applications are represented by executable binary files that will be installed into devices without the need for other abstraction layering to the operating system. They are able to call built in functionalities such as calendar, notifications, the dialer, email provider and others services and functionalities provided natively by the OS. Despite the fact that native applications development requires platform specific skill and expertise, this strategy delivers a higher quality user experience than other mobile application development methods (cross-platform or hybrid approach). Native apps are also best distributed through an app store.[10]

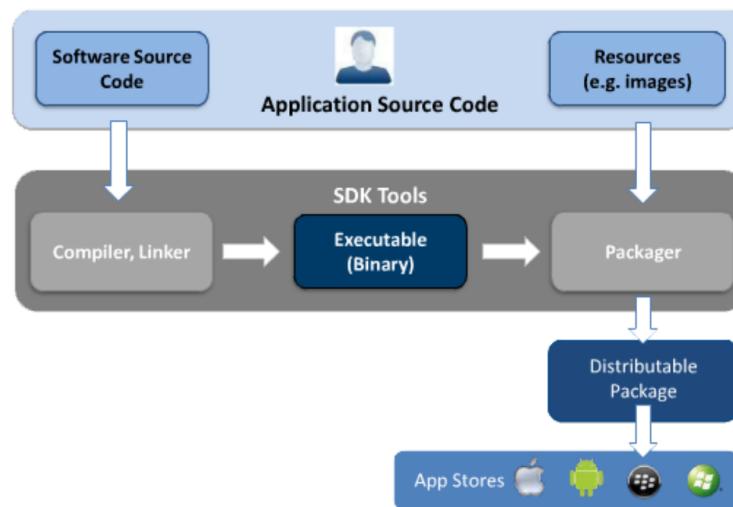


Figure 6.2: Native app development process.

6.2.2 Cross-Platform

Cross-platform development produces one code base to maintain and write targeting multiple devices and platforms. It promises lower time of development and costs. The main categories of this group are web, hybrid, interpreted and generated apps. None of the previous is neither prevalent nor the best solution to the problem of developing cross-platform mobile applications. The main benefits of cross-platform development

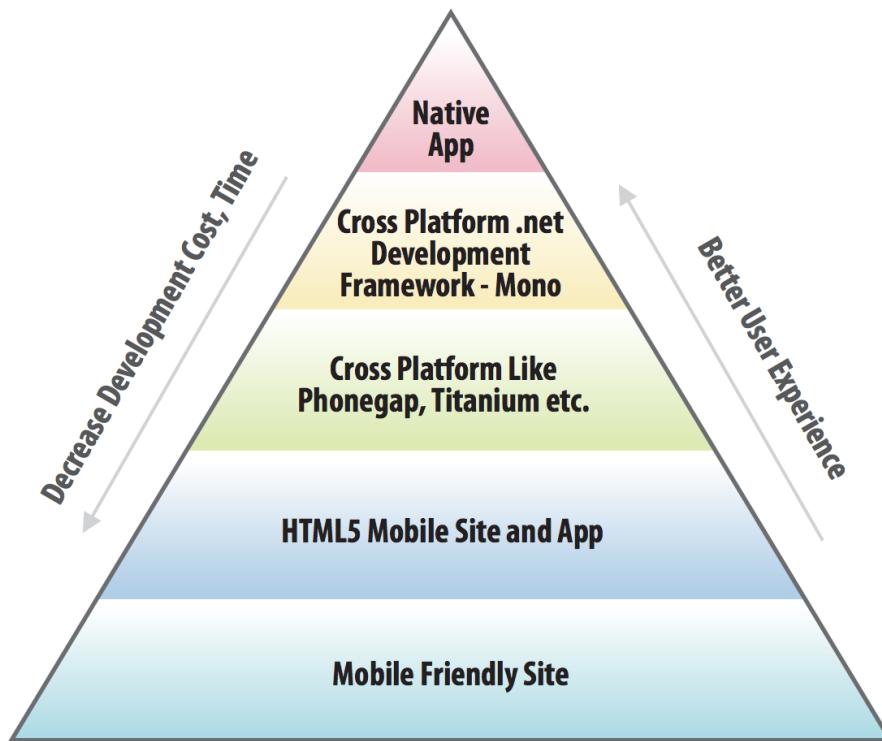


Figure 6.3: Native vs Cross-Platform cost and time factors.[2]

are:

- Reduction of required skills to develop applications due to the use of common programming languages;
- Reduction of coding work, because the source code is written once and it is compiled for each supported OS;
- Reduction of development time and long term maintenance costs;
- Decrement of API knowledge, because with these tools is not needed to know the API's of each OS, but only the API's provided by the selected tool.[11]

6.3 Mobile hardware limitations

As this project started from a previously developed desktop application, we needed to take into account all the limitation of mobile devices hardware. Despite the trend of mobile hardware is a constant increasing in computational power and memory availability, we cannot compare it with the one available nowadays in PC architectures, as the architectures on this two platforms are totally different: most of all pc are based on x86 architectures, whereas mobiles rely on ARM architectures. Furthermore, we could not commit the error of targeting only recent devices as they represent only a small percentage of all devices in the market, and it would have been in contrast to our main goal, which is a wide device compatibility. We will discuss briefly about memory and computational differences in these type of architecture in the next paragraphs.

6.3.1 Memory limitations

The first point to consider is the device memory limitation. In the previous desktop application not so much effort was spent to reduce memory allocation: all the ECG record samples were allocated in memory both during acquisition and record opening. This approach could result convenient considering its low implementation effort. Nevertheless it could be source of lots of problem when moving to a mobile device, where not only the physical memory is often small, but also the memory reserved to an application is limited (figure 6.4). These limitations are both due to the coexistent apps running in the system and the power management of the mobile OSs. Therefore we needed to find a good approach in order to manage these problems and not to face with memory leakages.

6.3.2 Performance limitations

Talking about CPU computational power there is a huge gap between the two architectures.. Not considering the evident difference for what regarding the space availability for the CPU package, the main difference is that ARM architectures are focused on reducing the power consumption, a crucial characteristic in mobile devices. A technical report analyzed the execution time of a benchmark using a cross-platform version of WebKit showed that the tested ARM A9 architecture had a ratio of 5.8 (normalized time) compared to the tested x86 architecture.[12]

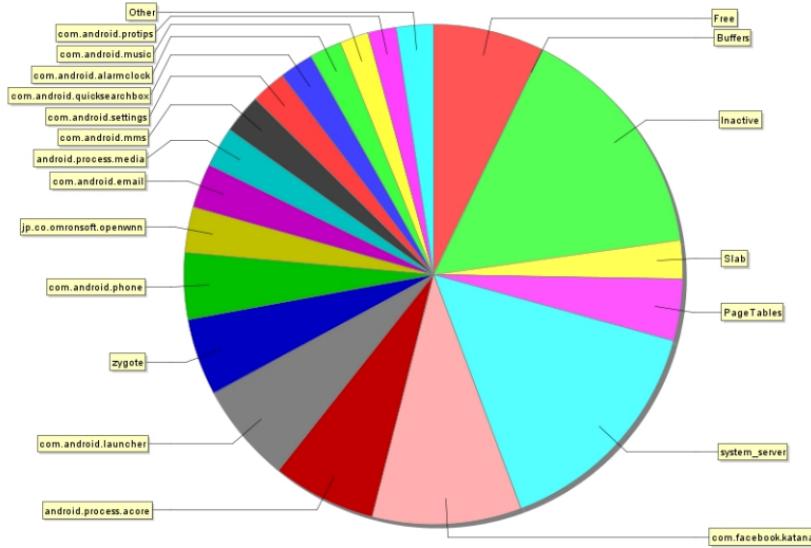


Figure 6.4: Memory reserved to an app in Android.

6.4 ECG baseline wander

For a right interpretation and identification of physiological and pathological phenomena, a low error is required in the ECG signal. However often ECG recordings are distorted by two main artifacts:

- high-frequency noise caused by electromyogram induced noise, power line interferences, or mechanical forces acting on the electrodes;
- baseline wander (BW) that may be due to respiration or the motion of the patients or an instrument fault (Figure 6.5).

These artifacts strongly limit the utility, readability and interpretation of recorded ECGs. In ECG enhancement, the goal is to separate the valid ECG from the undesired artifacts so that we can extract a signal that allows an easy visual interpretation.[13] The first class of artifacts is mostly solved within the acquisition device at hardware low-level through filters and a proper settings. In this project we will deal with the second class of artifact which can be managed at software level.

6.4.1 State of Art

In recent years, ,two methods have often been applied to remove BW:

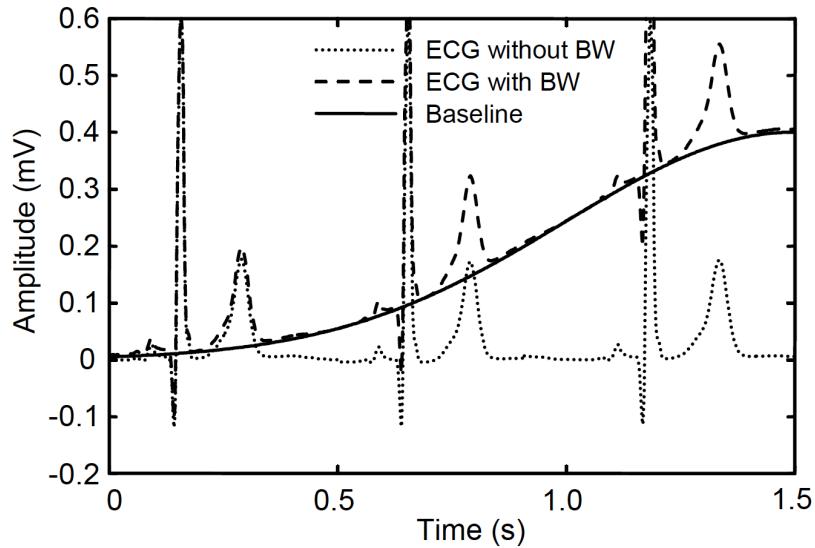


Figure 6.5: The baseline wander artifact in the ECG signal.

- Polynomial fitting: in order to assess the baseline, it uses a polynomial interpolation. The baseline is fit from some fiducial points that are determined from P-R intervals, whereas these fiducial points are difficult to accurately locate before noise is removed from the ECG signal. As result this approach is ineffective if the ECG signal is contaminated by noise.[14]
- High-pass filtering: to implement this type of filtering, usually a moving average filter[15] and wavelet translation[16]. This approach however would unavoidably introduce distortions in various parts of the ECG signal, especially in the ST segment due to the spectra of the ST segment that overlaps the spectra of BW.

Simple Moving Average Filter

This type of filtering is defined as

$$y(n) = \frac{1}{2N+1} \sum_{i=-N}^N x(n+i) \quad (6.1)$$

where $x(n)$ and $y(n)$ are input signal and output signal of the moving average respectively, and N specifies the observation window length equal to $2N + 1$. So the

baseline can be estimated as

$$z(n) = x(n) - y(n) \quad (6.2)$$

where $z(n)$ is the output signal of the high-pass filter.

Distortion using Simple Moving Average Filtering

Obviously the baseline values estimated from the P-R segment (between 0.3 s and 0.6 s) are very close to real baseline values, while the baseline values estimated from segments including QRS complex and T wave are far away from the real baseline. Therefore, if an observation window covers some sample points with extreme amplitudes, an ECG signal would be distorted.

CPU intensive operation

Many methods are known in the literature that perform better and reduce the error in the ECG signal due to the filtering, like wavelet package translation filter[17], or the statistical weighted moving average filter[18]. Defined the normalized root mean square error ($NRMSE$) and maximum error (ME) as

$$NRMSE = \sqrt{\frac{\sum_{i=1}^L (ecg_{in}(i) - ecg_{out}(i))^2}{\sum_{i=1}^L ecg_{in}(i)}} \quad (6.3)$$

$$ME = \max_{i=1,2,\dots,L} |ecg_{in}(i) - ecg_{out}(i)| \quad (6.4)$$

The figure 6.6 shows a comparison between the methods. Of course these advanced methods' computational complexity is really high. Considering that this type of filtering is supposed to be used during a real-time acquisition in our project, they are obviously prohibitive. Hence we need to find out new approaches which are less CPU intensive.

6.5 Signal visualization

The problem of an efficient ECG signal visualization shows up when the application has to plot samples of an acquisition device with a rate of 500Hz. Because of this the app has to be able to draw $500 - 1$ lines (connecting two samples) for each mapped

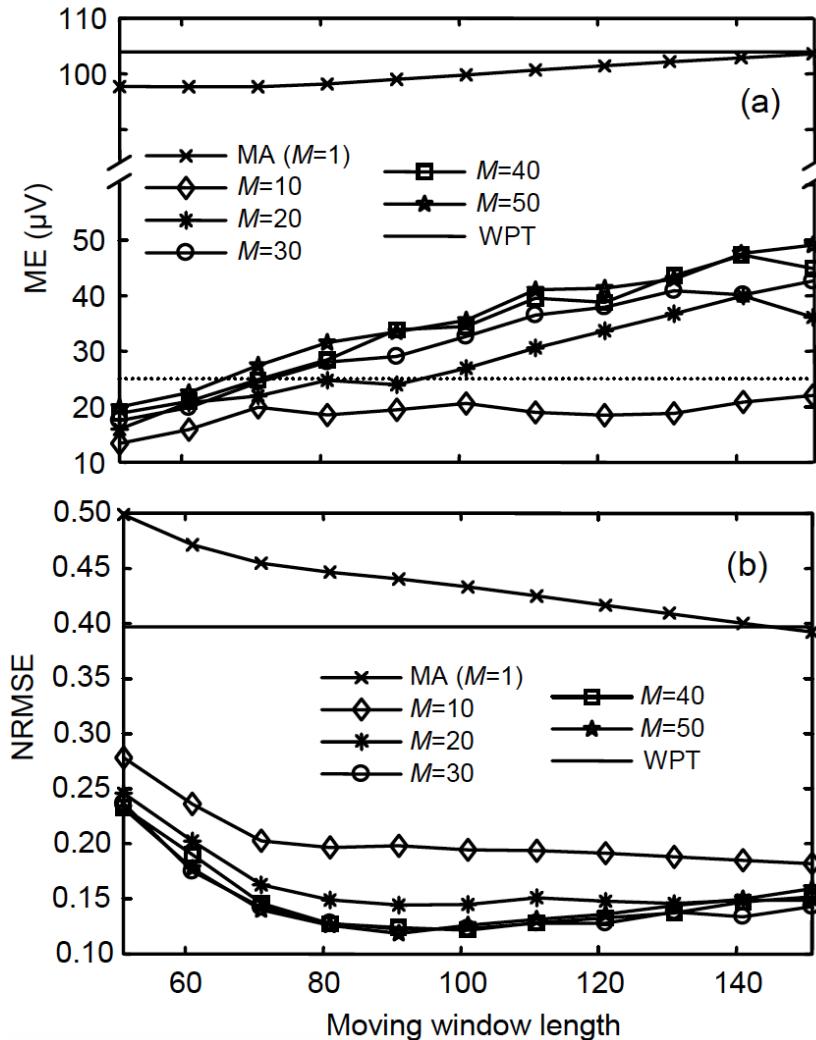


Figure 6.6: Relation between distortion and moving window length for wavelet package translation (WPT), traditionally used MA, and our proposed high-pass filter based on a statistical weighted moving average (SMA).

(a) Maximum error (ME) vs. moving window length;

(b) Normalized root mean square error (NRMSE) vs. moving window length. M is the number of sub-bounds. SMA is the same as MA when M=1.

second in the ECG paper grid. Taking into account the standard size of a second in the standard ECG paper of 5mm, and the fact that the visualization size has to respect this measure, if we consider a device screen of 4.6 inch in portrait, we have around 11 seconds of samples to visualize. Hence we have $(500 - 1) * 11 = 5489$ lines to plot. If we are supposed to plot ideally all the ECG derivation (12), the number of total lines become $5489 * 12 = 65868$. Now, this is the number of lines we

need to plot in order to cover the total visible window screen space. If we want to have a good display smoothness we must guarantee at least 30FPS (frame rate per seconds), ideally 60 frame/s (FPS). Assuming a minimum of 30 FPS we have a total of $65868 * 30 = 1976040$.

These numbers gives us an idea of the computational effort that is spent only to plot the samples.

Chapter 7

Solution choices

After a long research period of time and direct experience on developing mobile application using the most known cross-platform (Xamarin) and hybrid (Cordova Phonegap) solution, we decided to go native. This decision was based mostly on the needs and the strict performance requirements related to the project. Thanks to a native implementation we can achieve better results in term of performance and in term of user experience of the application. At the choice of a native platform we picked Android because it is the most spread mobile OS over mobile devices and it is open source even if mostly maintained by Google.

7.1 Android Platform

Android is a mobile operating system (OS) currently developed by Google, based on the Linux Kernel and designed primarily for touchscreen mobile devices such as smartphones and tablets. Android has the largest installed base of all operating systems of any kind. Android has been the best selling OS on tablets since 2013, and on smartphones it is dominant by any metric.[19] We have chosen to develop ECG-ira firstly on this platform because of it is widely spread over the world and its nature of being from an open source software made it stable and widely supported. The Android OS programming language is Java which is one of the most known and used OOP language for application and web development.

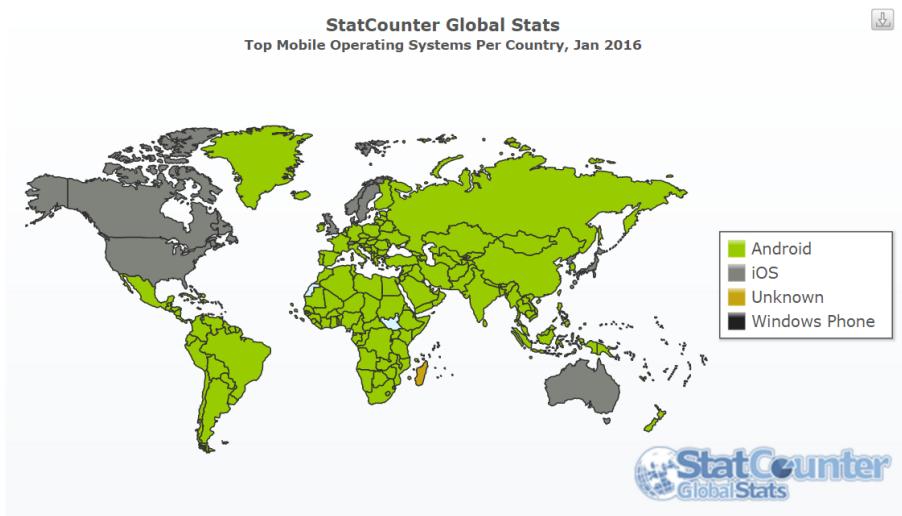


Figure 7.1: Top Mobile Operating Systems Per Country, Jan 2016. Statcounter.com

7.2 Why native?

The advantages of cross-platform (and hybrid) solution is mostly connected to code maintenance and to faster development, because most of the written codes stands for all the supported OS at building phase. Business logic and data structure can be easily share among the many OS for example by using Xamarin we can write unique code using C# (c-sharp) and abstract the business logic, the web services and the database management independently from the specific platform we want to target. Most of the time it is just a matter of working on different user interface, one for any supported OS. All of this look awesome and it is, but when it comes to performance metrics, custom interfaces and user experience, here we meet its weaknesses and limitations. ECG-ira main goal is to build up an usable and stable mobile medical application and to fulfill it we needed to exploit the native platform in order to achieve the best performance and the best user experience. As Android is the most spread mobile OS over smartphones and tablets it results in an obvious pick.

We believe native may not be the best pick for any kind of application. The choices has to be done according to the project requirements and goals. Pitfall for going native is the long development time and a deep (if not full) knowledge of that specific platform.

7.3 Android concurrency exploitation

As we have seen previously, a considerable computational effort is required to the app, especially during record acquisition. For this reason, the only way in order to guarantee a reasonable performance was to exploit the concurrency mechanisms available in the chosen platform. This decision will obviously increase the complexity of the execution: analyzing the execution of a single-threaded application is relatively simple because the order of execution is known. In multi-threaded applications, it is a lot more difficult to analyze how the program is executed and in which order the code is processed.

In the following paragraphs we will start from the basic mechanisms provided by Java language, and we will then analyse the ones, given available by the Android OS, that we have chosen to use in our application.

7.3.1 Thread Overview

Software programming is all about instructing the hardware to perform an action. The instructions are defined by the application code that the CPU processes in an ordered sequence, which is the high-level definition of a thread. From an application perspective, a thread is execution along a code path of Java statements that are performed sequentially. A code path that is sequentially executed on a thread is referred to as a task, a unit of work that coherently executes on one thread. A thread can either execute one or multiple tasks in sequence.

Thread execution

A thread in Java machine is represented by `java.lang.Thread`. It is the most basic execution environment in Android that executes tasks when it starts and terminates when the task is finished or there are no more tasks to execute; the alive time of the thread is determined by the length of the task. `Thread` supports execution of tasks that are implementations of the `java.lang.Runnable` interface. An implementation defines the task in the `run` method:

```
private class MyTask implements Runnable {
    public void run() {
        int i = 0; // Stored on the thread local stack.
    }
}
```

```
}
```

All the local variables in the method calls from within a run() method—direct or indirect—will be stored on the local memory stack of the thread. The task’s execution is started by instantiating and starting a Thread:

```
Thread myThread = new Thread(new MyTask());
myThread.start();
```

On the operating system level, the thread has both an instruction and a stack pointer. The instruction pointer references the next instruction to be processed, and the stack pointer references a private memory area—not available to other threads—where thread-local data is stored. Thread local data is typically variable literals that are defined in the Java methods of the application.

A CPU can process instructions from one thread at a time, but a system normally has multiple threads that require processing at the same time, such as a system with multiple simultaneously running applications. For the user to perceive that applications can run in parallel, the CPU has to share its processing time between the application threads. The sharing of a CPU’s processing time is handled by a scheduler. That determines what thread the CPU should process and for how long. The scheduling strategy can be implemented in various ways, but it is mainly based on the thread priority: a high-priority thread gets the CPU allocation before a low-priority thread and receive more execution time with respect to low-priority threads. The execution of two concurrent threads can be done in java just declaring two Thread objects and then starting them by calling the method Thread .start():

```
Thread T1 = new Thread(new MyTask());
T1.start();
```

7.3.2 Threads in Android

In Android there are basically three thread types:

- **UI thread** (or main thread): it is started on application start and stays alive during the lifetime of the application process. The UI thread is the main thread of the application, used for executing Android components and updating the UI elements on the screen. If the platform detects that UI updates are attempted

from any other thread, it will promptly notify the application by throwing a `CalledFromWrongThreadException`. This harsh platform behaviour is required because the Android UI Toolkit is not thread safe, so the runtime allows access to the UI elements from one thread only.

- **Binder threads:** they are used for communicating between threads in different processes. Each process maintains a set of threads, called a thread pool, that is never terminated or recreated, but can run tasks at the request of another thread in the process. These threads handle incoming requests from other processes, including system services, intents, content providers, and services.
- **Background threads:** All the threads that an application explicitly creates are background threads. This means that they have no predefined purpose, but are empty execution environments waiting to execute any task. The background threads are descendants of the UI thread, so they inherit the UI thread properties, such as its priority. By default, a newly created process does not contain any background threads. It is always up to the application itself to create them when needed.

The UI thread is the most important thread, but it gets no special scheduling advantage compared to the other threads—the scheduler is unaware of which thread is the UI thread. Instead, it is up to the application to not let the background threads interfere more than necessary with the UI thread.

7.3.3 Thread communication in Android

In multithreaded applications, tasks can run in parallel and collaborate to produce a result. Hence, threads have to be able to communicate to enable true asynchronous processing.

The most common thread communication use case in Android is between the UI thread and worker threads. Hence, the Android platform defines its own message passing mechanism for communication between threads. The UI thread can offload long tasks by sending data messages to be processed on background threads. The message passing mechanism is a nonblocking consumer-producer pattern, where neither the producer thread nor the consumer thread will block during the message handoff.

The message handling mechanism in android is implemented with the following classes:

- **android.os.Looper**: A message dispatcher associated with the one and only consumer thread.
- **android.os.Handler**: Consumer thread message processor and the interface for a producer thread to insert messages into the queue. A Looper can have many associated handlers, but they all insert messages into the same queue.
- **android.os.MessageQueue**: Unbounded linked list of messages to be processed on the consumer thread. Every Looper—and Thread—has at most one MessageQueue.
- **android.os.Message**: Message to be executed on the consumer thread.

The mechanism is summarized in the figure 7.2.

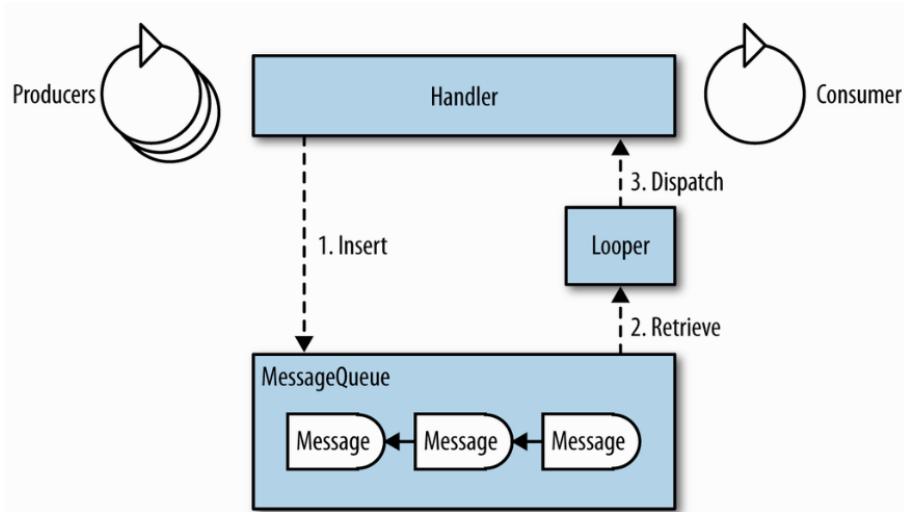


Figure 7.2: Overview of the message-passing mechanism between multiple producer threads and one consumer thread.

- **Insert**: The producer thread inserts messages in the queue by using the Handler connected to the consumer thread.
- **Retrieve**: The Looper runs in the consumer thread and retrieves messages from the queue in a sequential order.
- **Dispatch**: The handlers are responsible for processing the messages on the consumer thread. A thread may have multiple Handler instances for processing

messages; the Looper ensures that messages are dispatched to the correct Handler.

7.3.4 HandlerThread

Now we will describe a component that we have heavily exploited in our application, and as you will see in implementation details section, it represents the base of two main operations: the ECG signal drawing during both acquisition and record opening, and the ECG signal reading during record opening.

HandlerThread is a thread with a message queue that incorporates a Thread, a Looper, and a MessageQueue. It is constructed and started in the same way as a Thread. Once it is started, HandlerThread sets up queuing through a Looper and MessageQueue and then waits for incoming messages to process:

```
HandlerThread handlerThread = new HandlerThread("HandlerThread");
handlerThread.start();

mHandler = new Handler(handlerThread.getLooper()) {
    @Override
    public void handleMessage(Message msg) {
        super.handleMessage(msg);
        // Process messages here
    }
};
```

There is only one queue to store messages, so execution is guaranteed to be sequential, and therefore thread safe. The HandlerThread sets up the Looper internally and prepares the thread for receiving messages.

Here is a simple example of an implementation:

```
public class MyHandlerThread extends HandlerThread {
    private Handler mHandler;
    public MyHandlerThread() {
        super("MyHandlerThread", Process.THREAD_PRIORITY_BACKGROUND);
    }
    @Override
    protected void onLooperPrepared() {
```

```

super.onLooperPrepared();
mHandler = new Handler(getLooper()) {
    @Override
    public void handleMessage(Message msg) {
        switch(msg.what) {
            case 1:
                // Handle message
                break;
            case 2:
                // Handle message
                break;
        }
    }
};

public void publishedMethod1() {
    mHandler.sendEmptyMessage(1);
}

public void publishedMethod2() {
    mHandler.sendEmptyMessage(2);
}
}

```

Lifecycle

A running HandlerThread instance processes messages that it receives until it is terminated. A terminated HandlerThread can not be reused. To process more messages after termination, create a new instance of HandlerThread. The lifecycle can be described in a set of states:

- **Creation:** The constructor for HandlerThread takes a mandatory name argument and an optional priority for the thread:

```

HandlerThread(String name)
HandlerThread(String name, int priority)

```

The name argument simplifies debugging, because the thread can be found

more easily in both thread analysis and logging. The priority argument is optional and should be set with the same Linux thread priority values used in `Process.setThreadPriority`. The default priority is

`Process.THREAD_PRIORITY_DEFAULT,`

the same priority as the UI thread, and can be lowered to

`Process.THREAD_PRIORITY_BACKGROUND`

to execute non-critical tasks.

- **Execution:** The HandlerThread is active while it can process messages; i.e., as long as the Looper can dispatch messages to the thread. The dispatch mechanism is set up when the thread is started through `HandlerThread.start` and it is ready when either `HandlerThread.getLooper` returns or on the `onLooperPrepared` callback. A HandlerThread is always ready to receive messages when the Handler can be created, as `getLooper` blocks until the Looper is prepared.
- **Reset:** The message queue can be reset so that no more of the queued messages will be processed, but the thread remains alive and can process new messages. The reset will remove all pending messages in the queue, but not affect a message that has been dispatched and is executing on the thread:

```
public void resetHandlerThread() {
    mHandler.removeCallbacksAndMessages(null);
}
```

The argument to `removeCallbacksAndMessages` removes the message with that specific identifier. `null`, shown here, removes all the messages in the queue.

- **Termination:** A HandlerThread is terminated either with `quit` or `quitSafely`, which corresponds to the termination of the Looper. With `quit`, no further messages will be dispatched to the HandlerThread, whereas `quitSafely` ensures that messages that have passed the dispatch barrier are processed before the thread is terminated. You can also send an interrupt to the HandlerThread to cancel the currently executing message:

```
public void stopHandlerThread(HandlerThread handlerThread) {
    handlerThread.quit();
    handlerThread.interrupt();
}
```

A terminated HandlerThread instance has reached its final state and it cannot be restarted.

7.3.5 Thread Pools

A thread pool is the combination of a task queue and a set of worker threads that forms a producer-consumer setup. Producers add tasks to the queue and worker threads consume them whenever there is an idle thread ready to perform a new background execution. Therefore, the worker thread pool can contain both active threads executing tasks, and idle threads waiting for tasks to execute. There are several advantages with thread pools over executing every task on a new thread (thread-per-task pattern):

- The worker threads can be kept alive to wait for new tasks to execute. This means that threads are not created and destroyed for every task, which compromises performance.
- The thread pool is defined with a maximum number of threads so that the platform is not overloaded with background threads—that consume application memory—due to many background tasks.
- The lifecycle of all worker threads are controlled by the thread-pool lifecycle.

ThreadPoolExecutor

A thread pool's behaviour is based on a set of properties concerning the threads and the task queue, which you can set to control the pool. The properties are used by the ThreadPoolExecutor to define thread creation and termination as well as the queuing of tasks. The configuration is done in the constructor,

```
ThreadPoolExecutor executor = new ThreadPoolExecutor(
    int corePoolSize,
```

```
int maximumPoolSize,  
long keepAliveTime,  
TimeUnit unit,  
BlockingQueue<Runnable> workQueue);
```

where:

- **corePoolSize**: The lower limit of threads that are contained in the thread pool. Actually, the thread pool starts with zero threads, but once the core pool size is reached, the number of threads does not fall below this lower limit. If a task is added to the queue when the number of worker threads in the pool is lower than the core pool size, a new thread will be created even if there are idle threads waiting for tasks. Once the number of worker threads is equal to or higher than the core pool size, new worker threads are only created if the queue is full.
- **maximumPoolSize**: The maximum number of threads that can be executed concurrently. Tasks that are added to the queue when the maximum pool size is reached will wait in the queue until there is an idle thread available to process the task.
- **keepAliveTime**: Idle threads are kept alive in the thread pool to be prepared for incoming tasks to process, but if the alive time is set, the system can reclaim noncore pool threads. The alive time is configured in TimeUnit, the unit the time is measured in.
- **workQueue**: An implementation of BlockingQueue that holds tasks added by the consumer until they can be processed by a worker thread. Depending on the requirements, the queuing policy can vary.

ScheduledThreadPoolExecutor

This is an extension of the ThreadPoolExecutor, which can schedule commands to run after a given delay, or to execute periodically. This class will be really useful in our application because of its capability of scheduling task at a fixed rate through the method:

```
scheduleAtFixedRate(Runnable command, long initialDelay,  
long period, TimeUnit unit)
```

where:

- **command**: the task to execute
- **initialDelay**: the time to delay first execution
- **period**: the period between successive executions
- **unit**: the time unit of the initialDelay and period parameters

7.3.6 AsyncTask

As the name indicates, an `AsyncTask` is an asynchronous task that is executed on a background thread. The only method you need to override in the class is `doInBackground()`. Hence, a minimal implementation of an `AsyncTask` looks like this:

```
public class MinimalTask extends AsyncTask {
    @Override
    protected Object doInBackground(Object... objects) {
        // Implement task to execute on background thread.
    }
}
```

The task is executed by calling the `execute` method, which triggers a callback to `doInBackground` on a background thread:

```
new MinimalTask().execute(Object... objects);
```

When an `AsyncTask` finishes executing, it cannot be executed again—i.e., `execute` is a one-shot operation and can be called only once per `AsyncTask` instance, the same behaviour as a `Thread`.

In addition to background execution, `AsyncTask` offers a data passing mechanism from `execute` to `doInBackground`. Objects of any type can be passed from the initiating thread to the background thread. This is like `HandlerThread`, but with `AsyncTask` you do not have to be concerned about sending and processing `Message` instances with a `Handler`.

In the common case where you want to execute a task in the background and deliver a result back to the UI thread, `AsyncTask` shines; it is all about handling the flow of preparing the UI before executing a long task, executing the task, reporting progress of

the task, and finally returning the result. All of this is available as optional callbacks to subclasses of the AsyncTask, which look like this:

```
public class FullTask extends AsyncTask<Params, Progress, Result> {
    @Override
    protected void onPreExecute() { ... }
    @Override
    protected Result doInBackground(Params... params) { ... }
    @Override
    protected void onProgressUpdate(Progress... progress) { ... }
    @Override
    protected void onPostExecute(Result result) { ... }
    @Override
    protected void onCancelled(Result result) { ... }
}
```

This implementation extends the AsyncTask and defines the arguments of the objects that are passed between threads:

- **Params:** Input data to the task executed in the background.
- **Progress:** Progress data reported from the background thread (i.e. from doInBackground) to the UI thread in onProgressUpdate.
- **Result:** The result produced from the background thread and sent to the UI thread.

All callback methods are executed sequentially, except onProgressUpdate, which is initiated by and runs concurrently with doInBackground. Figure 7.3 shows the lifecycle of an AsyncTask and its callback sequence. The different steps are:

1. Create the AsyncTask instance.
2. Start execution of the task.
3. First callback on the UI thread: onPreExecute. This usually prepares the UI for the long operation—e.g., by displaying a progress indicator on the screen.
4. Callback on a background thread: doInBackground. This executes the long-running task.

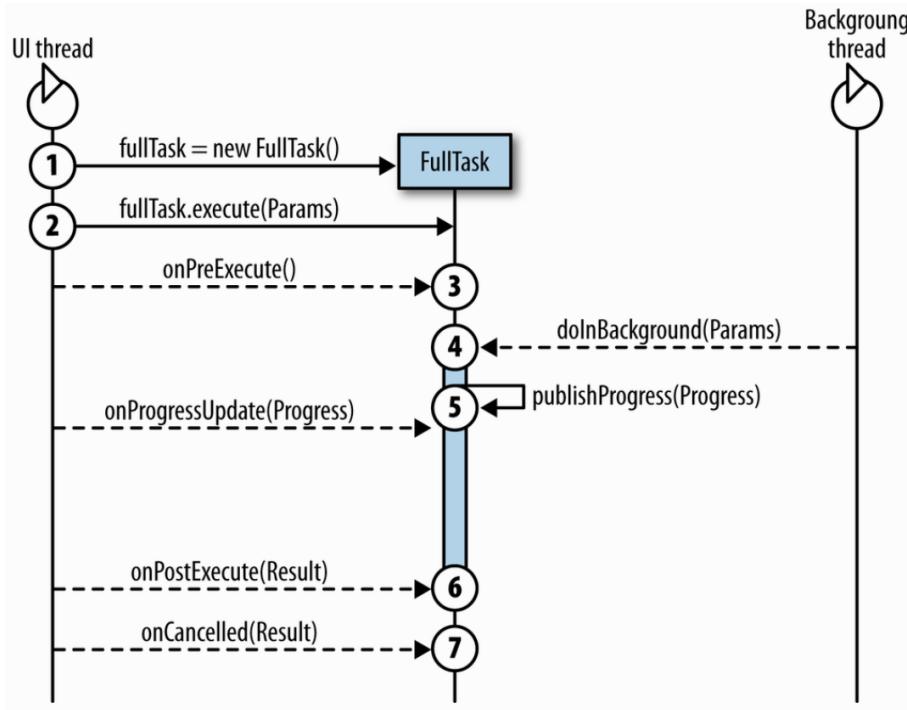


Figure 7.3: The execution lifecycle of `AsyncTask`.

- Report progress updates from the `publishProgress` method on the background thread. These trigger the `onProgressUpdate` callback on the UI thread, which typically handles the update by changing a progress indicator on the screen. The progress is defined by the `Progress` parameter.
- The background execution is done and is followed by running a callback on the UI thread to report the result. There are two possible callbacks: `onPostExecute` is called by default, but if the `AsyncTask` has been cancelled, the callback `onCancelled` gets the result instead. It is guaranteed that only one of the callbacks can occur.

The progress update mechanism solves two use cases:

- Displaying to the user how the long-running operation is progressing, by continuously reporting how many of the total tasks are executed.
- Delivering the result in portions, instead of delivering everything at the end in `onPostExecute`. For example, if the task downloads multiple images over the network, the `AsyncTask` does not have to wait and deliver all images to the UI

thread when they are all downloaded; it can utilize publishProgress to send one image at the time to the UI thread. In that way, the user gets a continuous update of the UI.[20]

7.4 Baseline wander solutions

In order to solve this problem, taking into account all the related problematic, we have chosen two kinds of approach. Each one will be described afterwards. The first approach will be always active, and will consist in the dynamic calculation of samples vertical axis during drawing iteration. The second is the usage of a simple moving average filter that could be activated in the app settings for the acquisition phase.

7.4.1 Adaptive vertical displaying

We decided to hold this type of solution active by default in order to maintain a solid and versatile way to overcome the worst scenario caused by a strong baseline wander. The approach works like this: at each frame drawing, we have a visible window of samples, with a length dictated from the space available on the device's screen, that we have to plot. Given that window, we know that we have to fit as many samples as we can, inside the space dedicated to that signal, the signal ECG strip. In a normal scenario, the signal will be aligned to a baseline, and so we can easily plot all the samples inside the relative strip. However in some other scenarios, it could happen that because of movement of the patient, the signal could immediately drop down. For this reason, the signal could easily go out of the available vertical space. To avoid this, we have to provide a mechanism in order to hit these cases and accordingly respond. We do this by not fixing the vertical baseline of the ECG strip, and leave it dynamic. Therefore this baseline will go up and down according to the position of the values inside the samples window. So every time, before we redraw the updated window on the screen, we compute the baseline of the signal at that moment by computing the mean of all samples. After that, we can shift the signal to plot up or down, trying to include the majority of samples on the screen. An example of the mechanism is showed in the figure 7.4. As you can see, starting from the second 23, a patient movement caused the baseline wander artifact, causing a drop of the D1 signal. The app by applying the dynamic displaying, it computed the new baseline of the signal, represented by the mean of all samples, and shifted all sample window

accordingly. In the figure, the variation of the vertical axis is represented by Δy . In this way we can overcome this type of scenario, avoiding the appliance of any kind of filtering. The latter, depending on the technique, can result in some percentage of error on the filtered ECG signal.

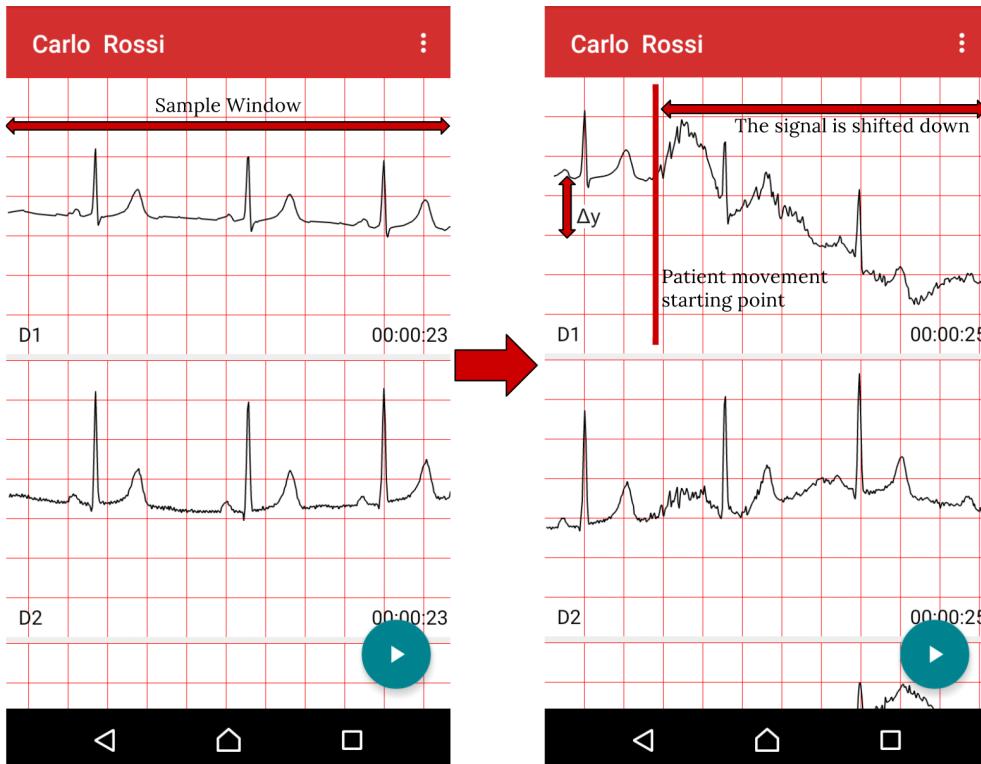


Figure 7.4: The dynamic displaying result after a patient movement, causing the baseline wander artifact.

7.4.2 Moving average filter

As discussed in the requirements section, many types of filtering are known to overcome to the baseline wander artifact. Some of them are capable to reduce at minimum the error on the output signal. Nevertheless this comes with a cost in computational effort, and so there is the need to mediate between the type of filtering and its related complexity. Unfortunately in our case, putting in all the required operation, especially during real-time acquisition, where the app has to hold the bluetooth channel for transmission, interpret the transmitted signal, write to a file, derive the missing ECG lead, and plot at a reasonable rate the acquired signal, we had very limited

computational availability to spend in any type of filtering in order to remove the eventual baseline wander. Therefore we decided to apply the most basic type of filtering during acquisition that is the simple moving average filtering. Given that, we are conscious about the possible distortion introduced by this filter, that's why we decided to:

- Take a reasonable size for the moving average window (two seconds at least), inasmuch if it is true that as much as the window size grows, the effectiveness of the filter decreases, by doing this, we can keep the error rate low.
- Keep the filter deactivated by default, so that the doctor will decide when will be opportune to use it.

7.5 Custom View Drawing

Now we will talk about the most costly operation performed about our application. This was the thing on which we have spent lots of work, investigating all the problematic and different possibilities that we had in order to make the best possible implementation choice. We have mentioned earlier the computational effort that needs to be spent on signal drawing and for this reason we tried all the possible ways in order to discard the bad implementation, always having performance in our mind.

7.5.1 Custom Libraries

This was our first trial: we tried to find some libraries that could have permitted us to avoid an implementation from scratch of our drawing classes. For sure this possibility was the easiest possible. Given that our signal was not so different from other types of signals, as could be interpreted as a generic function plotted on a two dimensional system, we had quite sure that we could have found a nice plotting library and avoid useless implementation. Actually, we were able to find some well-realized libraries for handling plottings, but all of them clashed with one characteristic that we was seeking for: the customization of the rendered views, as we have said previously, has to mimic as much as possible the ECG paper on the look and also respect the required standard sizing of the same. Therefore, for this reason, we needed to discard this solution, given that some libraries permitted us to have very good displaying performance.

7.5.2 Hardware Accelerated Drawing (GPU)

Another solution, and potentially the best one, was to exploit the graphic hardware acceleration. Android is possible by using OpenGL ES, a subset of the OpenGL API designed for embedded system. The use of OpenGL can move all the graphics computations to the GPU, and so freeing up precious computing resources on the CPU. But unfortunately, in Android the usage of these libraries is not so integrated: they are written in hardware native code, which is C. This characteristic, while could of course guarantee the best performance[21], introduce a misalignment with the language used for developing Android application, which is Java. As a result, it is a common belief that the usage of OpenGL ES in Android is quite painful, forcing many developers to switch to better alternative libraries and frameworks, like Unity, LibGDX, Cocos2D or others. Putting aside all the problematic related to its implementation's effort, we decided to try OpenGL ES for the drawing part. With much surprise, this led us to an unexpected result: the results in performances were quite beneath the performance achieved using the CPU also for the application drawing. This relies on the fact that, as said earlier, the OpenGL libraries on Android are not so integrated, and developer are required to represent datatypes of the C language in the Java language. This may not seem problematic, but in a situation where all ECG samples need to be represented as classes, holding their coordinates in the ECG paper space in a corresponding matrix, and at each draw update there is the need to reallocate all the samples matrix, causing a remapping to the C data types, the performance improvements are quickly drop out. For this reason, we realized that our best chance was to relying on the CPU also for the drawing, and trying as much as possible to optimize the algorithms in order to achieve the best performances.

7.5.3 Not hardware Accelerated Drawing (CPU)

Having underlined the downsides of the previous solution, we decided to spend all our energies to implement the best possible drawing code, relying on the mechanisms provided by Android for drawing operations using the CPU. This is possible using Canvas.

Android Canvas provides the developer with the ability to create and modify 2D images and shapes. Moreover, the Canvas can be used to create and render our own 2D objects as this class provides various drawing methods to do so. Canvas can also

be used to create some basic animations such as frame-by-frame animations or to create certain Drawable objects such as buttons with textures and shapes such as circles, ovals, squares, polygons, and lines.[22]

As we mentioned in the chapter about Android concurrency exploitation, all applications run on a single thread in Android. All instructions run in a sequence on the UI thread, meaning that the second instruction will not start unless the first one is finished. The UI thread as it is responsible for drawing all the objects or views on the screen and processing all events, such as screen touches and button clicks. Now the problem is that, if we have two operations scheduled to run in the same default thread or UI thread and the first operation takes too long to finish, the system will ask the user to forcibly close the application or wait for the process to complete. This scenario is called ANR (Application Not Responding).

Given that we wanted also to provide scrolling of the different ECG leads in our app, we could not hold this computational effort on the UI thread, which as result of drawing operations, would be blocked. Therefore we decided to assign all the drawing operations on different threads created ad-hoc. The number of threads dedicated to drawing will be decided at run time by the application, depending on the device availability, and thus allowing device scalability.

The use of Canvas permitted us to achieve both reasonable performances, after a deep code optimization, and high customization of the rendered view.

Chapter 8

System architecture

The entire system is based on an acquisition device named ZEcg, and a native mobile application on the Android Platform. Our focus and main effort were on the developing of the mobile application fulfilling all the requirements. The acquisition device was instead developed and designed by Crespi Alessandro and Ulisse Pizzagalli during their thesis work at the Politecnico of Milan. [23]

8.1 Acquisition device

ZEcg is composed by the following different modules:

1. OVP: used to protect the patient from high voltage or voltage leakage.
2. LFP EMI Filter: anti-aliasing RC filter used to remove noises due to high frequencies.
3. RLD: Active electrode driver for the right leg
4. WCT: Derivator used to compute the precordials (Wilson Central Terminator)
5. PGA: 8 gain amplifier with programmable inputs
6. ADC: 8 analog-digital 16 bit 8KSa/s converter.
7. MCU: microcontroller
8. Bluetooth: bluetooth module for data transmission

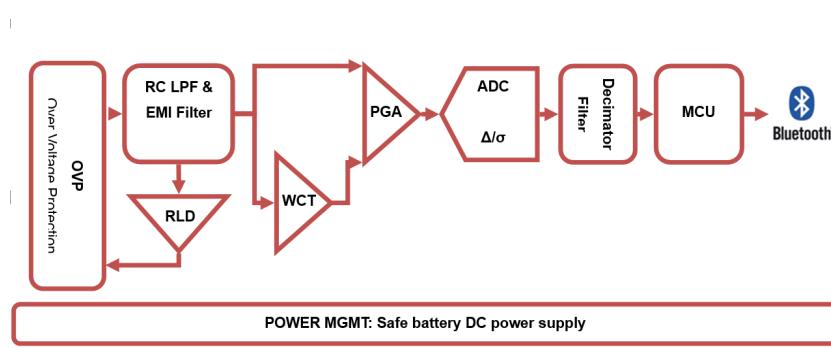


Figure 8.1: ZEcg device block diagram including all the modules components.

9. Power MGMT: to manage battery recharge and stabilizer

The core component is the Texas Instrument system on chip ADS1198. This chip has 8 input bipolar channels, representing the 8 clinical leads.

The channels 1 and 2 produce the lead I and lead II. Channel 1 measures the potential difference between the electrode RA(-) and the electrode LA(+), the channel 2 the difference between RA(-) and LL(+). Lead III and the augmented leads are obtained from a combination of lead I and lead II at software level.

V1, V2,...,V6 are computed as difference between the respective electrode and the signal related to the negative value of the WCT (Wilson Central Terminator).

The WCT signal comes from the average between RA, LA and LL and it's connected to the channels 3,4,5,6,7,8. Each channel is amplified using a programmable gain (PGA) and a CMRR, before being converted into digital. For more details about the entire device architecture we invite you to read the thesis of our colleagues Crespi Alessandro and Ulisse Pizzagalli[23] who designed and developed ZEcg.

8.1.1 Mobile app

The mobile application we developed for this thesis work was implemented having in mind all the user best interface design principles and the best practice starting from the planning and design phase to the final coding phase. As this application was designed for Android OS, we strictly followed Google specific standards and procedures.

We made use of Android Studio as IDE (strongly suggested by Google as main IDE to develop Android native applications). Starting from 2013 the development of native Android application moved from Eclipse + Android Plugin to Android Studio. The

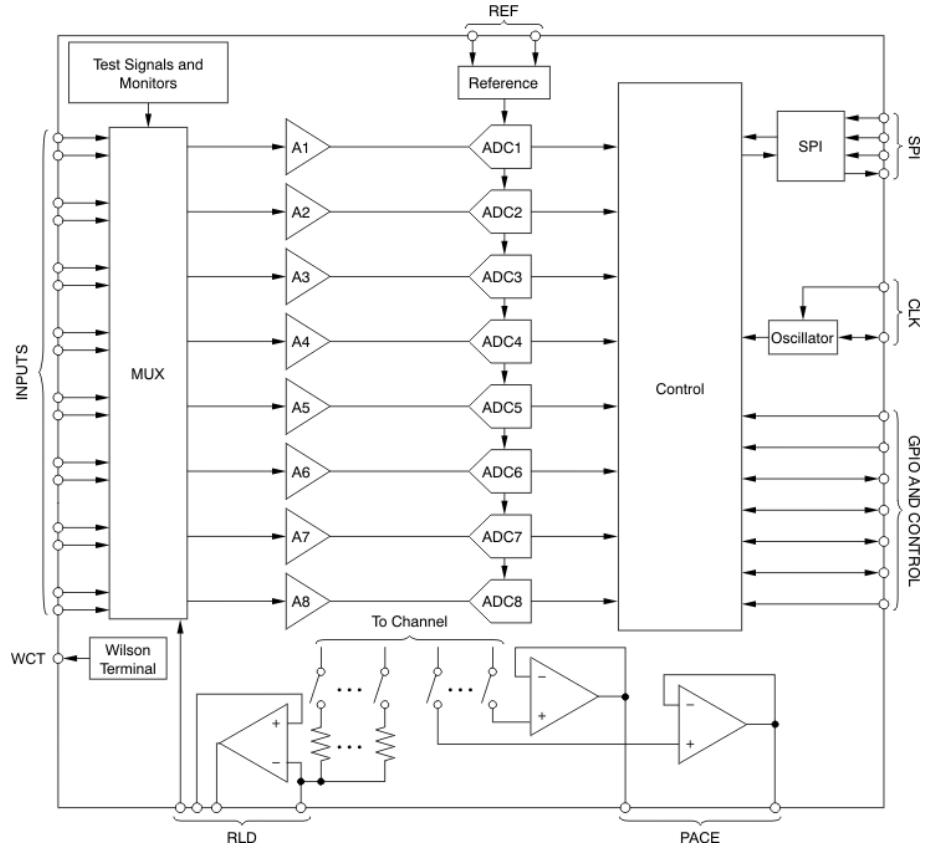


Figure 8.2: ADS1198 functional diagram showing the 8 channels representing the 8 leads used during the ECG record acquisition.

entire project building system has changed and moved to use the Gradle building system[24]. The entire process during project building to compilation can be resumed in the following image: The most important steps along the building process are:

1. The Android Asset Packaging Tool (aapt) takes the application resource files, such as `AndroidManifest.xml` file and the XML files for the Activities, and compiles them. A `R.java` file is produced so that all the resources can be easily accessed within your application.
2. The aidl tool converts any `.aidl` interfaces into Java interfaces.
3. The Java compiler will compile the `R.java` and `.aidl` files generating the `.class` files.
4. The Dex tool will convert the `.class` files to Dalvik bytecode. Any third libraries and `.class` files included in the project build will be also converted into `.dex` files

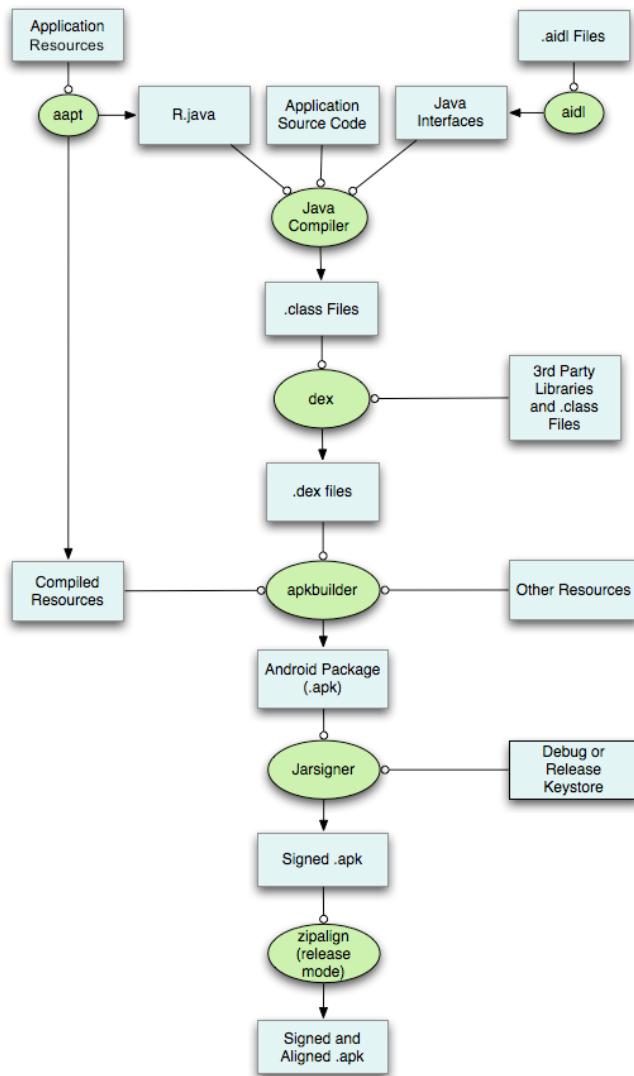


Figure 8.3: Android Build system process. How and which component are involved during an android application build and compilation.

so that they can be later packed into the final .apk.

5. All non-compiled resources(such as images), compiled resources, and the .dex files are sent to the apkbuilder tool that will output the .apk file.
6. Once the .apk is built, it must be signed with either a debug or release key before it can be installed to a device.
7. To reduce the size of the .apk and to decrease the memory usage for releasing mode the zipalign tool is launched.

We split the application functionalities into different packages. The packages content

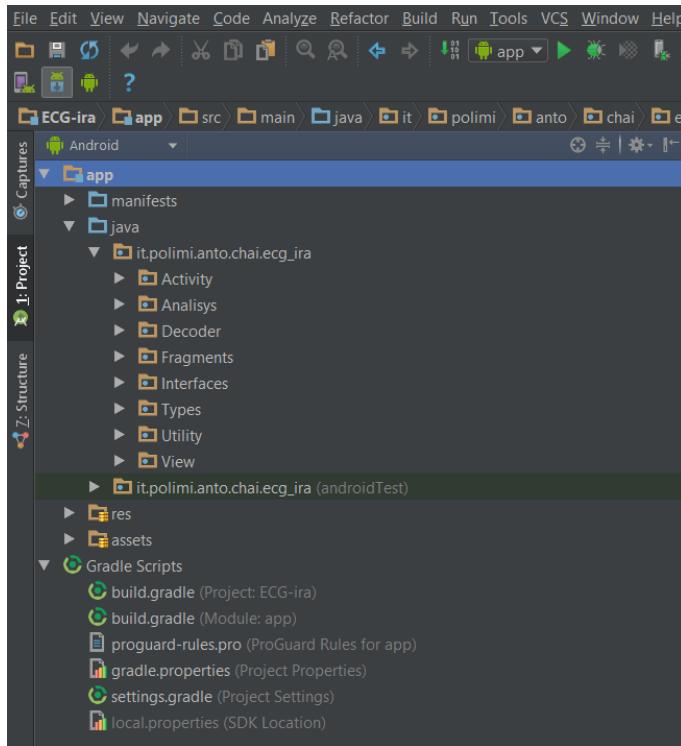


Figure 8.4: ECG-ira project structure and packages inside Android Studio IDE.

are divided as follow:

- Activity: it contains all the Activity used inside the Application.
- Analysis: All the java classes used for the analysis of the ECG records; they include the Neural Network, the QRSDetector, the PWaveDetector and some signal Filters implementations.
- Decoder: this package contains the java classes in charge to decode and parse the .hea and .dat files.
- Fragments: contains some fragments used inside the application, for example the ones used to show the chronograms (Istogram, Tacogram, ST graphs).
- Interfaces: all interfaces declarations to abstract the specific application behaviour with more general one.

- Utility: some utility class such as the in application FileManager.java class, and the Adapters used within the application (for example to show ListViews or RecycleViews).
- View: it contains all the classes working directly on the View and the custom view themselves such as the View used to show the ECG grid.

The key points of this application are based on three main concepts:

- Code reuse and maintainability
- Maximize performance depending on the device constraints (CPU, memory, storage capacity)
- Usability in term of interface and interaction

Code reuse and maintainability

The most important classes representing the core of the application functionalities are all abstracted through interfaces or in case of the classes related to the View through a set of parameters. This makes the application highly customizable and easily to extend. For example if in a next future there will be a new ECG data format, it can be possible to give the app the capability to read and parse such a format just by implementing the interfaces and writing the specific code to parse such a new data format.

A concrete example is given in our code by observing that both the MITDataReader and the ZEcgDataReader extends the abstract class DatReader, so the next format reader has just to extend it as well. We know that MIT-BIH format is completely different from the ZEcg format, that why specific code to parse the file content to extract the signal is mandatory.

Any new format reader has just to override the method:

```
public void readAndAddSample(int numOfSample, int addPosition) {  
}
```

More concrete details about the effective implementation will be exploited later on in the next chapters.

Maximize performance depending on the device constraints

We do well know about the great number of constraints due to the huge number of different devices with different hardware on. We decide to make our application available starting from Android API 16 (the minimum sdk API recommended by Google to support). To overcome the problem of the difference devices and Android OS version starting from Jelly Bean (API 16) we took maximum advantages from the device hardware by splitting the thread jobs in between the maximum number of cores available. In the class CpuInfoExtractor we discover the hardware capabilities (number of cores) and according to it we split the calculus between a certain number of threads. More the cores, more the threads we can take advantage of.

From an interface point of view instead, we are forced to depend on the density of pixel of the device screen and its size, but at the same time to accomplish our requirements related to achieve a perfect grid of squares of centimeters. We found a way to always achieve the same dimensions of square independently by the screen size of the devices by retrieving and taking in account the display exact pixels per inch size in the X dimension. Then, we computed the number of pixels needed to achieve the right sizing. In this way, we solved the problem of having same dimensions and so metrics on different devices. On the other hand, if this method is quite functional and device independent, on devices with different screen size (width for example) the number of grid cells may vary from just a few on small screen, to many of them on tablets. This is a direct consequence due to the difference in number of pixels and the pixels size itself (some are squared but most of them are rectangles).

Usability in term of interface and interaction

One mobile application is usable if it does what the user expects it to do when interacting with it.

We followed all Google guidelines in term of user experience using native view and patterns. We took advantage of the last API features, such as RecycleViews instead of the old classic ListViews. We made use of the button ripple effects available starting from Lollipop (Android API 21), but at the same time we provided to the older Android OS version the selector effect that still gives a nice response to the user interaction with command and buttons. We made use of the typical android Preference Settings so familiar to Android users and most important we always inform

the user about the operations going on so he never feels lost inside the application.

Chapter 9

Implementation details

In this chapter we will focus on the implementation details of the main components and operation, trying to give a clearer idea of the mechanism under the hood of the app. We will start by explaining the general idea and purpose of each included component and their main functions. Then we will move to the sequence flow section, where we will put all pieces together in order to explain their iterations and give a picture of what is happening in the different use cases.

9.1 Main components

We decided to include only, which are for us, the main components of the app, and exclude all the others that are not so relevant in order to understand the operations flow. For the first two, we will start from an abstract idea of the component, identifying the main functionalities that it has to provide.

If we think about our application in an high level, it usually will need to acquire the ECG signal from a source, decode and modify it, and finally display it in the screen. This operations are splitted in two components:

- A Data Source that will take care about acquisition and decoding
- A Display that will take care about all the operation related to drawing ECG signal

9.1.1 Data Sources

As introduced before, the Data Source components will take care about receiving the signal from a source, decode and modify it, and pass this information to a Display. Going down of a level, we know that ECG signal in our app will come from:

- a saved ECG record file (.dat file), during record opening;
- the acquisition device zecg, through a bluetooth connection during real-time acquisition.

From this last distinction, we could go down to another level on both possibilities:

- in saved ECG record opening, there are potentially different format for encoding the ECG signal in a file;
- the acquisition could support many acquisition device, and not only zecg.

For these reasons it can be very effective to abstract as much as possible all the characteristics that different acquisition devices or file formats share. In this way we can treat them on the same way in many situations.

Actually we were not required to support many different acquisition devices, and the high parametrization of the device in the app settings was quite enough for our goal. For this reason the acquisition will be handled by the class ZEcgReceiver, which does not implement any interfaces or extends any abstract classes.

We wanted instead to set that abstraction level for which regards different ECG record format. For this reason we will start from the description of the SampleSource interface, and then we will talk about its implementing abstract class, the DatReader, and finally how to deal with different ECG record formats.

SampleSource

SampleSource is an interface that is useful for the saved ECG record opening. It will include methods for getting sample from a source, starting and pausing a ECG playing animation (two types: scrolling or oscilloscope), adding listeners and others. An overview of all methods is given in the figure 9.1. All its methods are:

- start(): it starts the SampleSource and prepare all resources that it needs;
- quit(): it stop the SampleSource and destroys all resources used;

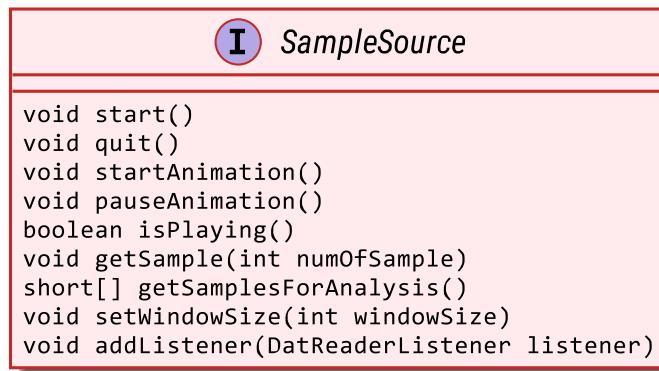


Figure 9.1: Class diagram of the SampleSource interface.

- `startAnimation()`: it starts the animation of the ECG signal (scrolling or oscilloscope);
- `pauseAnimation()`: it pauses the animation of the ECG signal;
- `isPlaying()`: it returns true if the animation is active, false otherwise;
- `getSamples(numOfSamples)`: it gets from the source `numOfSamples` samples and stores them;
- `getSamplesForAnalysis()`: it returns D2 lead, used for the ECG analysis;
- `setWindowSize(windowSize)`: it sets the size of the display window size, specified in number of samples, that depend on the device screen size and orientation;
- `addListener(listener)`: it registers a listener that will be notified if the SampleSource changes animation state.

The SampleSource interface that we introduced will be the basis for the following class description, the DatReader, the core of our record ECG opening.

DatReader

The DatReader is an abstract class that handles all the most sensitive operations for which regarding ECG record opening. It implements the SampleSource interface, described before, and permits to support many kind of ECG record format, with the only need of overriding one of their methods. There are many things under the hood of this class, so they will be described one by one.

The first characteristic of this class is that, all its operation are executed asynchronously from the UI thread: this is done by extending the class HandlerThread, described in the previous chapter. By doing this, we can exploit all the functionalities provided by this class, as the Looper and Message passing mechanism. So the DatReader once started, it will stay watching at its MessageQueue, though its Looper, waiting for a new message. The messages in this class will be basically the number of samples that it will need to read. To read the ECG samples, the DatReader will hold a stream to the record .dat file (file containing all samples). Reading constantly from a file, therefore avoiding the allocation in memory of the record, permits us to maintain the app memory low.

But who is sending these messages to the DatReader? Well, there are basically two scenarios:

- The user that is using the app scroll the ECG paper with its finger, causing the app to load the right number of samples depending on the movement size and direction;
- The user activated the ECG scrolling animation, and so some samples need to be loaded after a certain period, depending on the ECG record rate.

Focusing on the first scenario, the entity that will handle ECG visualization and user inputs will be the SampleDisplay, which will be described in details in the next section. And so it will respond to the user scroll, calculating the right number of samples and request them to the DatReader, which is viewed as a SampleSource. The request will be represented by a Message, containing all useful information, like of course the number of samples to load, but also others.

There is a problematic where we want to provide scrolling of the ECG paper in both direction, while we are reading constantly from a file: we need to be able to move the file pointer in both direction, not a trivial functionality. The class that we rely on is the RandomAccessFile class: using this class we can always retrieve the actual position of the file pointer and then move it with the method seek(). But this is not the only problematic. We need to take into account that the application will always show on screen a window of the ECG record, with a dynamic size, dependent on the space available on screen. Therefore, the file pointer will refer to a specific point of this window. This point can be the tail or the head of the window, namely the point where the samples were added last time.

Consider the following example, represented in the figure 9.2: an user opens a ECG record, the DatReader will be created and started, and will load all the samples need in order to fill the device screen. At this moment, it will have read all the samples sequentially, from the first one to the last required one, from the file. Thus, the file pointer is referring to the tail of the ECG record window. After that, the user starts scrolling to the right, in order to see the following two seconds of the record. So the DatReader continues to read sequentially the following samples in the file. But now suppose that the user decides to change its scrolling direction and go back. We cannot only read the file backwards from the last pointer position, because those samples will be already loaded and visualized in the sample window. What we need to do is to skip the entire window, and start reading from that point. So after the left scroll, the pointer will be moved to the head of the window.

To make this mechanism possible, it is clear that we have to store the previous

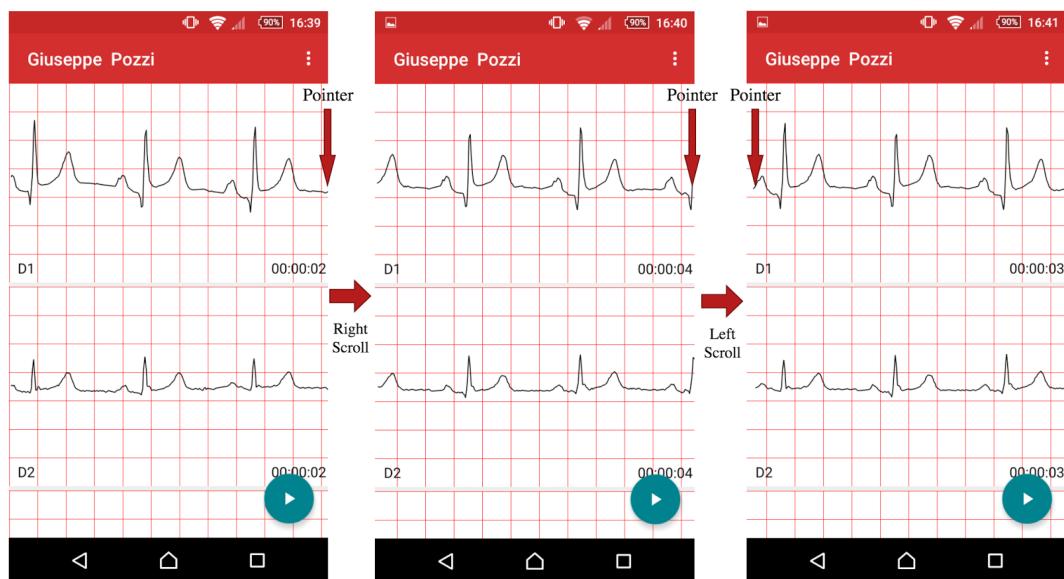


Figure 9.2: The movement of the file pointer in the DatReader, after the user change direction of scrolling the ECG paper.

reading direction, stored in the variable `prevDirection`. If the new direction will be different, the DatReader will move the file pointer, skipping the window. So we need also to store the size of the window: this is done with the variable `windowSize`. After each file reading operation, the DatReader will send to the SampleDisplay all new samples, specifying if add them to the tail of the window or to the head. The iteration between SampleSource and SampleDisplay in the user scrolling scenario

is summarized in the figure 9.3. Before being sent to the SampleDisplay, the read

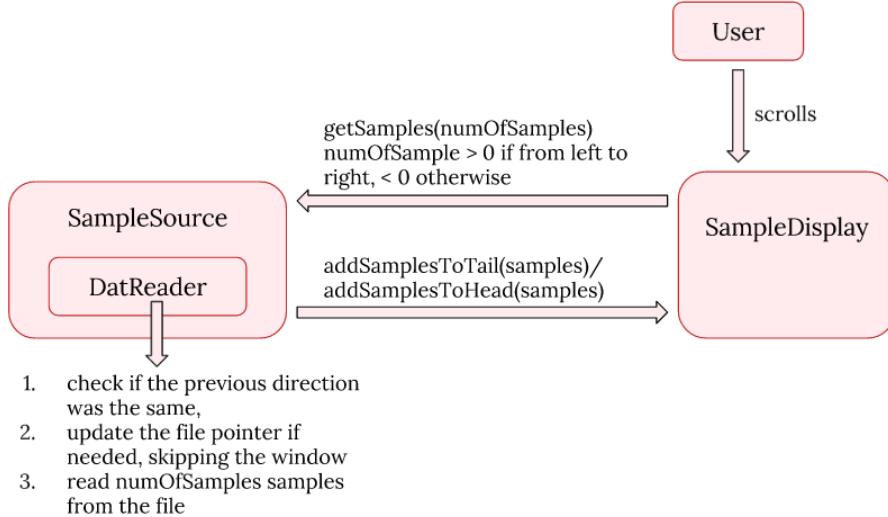


Figure 9.3: The iteration between SampleSource and SampleDisplay after the user scroll the ECG paper.

samples are reconverted to mV, by dividing all values by the gain of the signal. In this way, we can work with universal values during plotting, as different ECG record can have different gain.

Now we will move to the second scenario. Here, the user has activated the ECG scrolling animation, by clicking on the play button, which will be binded to the SampleSource, causing the call of startAnimation(). In this case, there will be another entity that will request samples from the SampleSource. The crucial aspect is that, it will continuously requests new samples, but doing so respecting the rate of the original ECG record. This is not a trivial problem, as some requests could be delayed because of another operation that is not finished. So these requests need to be concurrent, reliant on a mechanism to guarantee the right rate, with a good precision. The solution was the adoption of a ScheduledThreadPoolExecutor, a pool of thread that we have described in the chapter about concurrency. Thanks to its method:

```

scheduleAtFixedRate(Runnable command, long initialDelay,
long period, TimeUnit unit)

```

we can program all the requests in the pool by:

```

scheduledThreadPoolExecutor.scheduleAtFixedRate(new Runnable() {
    @Override

```

```

public void run() {
    getSample(1);
}
}, 0, 1000000 / rate, TimeUnit.MICROSECONDS);

```

having a precision of microseconds. Being this type of operation not so expensive, we initialized our pool with only one thread.

We have omitted that the method `getSamples(numOfSamples)` doesn't have an implementation for getting the samples from the source, but it just acts like a request to the SampleSource: by executing it, the implementation in the DatReader will create a message and enqueue it to its MessageQueue. The DatReader will continuously read messages from its MessageQueue and sooner or later, it will find that message, read the sample from the source, and update the SampleDisplay with the new sample. An summarization of the mechanism is showed in the figure 9.4. Furthermore, in the figure 9.5, you can find an overview of the DatReader class, including all variables introduced in this section.

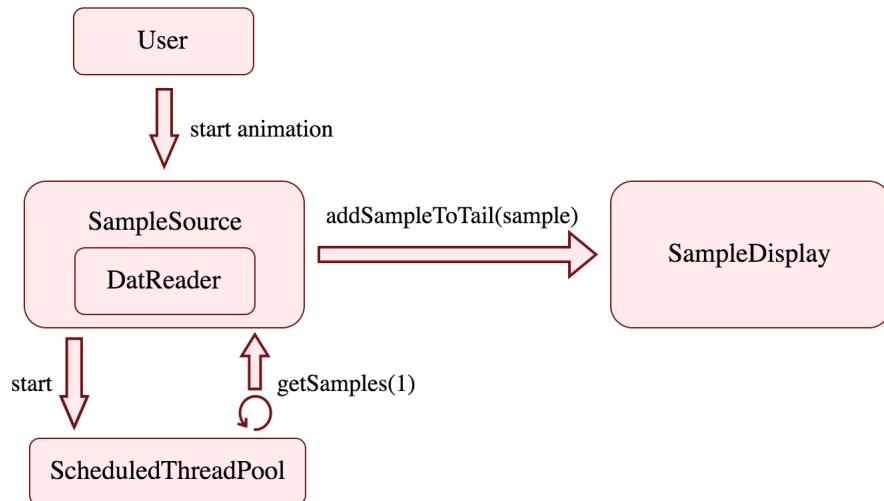


Figure 9.4: The animation mechanism of DatReader using ScheduledThreadPoolExecutor.

ZEcgReceiver

Now we describe the component that will be used during the ECG real-time acquisition. It will communicate with the BluetoothSerial class and interact, similarly to the DatReader, with the SampleDisplay. As the DatReader, it extends the HandlerThread,

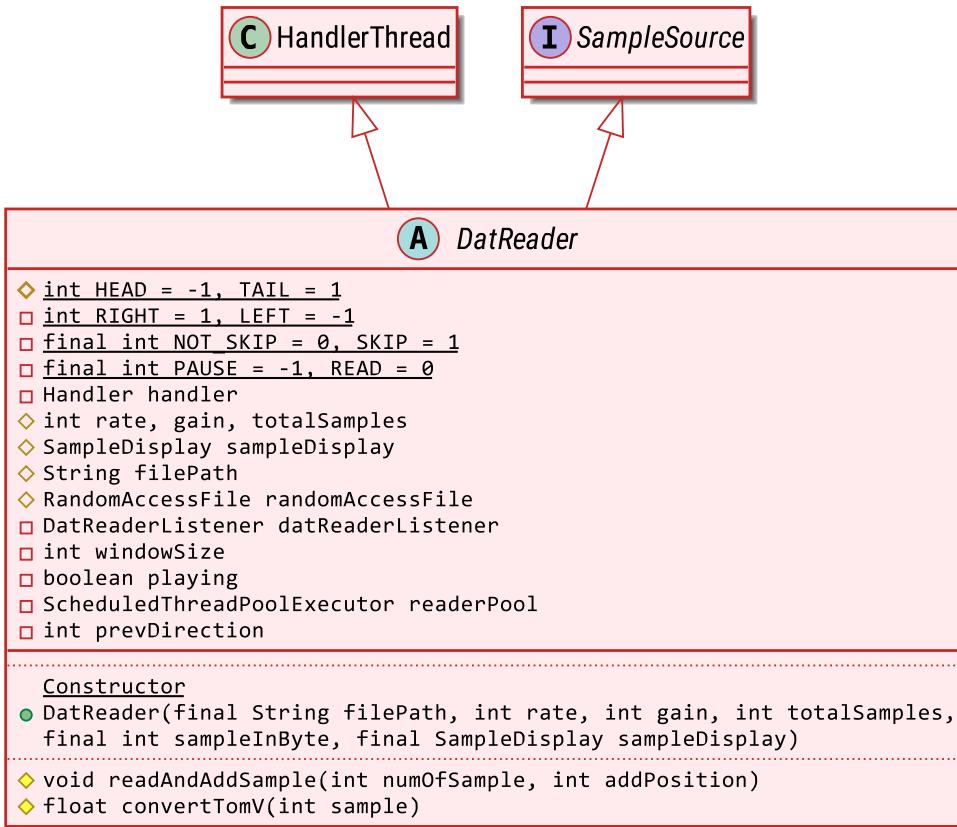


Figure 9.5: Class diagram of the `DatReader` class.

therefore it will have a personal Looper and MessageQueue. It will also have an acquisitionBuffer that will use to store partial data received from `BluetoothSerial`. So, what it will do is waiting for new messages on its MessageQueue, and react on a new message. A message coming from `BluetoothSerial` will contain a buffer of data received from the acquisition device through a bluetooth connection. At a new message receiving , `ZEcgReceiver` will read the buffer in the message and append it to its acquisitionBuffer. After that, it will scan its acquisitionBuffer and read all possible samples inside, and free it from all decoded samples. All new decoded samples are then sent to the `SampleDisplay` for the visualization.

`ZEcgReceiver` will also handle many other operation as:

- It will store a `heartRateBuffer` of the last sample received, with a size of 10 seconds, used to compute the real-time heart rate, to be notified to the `SampleDisplay`;
- It will include the calculation of the Moving Average Filter, if enabled, in order

to remove the baseline wander artifact;

- The data received from the acquisition device will also include information about its battery level. For this reason, it will parse it and send it to the SampleDisplay;
- It will write additional information in the file .hea of the ECG record.

9.1.2 Display

Before we described all the entities behind the generation of the data that will be sent then to the visualization part of the application. Now we will talk about all the components that make ECG visualization possible. Equivalently to the previous section, we will start describing the general interface the will represent the visualization component, the SampleDisplay, moving later to the component that will implement it.

SampleDisplay

The have already mentioned many time the SampleDisplay component when talking about data sources. This interface represent the general entity that expose some entry point to add new samples and handle ECG signal visualization. The figure 9.6 shows an overview of all methods. Methods description:

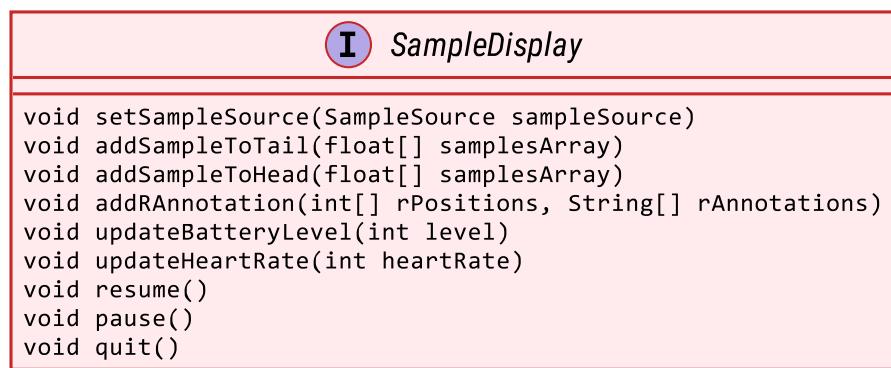


Figure 9.6: Class diagram of the SampleDisplay interface.

- `setSampleSource(SampleSource sampleSource)`: it binds the SampleSource passed as parameter;

- `addSampleToTail(float[] samplesArray)`: it adds all samples in the passed array to the tail of the visualized samples window;
- `addSampleToHead(float[] samplesArray)`: it adds all samples in the passed array to the head of the visualized samples window;
- `addRAnnotation(int[] rPositions, String[] rAnnotations)`: it stores all the R position (of the QRS complex), with their relative annotations coming from the ECG signal analysis;
- `updateBatteryLevel(int level)`: it updates the visualized battery level;
- `updateHeartRate(int heartRate)`: it updates the visualized heart rate in BPM;
- `resume()`: it resumes the ECG signal visualization;
- `pause()`: it pauses the ECG signal visualization;
- `quit()`: it will close and free all the components used for visualization.

SignalScrollView

For our visualization of the ECG signal we wanted to provide a scrollable view of different ECG paper strip, where in each one was displayed a different ECG lead. Hence, we selected as entity implementing the `SampleDisplay` interface, a customized version of the `ScrollView` class, a scrollable view provided in Android. This class will be the entry point for all data sources, and a container for some other components, described in details further on. These contained elements will be the all `SignalSurfaceViews`, where each one represent the view of an ECG lead. The `SignalScrollView` represents also the core for one of the feature described in the previous chapters: the dynamic display scaling. It will do all calculation inside its method `calculateMetrics()`. The goal of the method is calculate the size of the basic block of the standard ECG paper, which is 0.5 cm, in pixel. The problem is that the pixel size varies a lot, being dependent on the device screen pixel density. So for doing this, we need to retrieve the pixel density of the device screen, and then calculate properly the right size of the `blockInPixel` (figure 9.7). The measure used to represent the pixel density is the *dpi* (*dots per inch*), defined as “*the number of individual dots that can be placed in a*

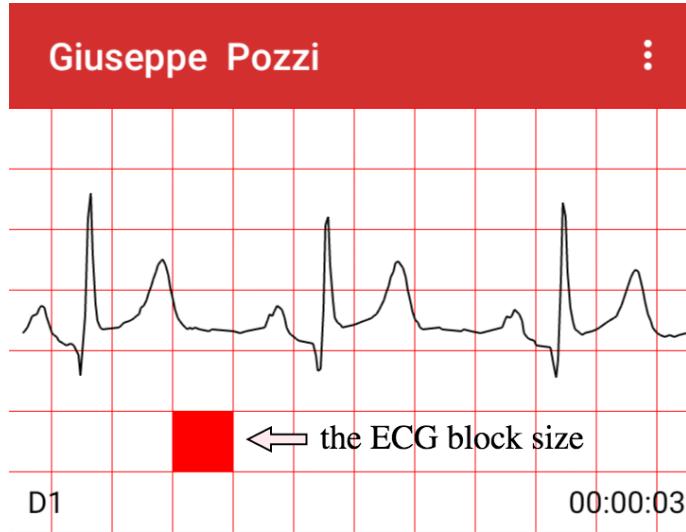


Figure 9.7: The size of the ECG paper block after the computation of the method calculateMetrics().

line within the span of 1 inch (2.54 cm)".[25] The formula for the computation of the blockInPixel will be:

$$\text{blockInPixel} = \frac{\text{dpi} * \text{blockInCm}}{\text{inchInCm}} \quad (9.1)$$

where:

- *dpi*: screen pixel density, dependant on the device specification;
- *blockInCm*: the size of the ECG paper block in centimeters, equals to 0.5 cm;
- *inchInCm*: the size of one inch in centimeters, equals to 2.54 cm.

With this calculation we can ensure a perfect scaling in many device, with a very small error percentage for high pixel density screens ($\sim 240\text{dpi}$ or higher).

The method calculateMetrics() take place in the initialization process of the class, together with the method fillScrollView(int heightInBlock). This method receives the parameter heightInBlock, which comes from the app settings and specifies the height of the ECG paper strip in ECG standard blocks, and create and add inside the SignalScrollView all the SignalSurfaceViews needed, depending on the number of ECG signal leads (one per lead). Moreover, during its initialization, the class will create another important component, the DrawingHelper, that we will describe next. This last will handle all the call of the SampleDisplay methods:

- addSampleToTail(float[] samplesArray);

- addSampleToHead(float[] samplesArray).

For this reason, SignalScrollView will route each adding request to the DrawingHelper, that will take care of storing the visualize samples window and adding all new samples to it.

Furthermore, SignalScrollView will perform a very important optimization for the visualization of the ECG paper strip: with its method pauseNotVisibleView() will be able to select only the strips that are present in the screen, and pause the drawing routine (described in the next sections) for all not visible strips, saving lot of computational power. This method will be triggered every time the user will scroll the view.

Lastly, we need to specify that SignalScrollView will also handle drag events coming from the user when he wants to scroll the ECg paper. But the class will be used as reference class both during real-time acquisition and ECG record opening. The difference in these scenario is that, during real-time acquisition, the dragging need to be disabled, while in the other, enabled. So to achieve this distinction, inside the constructor will be passed the parameter dragEnabled, which will be used to activate or not the overridden method onTouchEvent(MotionEvent ev), that will take care about decode the performed scroll and notify the SampleSource.

As usual, an overview of the class could be seen in the figure 9.8.

SignalSurfaceView

In order to achieve the best performance, we used the best mechanism available in Android for fast and custom drawing. For this purpose we extended the class SurfaceView. As explained in the Android developer site “*The SurfaceView is a special subclass of View that offers a dedicated drawing surface within the View hierarchy. The aim is to offer this drawing surface to an application’s secondary thread, so that the application isn’t required to wait until the system’s View hierarchy is ready to draw. Instead, a secondary thread that has reference to a SurfaceView can draw to its own Canvas at its own pace.*”[26]. We will focus on the thread responsible of drawing on the SignalSurfaceView when will describe the Drawer class. For now, we describe only the main operation of the SignalSurfaceView. This class will be liable of reacting after the creation of the view and on view change event. This last will occur typically after a rotation of the device screen, performed by the user. This will cause for example to change the orientation from portrait to landscape, and so the new SignalSurfaceView

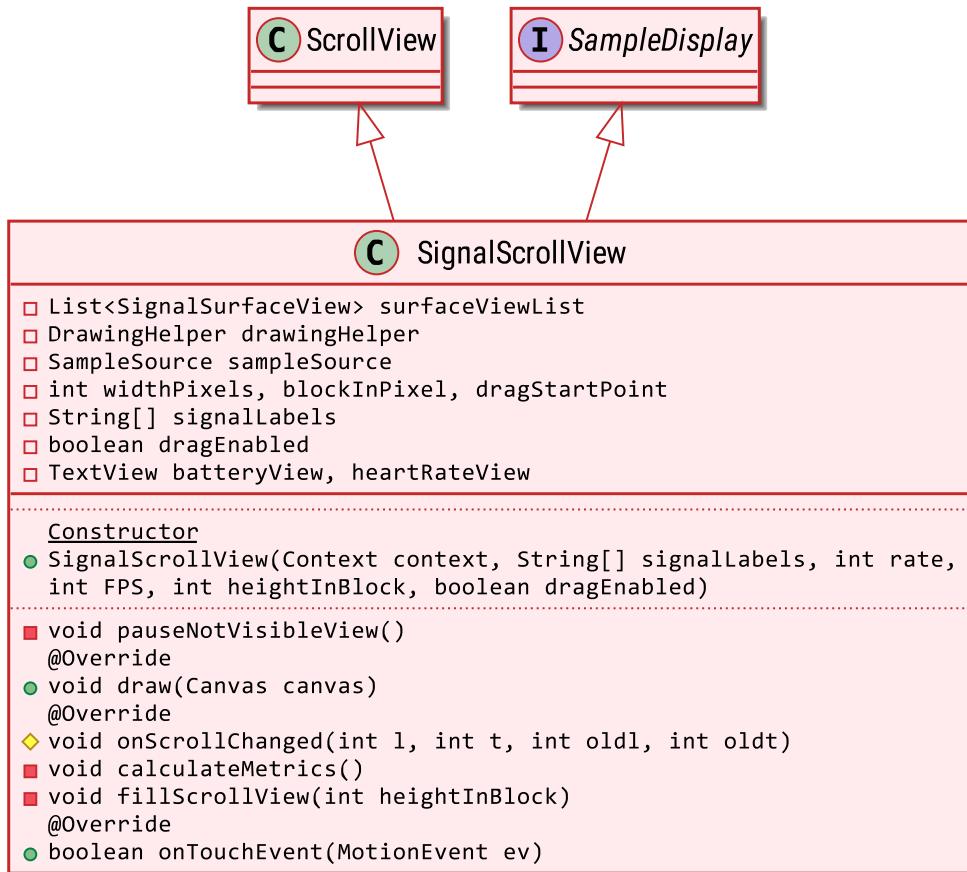


Figure 9.8: Class diagram of the SignalScrollView class.

will be required to fill all the new available space.

DrawingHelper

As the name suggests, this is the helper class for all drawing operation. Many of its methods and variables will be used by the Drawer for handle drawing. It is an abstract class, because we provide two different plotting type: a scrollable drawing and a oscilloscope drawing (figure 9.9). For each type, there is a proper extending class:

- DrawingHelperScroll, for scrollable drawing type;
- DrawingHelperOsc, for oscilloscope drawing type.

Each of these class will store the list of all sampleWindows, containing a sampleWindow for each ECG lead. The sampleWindow will store the samples window that is currently visualized on the screen. The collection type for the sampleWindow is different in the



Figure 9.9: The different ECG visualization types. The one on top is the scrolling drawing, the one in the bottom is the oscilloscope drawing. The time elapsed after the movement of the ECG paper is one second.

two classes, because each one use the most suitable implementation for its update type. For example, the DrawingHelperScroll, will treat the window as a queue: it will always add new sample to the head or to the tail. Thus, it uses a ArrayDeque, an very efficient and fast class when applied to this type of scenario.

During its initialization, with the method `prepare(args ...)` the DrawingHelper creates some Drawer. This last is the class used for drawing on the SignalSurfaceView, and will be described in the next section. The number of Drawer created depends on the computational power availability of the device used: in the method the DrawingHelper will retrieve the number of cores available in the device CPU, and will create accordingly the same number of Drawer (which, you will see, will extend the HandlerThread class). For this reason, with a single core CPU, only a Drawer will be created, and it will be the sole handler for all the SignalSurfaceView (and so all ECG strips). Instead, if we had four CPU core, and for example all the twelve ECG leads, there will be created four Drawer, where each will handle the drawing of three SignalSurfaceView. This characteristic provide a nice performance scalability, dependent on the CPU of the used device.

The DrawingHelper also will handle the lifecycle of all Drawers using:

- `startDrawing(int signalIndex)`, for starting the drawing operation of the Drawer which is the handler for the passed signalIndex;
- `pauseDrawing(int signalIndex)`, for pausing the drawing operation of the Drawer

which is the handler for the passed signalIndex;

- quit(), for stopping all Drawers currently active and destroying them.

Furthermore, the DrawingHelper will provide the helper methods, used by the Drawer, to compute the different drawing layers:

- drawGrid(Canvas c, Paint p): it draws the ECG grid in the passend Canvas using the passed Paint. It uses the variable blockInPixel for the right sizing;
- drawSignalLabelsAndTime(int signalIndex, Canvas c, Paint p): it draws the labels of the time and the ECG leads in the passend Canvas using the passed Paint;
- drawSamples(int signalIndex, Canvas c, Paint p): it draws all the samples inside the sampleWindow as lines in the passend Canvas using the passed Paint. It avoid the drawing of consecutive equals value when plotting the lines, in order to save useless operations;
- drawRAnnotations(Canvas c, Paint p): it highlights all the R peaks by drawing vertical lines on them and with the respective R annotation, in the passend Canvas using the passed Paint. This of course will be active only after the analysis.

In the figure 9.10 you can find the class diagram of the DrawingHelper class.

Drawer

As already mentioned earlier, the Drawer class will be a thread, responsible of drawing on one or more SignalSurfaceViews. It is a private static inner class of the DrawingHelper, and it hold a WeakReference to the DrawingHelper in order to avoid as much as possible memory leaks.[27] It extends the HandlerThread class and its basic functioning is reading and responding to all messages sent to its MessageQueue. The messages will contain the index of the ECG lead that need to be drawn. Actually, there will be always only one message per ECG lead on its MessageQueue: this is because, beside of the lead index, all messages are equals, and incorporate the same drawing routine including all drawing method seen in the previous section. But each message at the end will also include a recursive sending of the next message to

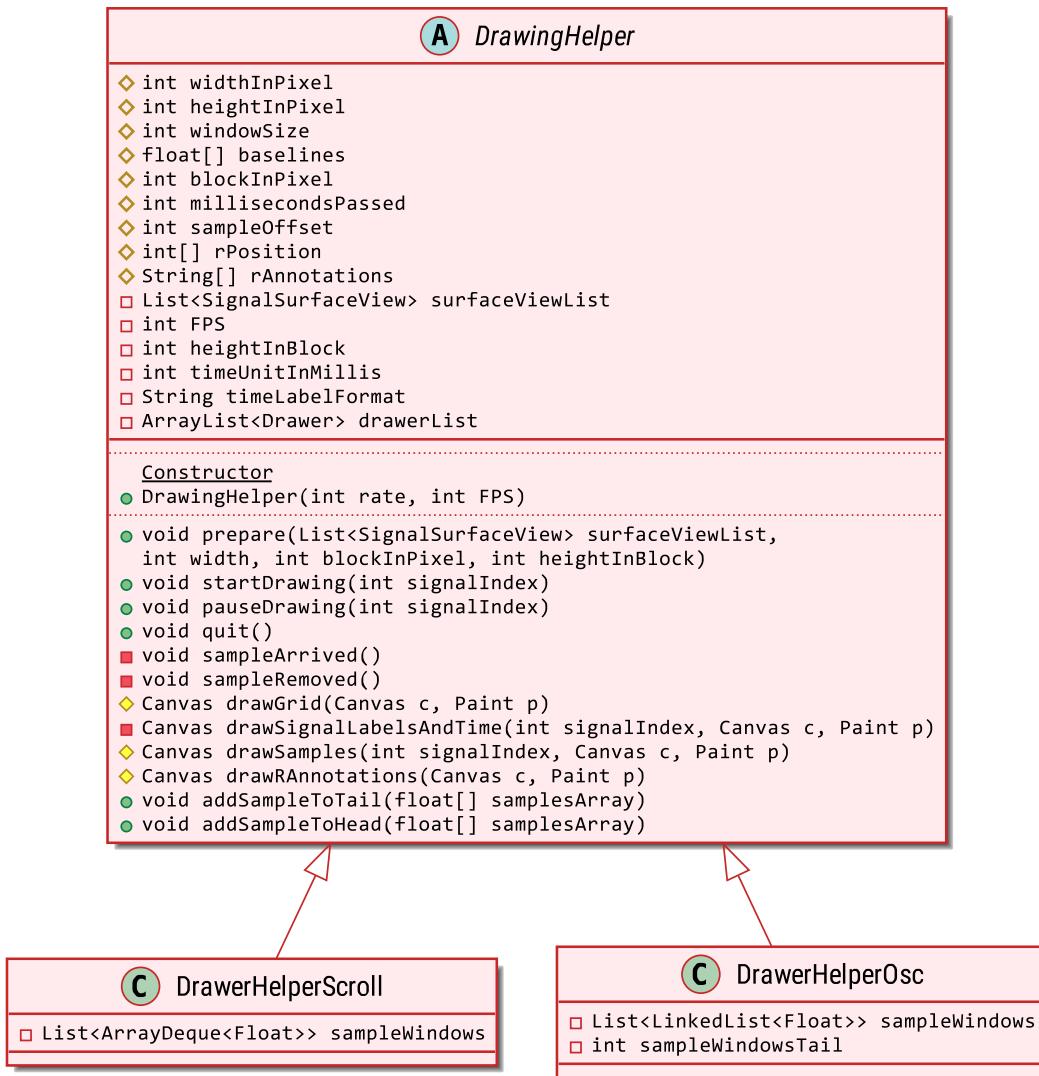


Figure 9.10: Class diagram of the DrawingHelper class.

the MessageQueue, delayed by the time unit dictated from the specified FPS of the application. So each message will cause a sending of the next message. In this way, the Drawer we can continuously draw and update the SignalSurfaceView, at a rate of FPS. The starting of drawing will be managed by the DrawingHelper, which will send the first recursive message. The pausing, manage by the same, will be easily achieved by the removal of all the messages present in the MessageQueue at that moment. In the figure 9.11 is showed the class diagram of the Drawer.

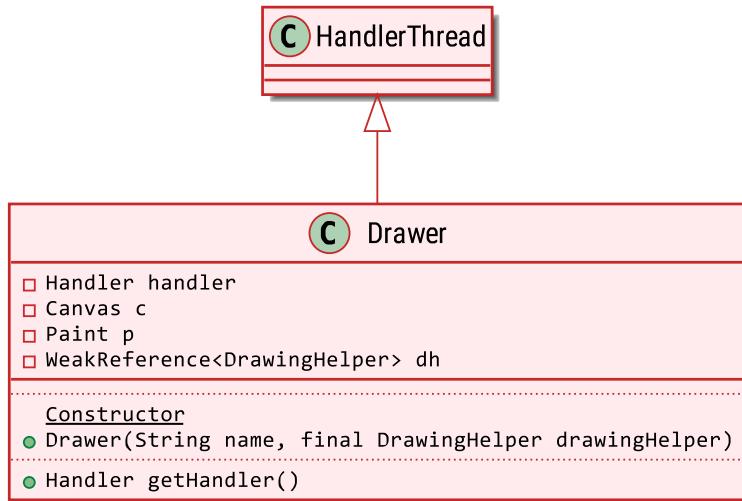


Figure 9.11: Class diagram of the Drawer class.

9.1.3 Connection

As we have often mentioned, the communication between the acquisition device and the app will rely on a bluetooth connection. There are mainly two different situations where we are using this type of connection:

- During the device connection menu, where we use bluetooth to scan all nearby devices in order to find ZEcg;
- During real-time acquisition, when we connect to the device and start receiving the patient ECG samples.

In the first case, we made use of a BroadcastReceiver, which will be notified during bluetooth scanning, when a new device will be found. As we don't want to include external devices in the device scan list, we filtered all devices by name, and included only devices with a name containing the word "ecg".

In the second case we made use of the class BluetoothSerial, which is an external library that we have modified and included in our application. The type of connection between two devices differs according to the specific bluetooth modules mounted on the different devices we tested our application with. For most of the devices we could establish an unsecure connection with Zecg, without any security protocols nor explicit pairing request. On other devices the connection was established only after an explicit pairing request. Either way the Zecg was always the server side and the smartphone

the client side of the Bluetooth connection. Whenever a client closes the connection with the server, the connection is automatically lost and any other device can ask data from the server.

BluetoothSerial

As we have said when we were talking about the ZEcgReceiver, the BluetoothSerial class will pass data received from the acquisition device to the last by sending messages to it. This is done by using an Handler instance that is passed in the BluetoothSerial constructor, that is pointing to the MessageQueue of the ZEcgReceiver. The messages will include the partial data received from the bluetooth connection, that will be parsed as we described in the ZEcgReceiver section. BluetoothSerial will create different thread, where each one will handle a different connection step:

- AcceptThread: this thread runs while listening for incoming connections. It behaves like a server-side client. It runs until a connection is accepted (not used in our case, as the smartphone always act like a client in the connection);
- ConnectThread: this thread runs while attempting to make an outgoing connection with a device. It runs straight through. The connection either succeeds or fails;
- ConnectedThread: this thread runs during a connection with a remote device. It handles all incoming and outgoing transmissions (only incoming in our case).

In the ConnectedThread the message sending will take place.

9.1.4 Operations Flow

In the previous sections we went deep to some component implementations in order to clarify some of the main functioning and iterations. Therefore, we move from a low level description, to an higher level, in order to explain the flow of the app operations. We will focus on our main two operation flow:

- Record acquisition: the user acquires a new ECG record by connecting to the acquisition device;
- Saved record opening: the user open a stored ECG record, choosing one of them from a list, and visualizes it.

For the sake of clarity, we put aside technical terminology for activity describing, as some methods call may be not so explanatory for the understanding. Instead, we used a general description of each activity, being all the details in the previous sections.

Record acquisition

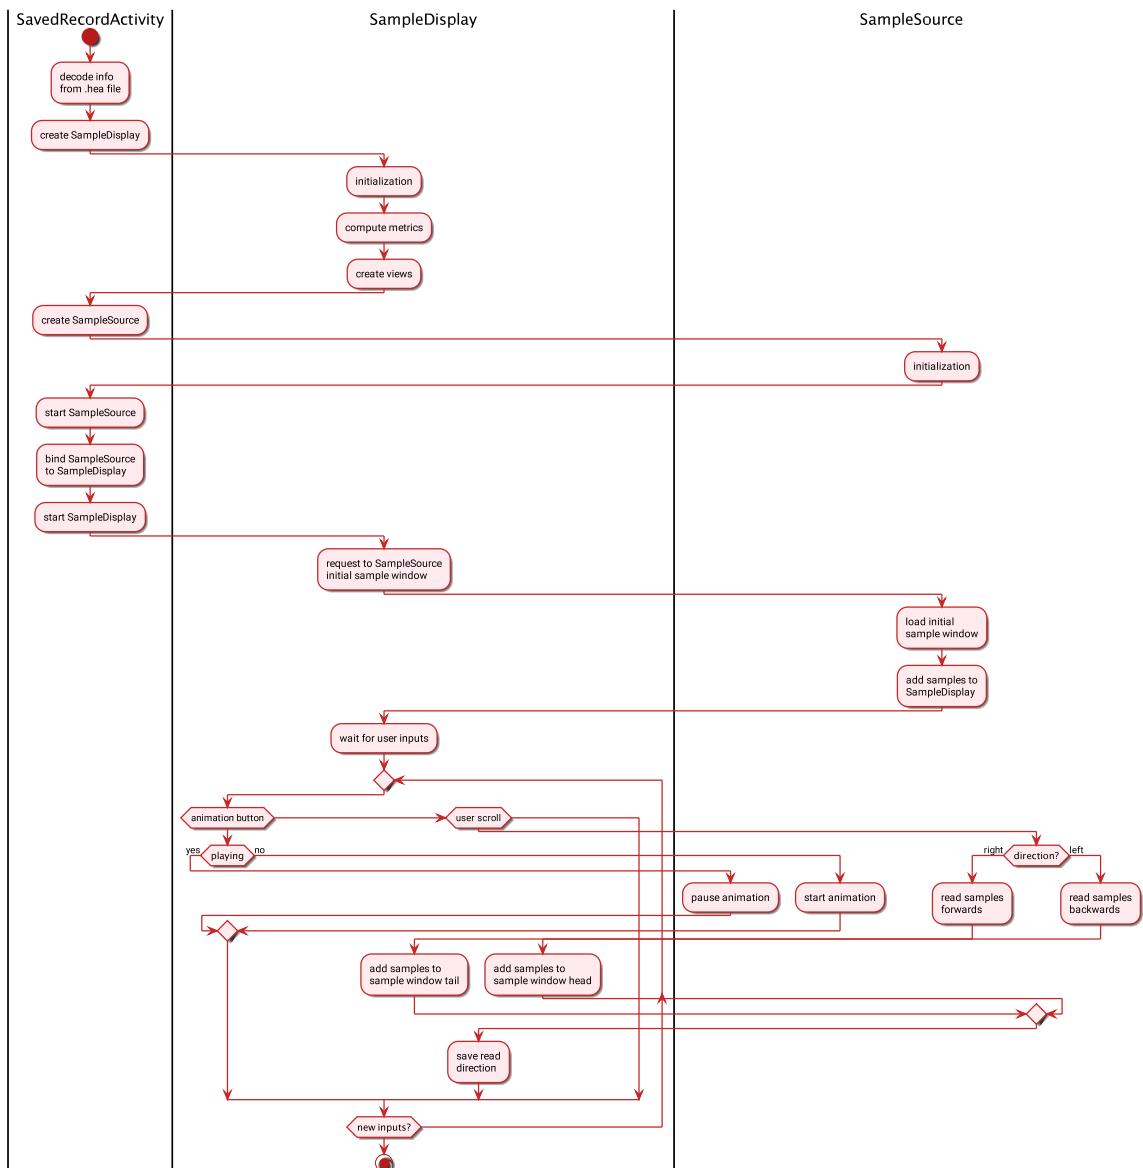


Figure 9.12: Activity diagram of a new ECG record acquisition operation.

The process starts with the selection of “*New Record*” in the app home. After that the user will move to the *CreateNewRecordActivity*, where will be displayed a

form where the user will be required to specify his:

- Name,
- Surname,
- Sex,
- Age,
- Weight.

If the user will omit some information, and clicks on the submit button, an error will be displayed, asking the user to complete the form. After a positive check of the form, there will be the creation of the .hea, including all the previous user data, and the .dat file, the container of all ECG record samples.

Then, the bluetooth connection of the device will be checked: if deactivated, the user will be jumped to the DeviceConnectionActivity, where will be able to activate it and search for the acquisition device; if activated, the BluetoothSerial will start the connection with the device, and wait for the first data. At each new data arrival, it will send a message with the new data to the ZEcgReceiver class. This last after having received the message, will extract the new data and insert it into a acquisition buffer. Then it will try to empty the buffer, removing all complete samples. All removed sample will be decoded and added to the sample window tail of the SampleDisplay. The process of acquiring new data and pass it to the ZEcgReceiver continues until the user will click on the stop button. Afterwards, the connection will be close and additional data will be written to the .hea record file. The operations flow are depicted as an activity diagram in the figure 9.12.

Saved record opening

The process starts with the selection of “*Open Record*” in the app home. A list of all stored ECG record will be presented to the user. The user will be able to scroll the list and select a ECG record to open it. After that it will be moved to the SavedRecordActivity, where the record .hea file will be decoded to extract all useful information, as sample rate, number of lead, lead labels, signal gain and others. Then, the SampleDisplay component will be created: this will compute all the metrics for which regarding the dynamic ECG paper sizing, and will create all ECG paper strip

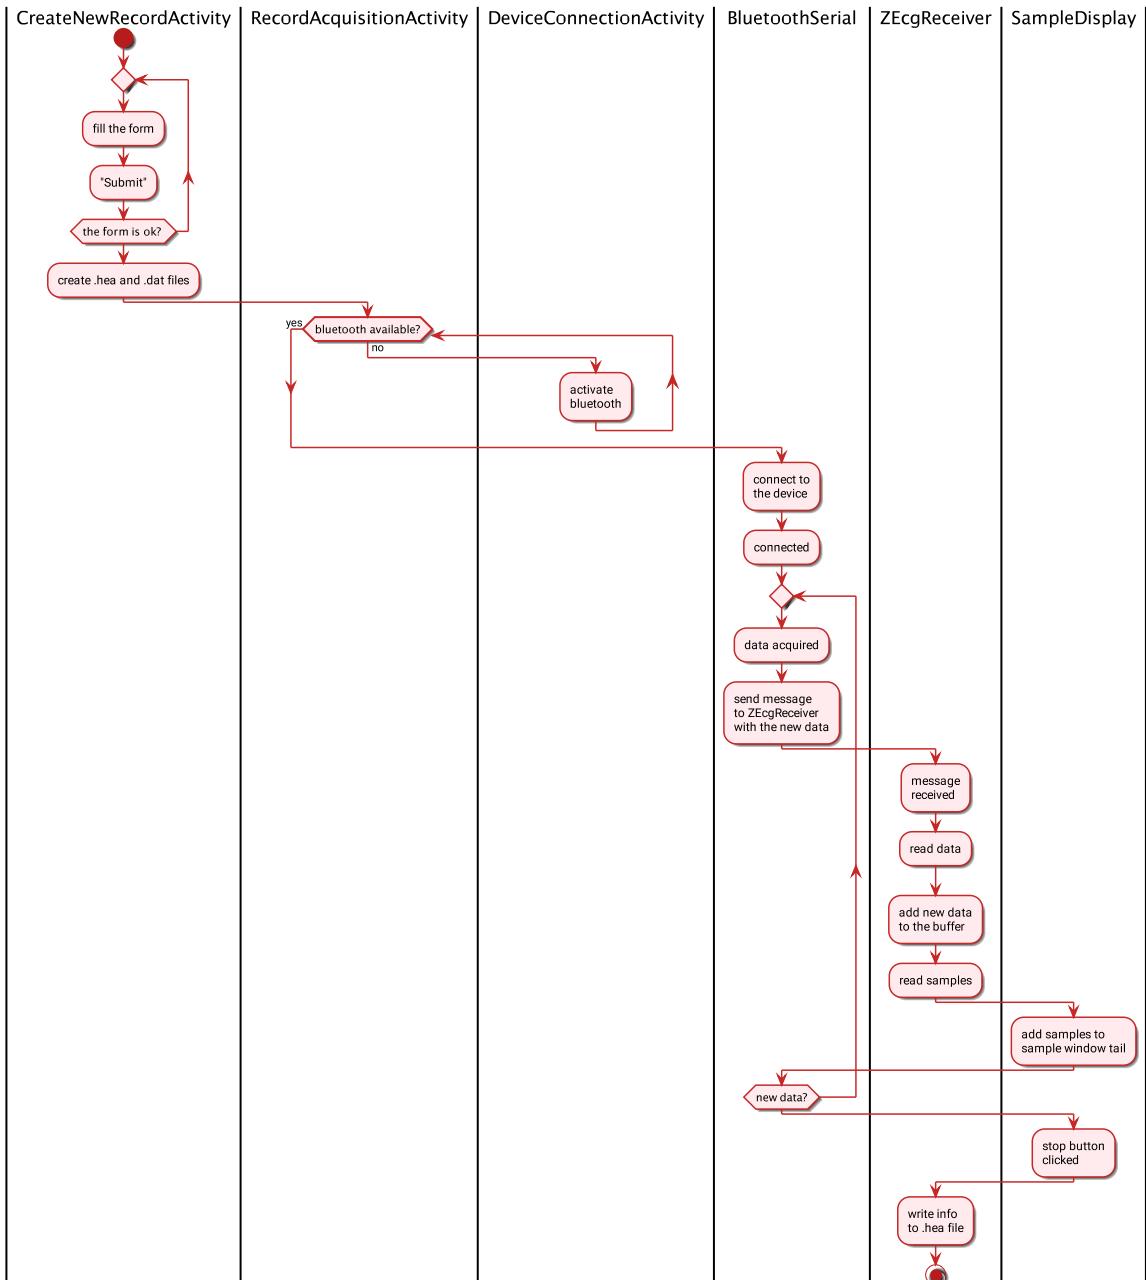


Figure 9.13: Activity diagram of a new ECG record acquisition operation.

views. Subsequently, the SampleSource will be created and initialized and binded to the SampleDisplay, which will be started. At start, the SampleDisplay will request from the SampleSource the initial sample window, in order to fill the ECG lead strips created previously. Later, the SampleDisplay will start waiting for any user input. The user input will be of two types:

- Animation button click;
- Scrolling.

In the first case, the state of the animation will be queried:

- If playing, the animation will be stopped;
- Otherwise, the animation will be started.

In the second case, according to the scrolling direction, the respective samples will be read and added to sample window tail or head, respectively if right or left.

The process of processing new user inputs continues until the user will close the record visualization. The operations flow are depicted as an activity diagram in the figure 9.13.

Chapter 10

Final result

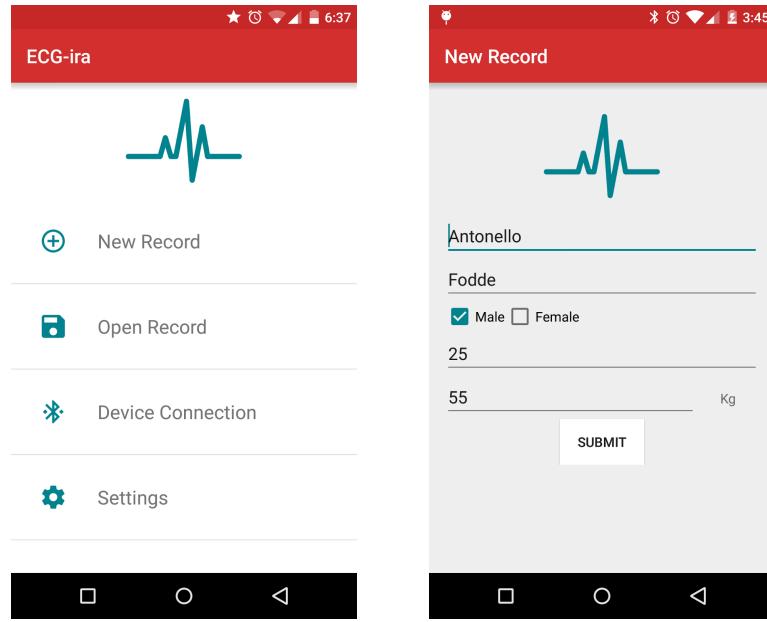
ECG-ira (ECG instant rapid analyzer) is a fully working mobile software application able to replace and fulfill most of the functionalities of an analog desktop application. The solution was designed according to the best mobile application development principles and it resulted working fine on all the smartphones¹ we tested it on. In this chapter we will show the final results providing significant screenshots of the application and metrics related to the performances evaluated and calculated with the application running on different devices.

10.1 App Screens

By launching the application the first screen is the home screen. It is the starting point, here the user can create a new record (starting a new acquisition with the ZEcg device), he can open the list of records (previously acquired or already present in the device system folder), he can discover and establish a connection with an acquisition device for future acquisitions and he can access to the application settings section where he can customize the application behaviour and tune some settings like the default folder to store in and open the records from.

By clicking on the “New Record” button option (see figure 10.1a), the user will be asked to fill a form with his data as a patient. This data will be then stored within a header file (with .hea extension) together with the file containing the effective record (.dat file extension). The personal data are only used by the doctor to identify the

¹Samsung Galaxy SIII Mini, Samsung Galaxy Note N7000, OnePlus X, Sony Xperia Z3 Compact, Sony Tablet Z, Huawei P8 Lite



(a) The home screen of the application ECG-ira.

(b) The form to be compiled before a record can start. It is necessary to bind any record to a patient data in order to avoid confusion between ecg records. This data is also useful for a fully understanding of the record itself

Figure 10.1: Two screens of the application, respectively the home screen and the form screen to be compiled by the patient

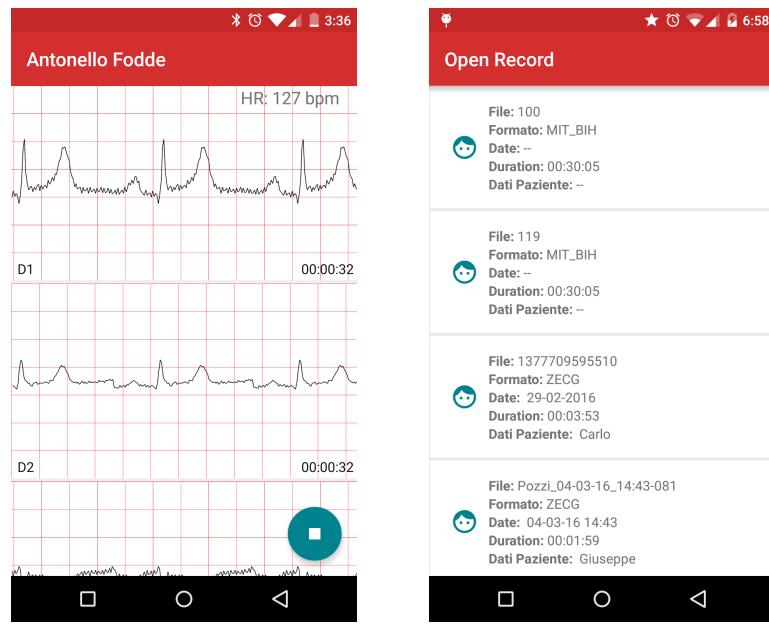
patient records.

Over the submission of the form the file .hea is immediately created to store all the information. If the device is already connected with an acquisition device ZEcg, the application tries immediately to establish a connection to receive data and visualize it (see the acquisition screen in the figure 10.2a). Otherwise it opens a connection request activity to enable the Bluetooth and start to discover nearby Bluetooth devices.

Going back to the home screen (figure 10.1a), by clicking on the “Open record” option button the user will access a list of ECG records saved (figure 10.2b) on its own device on a predefined root folder within the device system memory (the default folder can be changed within the setting section).

For each record the most relevant information are highlighted:

- File Name



(a) A real time ECG acquisition from the Zecg acquisition device.

(b) List of records within the device default folder

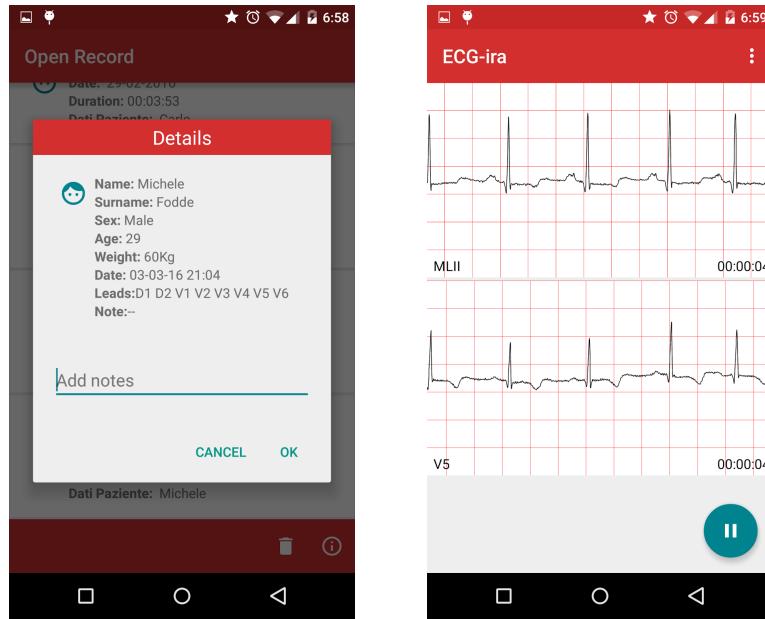
Figure 10.2: Two screens of the application representing the realtime acquisition and the list of records within the device

- File Format
- Acquisition Date
- Duration in time of the record
- Name of the patient

For more details about each record it is possible to long press on the desired record to show up a ToolBar shown in figure 10.3a.

The ToolBar is show at the bottom, it gives the possibility to show more details about the record and the patient information. In the figure 10.3a the user clicked on the “i” info button so the .hea file containing the user information are shown. Plus the user (in this case the doctor) can add personal notes to the specific record and the specific patient. The notes will be permanently stored within the .hea file.

The other option within the ToolBar is a delete record option in case the user decide to remove that specific record, assuming it was taken with too noises or it was just a test record.



(a) Screen triggered by long pressing on a list item. At the bottom of the screen there is a ToolBar showing the info and the delete options described above.

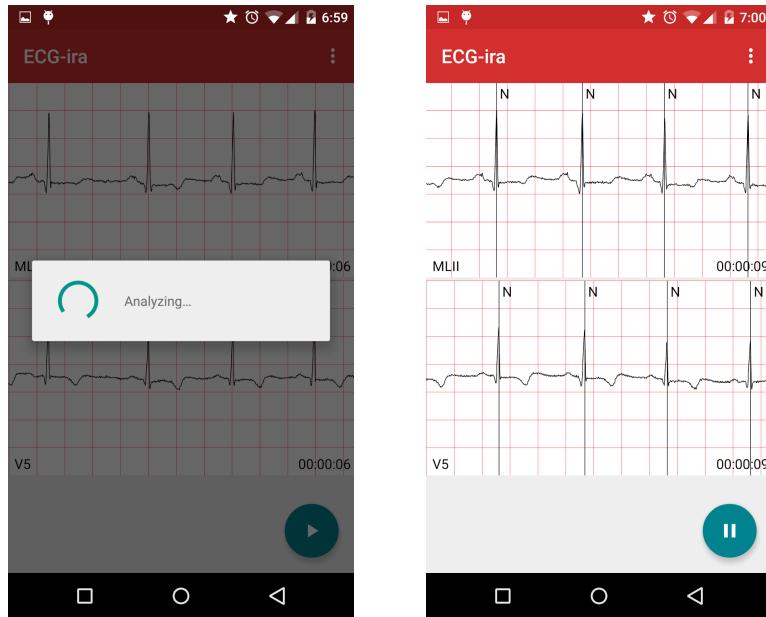
(b) An example of record opening and plotting. In the screen a MIT-BIH record format was opened.

Figure 10.3: Two screens of the application representing the info dialog for a certain record

By opening a record the application will fetch the .dat file and will plot its content over the ECG paper view (figure 10.3b). It is possible to scroll the record strip just by touching and dragging the screen with a finger, otherwise it is possible to reproduce the record as it was recorded with the exact acquisition speed time by clicking on the green round button at the bottom right.

The top ToolBar option menu is used to trigger the analysis algorithm over the opened record. If the user triggers this command, due to the great amount of operations behind the analysis process we force the user to wait for the result till the process ends.

The result of an analysis is the plot of the annotation letters (N=normal, V=ventricular, A=atrial) related to each detected beats and its diagnosis. In the figure 10.4b we can see four heart beats marked as Normal. The annotations covers all the record length and even if it is done on one lead only (D2 as default analysis signal), it is shown on



(a) Progress displaying during the analysis process over the record.

(b) Screen of the record plot with annotations results from an analysis over the first lead (MLII).

Figure 10.4: Two screens of the application representing the info dialog for a certain record and the visualization of that record

all the other leads.

Another result from the analysis is the generation of three different graphs (see figure 10.5), useful to support a better analysis over the heart overall behaviour as they sum up the total record “statistics”. They are:

- **Istogram**: on the X-axis we have an ordered values of the heart beats measures in bpm. On the Y-axis the occurrences of beats at that bpm value.
- **Tacogram**: on the X-axis we have the number of heartbeats retrieved by the analysis. On the Y-axis the bpm values of each heart beats.
- **ST+/ST-**: this graph shows instead, for each heartbeat, the amount of area below the ST segment and the area above that segment. This is useful especially to detect some specific arrhythmias.

About the customization setting for the application, we made a long setting screen with lots of customizable parameters (figure 10.6).

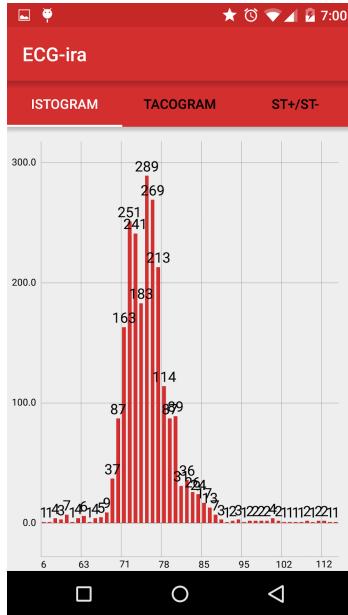


Figure 10.5: Screen of the three graphs generated after an analysis of the records. They are graphs summing up the complete record characteristics. Useful for a general overview of the patient record and health status.

The first is related to the application of a baseline filter during the record visualization in order to reduce the baseline wandering effect. The second category of settings are related to the acquisition phase. It includes the acquisition format, the gain, the sample rate, quantization and the name of the acquisition device (for example it can be: ZEcg vers 2/3/4).

The third category is related to the drawing settings. These settings will affect only the graphical views such as the height of each strips, the way the strip is reproduces (by scrolling or by using the old style oscilloscope effect) and the preferred frame rate to be used.

The last category is to tune the default folder to store the new acquired records and from which open the records within the application (figure 10.2b).

10.1.1 Performance

In this section we reported some performance results from different aspects such as Memory usage, CPU usage and response time during a run of the application performed on different devices with installed different Android OS and different hardwares. The tests were performed on the following devices dated respectibely on

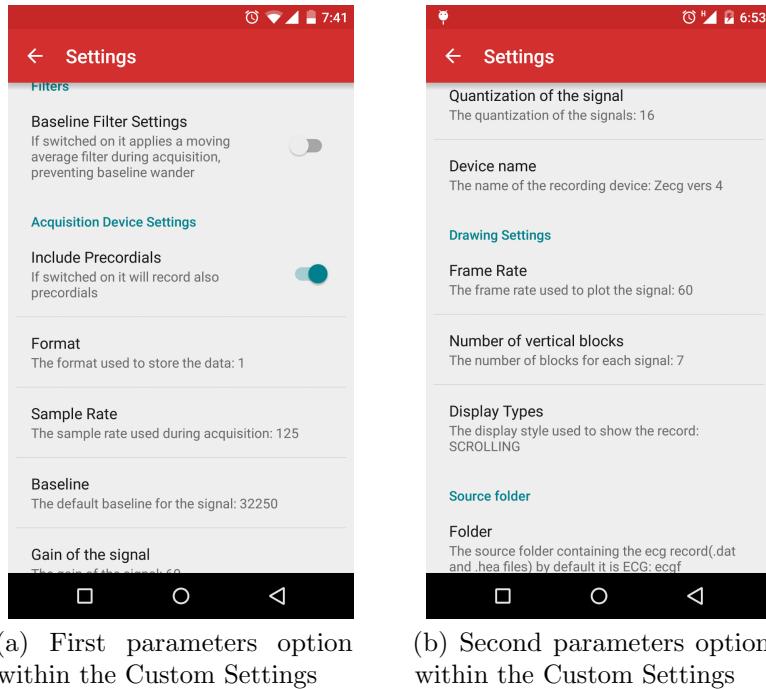


Figure 10.6: The customization section of the application. A list of available settings parameters to tune the application behaviour.

2011 and on 2015:

- Samsung Galaxy Note N7000
- OnePlus X

The table 10.1 shows the specification details of these 2 smartphones.

The results are not to be considered absolutely valid for all the above device families. The performance can be affected by many different factors that may not be related with the applications itself; hence the results may vary from a test to another. Since Android OS is in charge to manage the system memory and split the available resources between all the installed applications, having our application on foreground doesn't necessarily mean we can reserve all the power and memory resource for our business. Part of the system resources are allocated to run the Garbage Collector which itself has its own timing when cleaning and deallocating other application resources. Then all the applications having background tasks also need to be kept alive, for example notification services or messaging services.

It is not easy to measure the application performance in a clean absolute way, harder

Model	OnePlus X	Samsung Galaxy N7000
Brand	OnePlus	Samsung
Year	November 2015	October 2011
Display Type	AMOLED capacitive touchscreen, 16M colors	Super AMOLED capacitive touchscreen, 16M colors
Display Size	5.0 inches	5.3 inches
Resolution	1080 x 1920 pixels (441 ppi pixel density)	800 x 1280 pixels (285 ppi pixel density)
OS	Android OS, v5.1.1 (Lollipop)	Android OS, v4.1.2 (Jelly Bean) (upgraded from v2.3.5)
CPU	Quad-core 2.3 GHz	Dual-core 1.4 GHz Cortex-A9
Memory(internal)	16 GB	16 GB
RAM	3 GB	1 GB

Table 10.1: The two test devices specification details.

is to compare the results coming from different devices, that is why we will only report the results related to each devices and we will analyze how the application performs on each.

10.1.2 Evaluation

We evaluated the performance considering three application states:

- **Idle state:** the application is open on a simple basic screen such as the home screen or the record list screen.
- **Plotting state:** the application is performing a plotting of a record
- **Analysis state:** the analysis algorithm is performed over the record

Memory

Before going in details about the memory usage result for the application performing on the different devices, there is some basics knowledge about the analysis tools used and the way to read them that the reader should be aware of.

Android Studio provides a Memory Monitor Tool allowing the user to track how his

application performs in term of memory usage. Apart of a generic view of memory usage there is a more detailed view that let you observe how your app's memory is divided between different types of RAM allocation. The most important and relevant values to care of are:

- Private (Clean and Dirty) RAM: The memory being used by ONLY your process. This is the bulk of the RAM that the system can reclaim when your app's process is destroyed. The most important is the private dirty RAM, which is the most expensive because it is used by only your process and its content exists only in RAM so it can't be paged to storage. (Android does not use SWAP). All Dalvik and native heap allocations you make will be private dirty RAM; Dalvik and native allocations you share with the Zygote process are shared dirty RAM.
- Proportional Set Size (PSS): This is a measurement of your app's RAM use that takes into account sharing pages across processes. Any RAM pages that are unique to your process directly contribute to its PSS value, while pages that are shared with other processes contribute to the PSS value only in proportion to the amount of sharing. For example, a page that is shared between two processes will contribute half of its size to the PSS of each process.
- Dalvik Heap: The RAM used by Dalvik allocations in your application.
- .so mmap and .dex mmap: The RAM used for mapped .so (native) and .dex (Dalvik or ART) code.
- .art mmap: Amount of RAM used by the heap image which is based off the preloaded classes which are commonly used by multiple apps. It does not count towards your app heap size.
- .Heap: The amount of heap memory for your application.

The memory tests were executed running the application on the same screens and opening and plotting always the same record. It was used the 100.dat from MIT/BIH format because it is the longest and biggest record.

Samsung Galaxy N7000 This is the oldest device according to its release date. To make this device compatible with the minimum API requirement (API 16) the device was upgraded to android Jelly Bean (API 16).

```

Applications Memory Usage (kB):
Uptime: 135182 Realtime: 135180

** MEMINFO in pid 4578 [it.polimi.anto.chai.ecg_ir] **
              Shared   Private   Heap   Heap   Heap
              Pss     Dirty     Dirty    Size   Alloc   Free
-----+-----+-----+-----+-----+-----+-----+
  Native      40       32       40    3480    3451     24
  Dalvik    3409    15896    3108   13575   13008    567
  Cursor        0       0       0
  Ashmem        0       0       0
  Other dev      4      28       0
  .so mmap    1009    2388      592
  .jar mmap      0       0       0
  .apk mmap     137       0       0
  .ttf mmap       2       0       0
  .dex mmap    1948       0      12
  Other mmap    454      320      36
  Unknown    1373     692    1364
  TOTAL      8376    19356    5152   17055   16459    591

  Objects
  Views:          36      ViewRootImpl:         1
  AppContexts:      2      Activities:         1
  Assets:           3      AssetManagers:       3
  Local Binders:      6      Proxy Binders:     15
  Death Recipients:  0
  OpenSSL Sockets:  0

  SQL
  MEMORY_USED:      0
  PAGECACHE_OVERFLOW:  0      MALLOC_SIZE:        0

```

Figure 10.7: An example of detailed view of an application memory usage.

The hardware instead was the same from the manufacturer since its release on 2011. During the test the device performed quite well. It never crashed due to memory leakage or Out of Memory Exception.

Let's start observing the memory usage during the three application states:

Idle state The memory usage looked constant over the time as expected. The memory usage graph above shows the actual memory stack allocated for the application and how many of that is used by the application. As it can be seen on the top right the application is using 13.18 MB and 20.46 are still free. The total device memory is not shown. In case the application hits the upper bound, let's say it occupies the other free 20.46 MB, then the OS will allocate more memory for the actual application by garbage collecting resources from other background applications or actually killing other low priority processes (all the processes not in foreground are considered with lower priority).

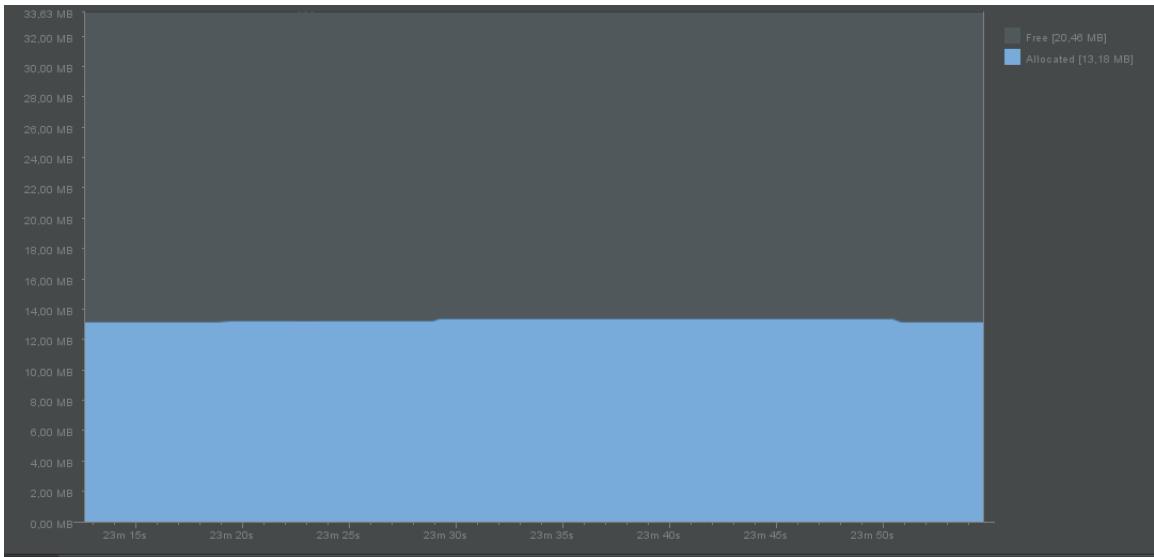


Figure 10.8: Memory usage during application idle state. It is quite constant (due to user interaction it can slightly increase or decrease if no interaction at all for a while).

Here are more detailed information about the memory usage and resources allocations during idle state:

Plotting state During a record plotting the memory usage increases since the drawing process itself consume memory by allocating new resources and by the resource manipulations and usage.

The waves in the memory usage graph (figure 10.10a are due to the garbage collector calls which free unused resources. This is really an interesting Garbage Collecting behaviour, because we can observe it collecting resources with a certain regularity rate (somehow very often).

The detailed information about memory usage divided by types of RAM used are shown in the figure 10.10b.

The Private Dirty Dalvik in increased from Idle state from 3108 kB to 5224 kB. The TOTAL Pss increased from 8376 kB in Idle state to 12343 kB during a plot.

```

Applications Memory Usage (kB):
Uptime: 135182 Realtime: 135180

** MEMINFO in pid 4578 [it.polimi.anto.chai.ecg_ira] **
              Shared   Private   Heap   Heap   Heap
              Pss     Dirty    Dirty   Size  Alloc   Free
-----+-----+-----+-----+-----+-----+-----+
  Native      40       32      40   3480   3451     24
  Dalvik    3409    15896    3108  13575  13008    567
  Cursor        0       0       0
  Ashmem        0       0       0
  Other dev      4      28      0
  .so mmap    1009    2388     592
  .jar mmap      0       0       0
  .apk mmap     137      0       0
  .ttf mmap       2      0       0
  .dex mmap    1948      0      12
  Other mmap    454     320      36
  Unknown    1373     692    1364
  TOTAL      8376    19356    5152   17055   16459    591

  Objects
  Views:          36      ViewRootImpl:         1
  AppContexts:      2      Activities:          1
  Assets:           3      AssetManagers:        3
  Local Binders:     6      Proxy Binders:       15
  Death Recipients:  0
  OpenSSL Sockets:  0

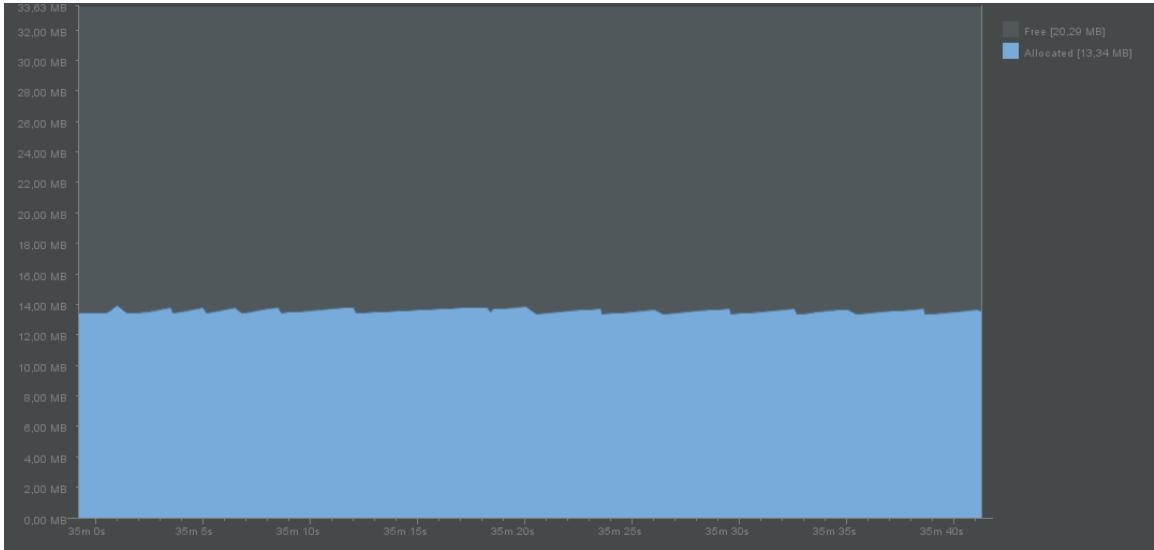
  SQL
  MEMORY_USED:      0
  PAGECACHE_OVERFLOW:  0      MALLOC_SIZE:        0

```

Figure 10.9: Details about memory usage and allocation over the different RAM section for the application ECира.

Analysis state This is the most interesting part of the memory usage results. We can clearly distinguish the application behaviour just by looking at the memory stack. The analysis is the more complex and intensive memory usage step. At seconds 0m 28s the user clicked on the analysis button. The peak on the memory usage graph is due to the creation of the buffer to store the entire record in order to perform the analysis algorithm on it. At 10m34s the analysis algorithms finish to perform. More resources are instantiated for computational purpose of the algorithms. As at 10m34s the algorithms terminate, all the resources can be freed and only results are returned.

As you can see in the memory usage graph, the OS kind of predicted or expected the application to require and allocate more resources, so it allocates more spaces for it even if the upper bound of 28.40 MB wasn't reached. But as soon as the resources allocated were of no use anymore, it triggered the Garbage Collector to free them. In figure 10.13 is the detailed information about the overall memory usage which is divided by RAM types. We can easily observe that the Private Dirty Dalvik became the triple with respect to its size during the plotting phase (5224 kB -> 15664 kB)



(a) Memory usage during a record plotting. The waves are due to Garbage collector calls to free resources. The final resource usage in this phase is rather constant as well.

```

Applications Memory Usage (kB):
Uptime: 282419 Realtime: 282416

** MEMINFO in pid 4578 [it.polimi.anto.chai.ecg_ira] **
              Shared   Private   Heap   Heap   Heap
              Pss     Dirty     Dirty    Size   Alloc   Free
-----+-----+-----+-----+-----+-----+
      Native      44       28       44   4444   4382      61
      Dalvik    5518   15576     5224  14599  13584   1015
      Cursor        0       0       0
      Ashmem        0       0       0
      Other dev      4       28       0
      .so mmap    1259   2488     784
      .jar mmap       0       0       0
      .apk mmap     149       0       0
      .ttf mmap      75       0       0
      .dex mmap   2448       0      12
      Other mmap     569     320      36
      Unknown    2277     516    2272
      TOTAL     12343   18956    8372   19043  17966   1076

      Objects
      Views:        205     ViewRootImpl:         3
      AppContexts:      4     Activities:         3
      Assets:         3     AssetManagers:       3
      Local Binders:    12     Proxy Binders:      19
      Death Recipients:  0
      OpenSSL Sockets:  0

      SQL
      MEMORY_USED:      0
      PAGECACHE_OVERFLOW:  0     MALLOC_SIZE:        0
  
```

(b) Memory usage by RAM types during a record plot.

Figure 10.10: Two memory usage views, relatively the realtime view with the application running and the per process memory view and allocation.

and the TOTAL Pss doubled. This behaviour is due to the record copy into memory and all the data structures used to compute the analysis.

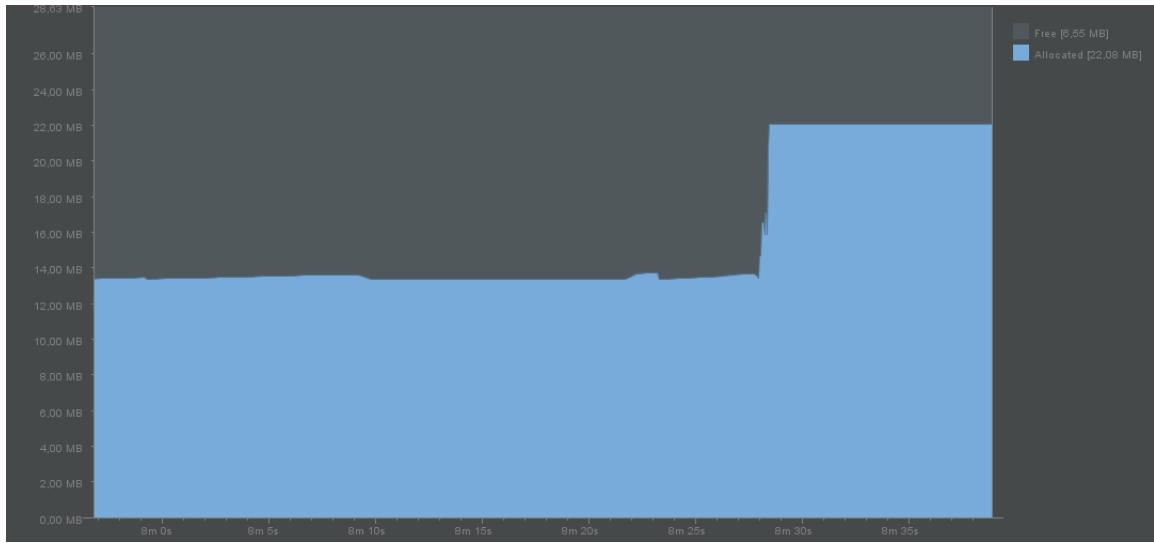


Figure 10.11: Memory usage during an analysis launched over a record. The step is due to the record allocation in a memory buffer.

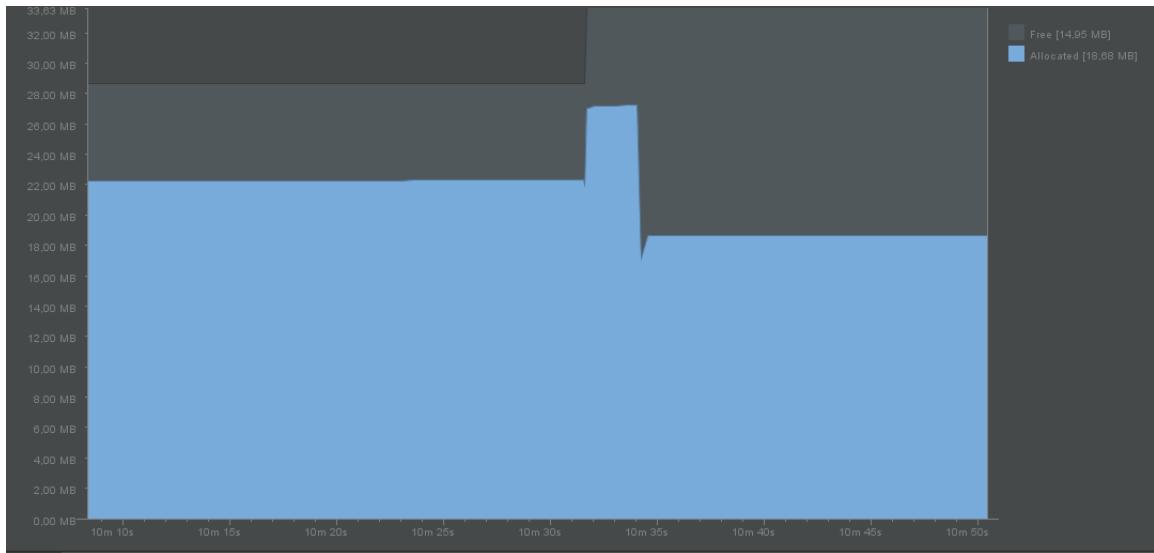


Figure 10.12: Memory usage at the exact time the algorithm performs the final steps by computing the ST+/St- areas (execution finish at 10m34s).

Much of the other field didn't changed their values nor allocation size, for example the .so mmap was the same both in the plotting state both in the analysis state.

As final result we can conclude that during the entire process and states of the application ECG-ira on the Samsung Galaxy N7000, the application performed fairly well consuming resources as it needed and releasing them as soon as possible. We

```

Applications Memory Usage (kB):
Uptime: 818527 Realtime: 818524

** MEMINFO in pid 4578 [it.polimi.anto.chai.ecg_ira] **
              Shared   Private    Heap     Heap     Heap
              Pss      Dirty     Dirty    Size   Alloc   Free
-----  -----  -----  -----  -----  -----  -----
Native      48       24       48    5604    5108    299
Dalvik    15941    15160    15664  34439   22620   11819
Cursor        0        0        0
Ashmem        0        0        0
Other dev     4       28       0
  .so mmap   1218    2476     796
  .jar mmap    0        0        0
  .apk mmap   231      0        0
  .ttf mmap    76      0        0
  .dex mmap   2580      0       20
Other mmap   880     320     256
Unknown    3255    428     3252
TOTAL      24233   18436    20036   40043   27728   12118

Objects
      Views:      230      ViewRootImpl:        4
  AppContexts:        4      Activities:        3
      Assets:        3      AssetManagers:      3
Local Binders:      23      Proxy Binders:     20
Death Recipients:    0
OpenSSL Sockets:    0

SQL
      MEMORY_USED:        0      PAGECACHE_OVERFLOW:        0      MALLOC_SIZE:        0

```

Figure 10.13: Memory usage details by RAM types.

observed also that the OS Garbage Collector overall behaviour was to trigger as often as possible as soon as the resources in no more used.

OnePlus X This device is the most advanced one we tested the application on. With it's 3GB RAM memory and its powerful quad-core processor we really expect no issues, but let's dig into the results.

For a better understand of some terms, you are invited to read the introduction of the Memory Chapter.

Idle state The idle state is as well predictable, the memory usage is stable. It is relevant to observe that the device consume much more memory (20.87 MB) with respect to the idle state values measured on the Samsung Galaxy N7000 (13.18 MB). This can be explained with the fact that firstly the two OS are firstly completely different and rely on different API. The OnePlus X device relies on the API 21 while the Samsung N7000 on API 16. Also the screen size and pixels density is not comparable since the OnePlus X doubles the old Samsung device.

Apart using more resources the final result is the same since the idle state doesn't

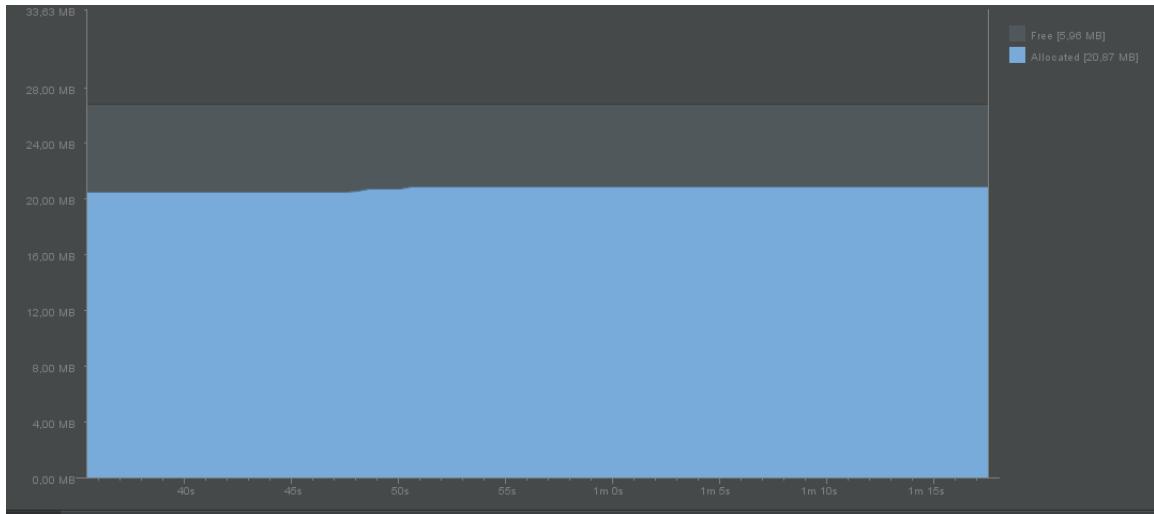


Figure 10.14: Memory usage of ECG-ira in Idle state.

show any memory leakage. More information are provided in the figure 10.15. An

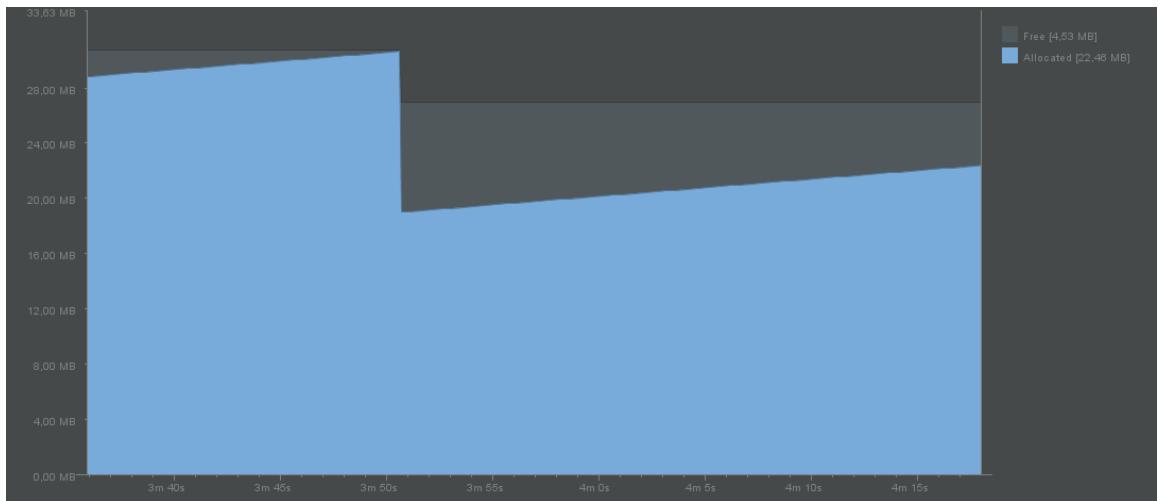
Applications Memory Usage (kB):							
Uptime: 9629164 Realtime: 12331506							
** MEMINFO in pid 7388 [it.polimi.anto.chai.ecg_ira] **							
	Pss	Private	Private	Swapped	Heap	Heap	Heap
	Total	Dirty	Clean	Dirty	Size	Alloc	Free
Native Heap	4281	4240	0	0	12288	6599	5688
Dalvik Heap	5552	5164	0	0	27470	21411	6059
Dalvik Other	576	576	0	0			
Stack	112	112	0	0			
Gfx dev	3182	2772	0	0			
Other dev	5	0	4	0			
.so mmap	857	336	120	0			
.apk mmap	136	0	24	0			
.ttf mmap	49	0	4	0			
.dex mmap	3156	0	3152	0			
.oat mmap	681	0	72	0			
.art mmap	802	584	20	0			
Other mmap	4	4	0	0			
EGL mtrack	39808	39808	0	0			
Unknown	156	156	0	0			
TOTAL	59357	53752	3396	0	39758	28010	11747
objects							
	Views:	158	ViewRootImpl:	2			
	AppContexts:	4	Activities:	2			
	Assets:	3	AssetManagers:	3			
	Local Binders:	10	Proxy Binders:	19			
	Parcel memory:	3	Parcel count:	14			
	Death Recipients:	0	OpenSSL Sockets:	0			
SQL							
	MEMORY_USED:	0	MALLOC_SIZE:	0			
	PAGECACHE_OVERFLOW:	0					

Figure 10.15: Memory usage of ECG-ira in Idle state by RAM types.

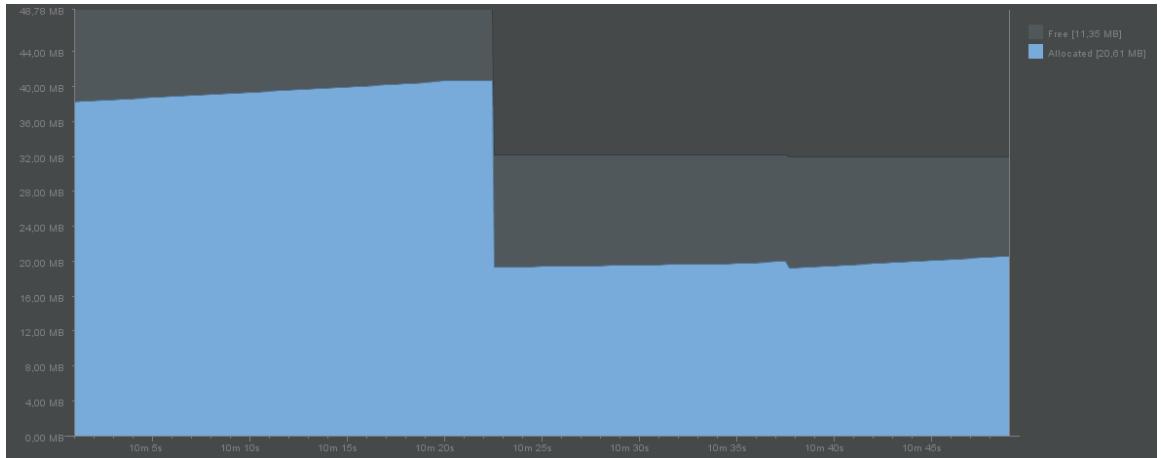
interesting data from the above table is that the Private Dirty Native Heap and the Private Dirty Dalvik Heap are really close to have the same value. This means that some views used inside the application makes use of the Graphical layer offered by

OpenGL (it means they are hardware accelerated). Actually most of the Total Pss and Private Dirty allocations comes from the EGL mtrack, in other words it refers to the memory consumed by Graphics layer.

Plotting state During the plotting phase we can observe a different Garbage Collector behaviour with respect to the one observed on the Samsung N7000. In the



(a) Memory usage during a record plotting. We can see that the step is due to a GC call.



(b) Memory usage during a record plotting. We force two GC calls as they can be seen considering the the first big step and the second small one.

Figure 10.16: Two memory usage views during a normal app execution, on the second figure a GC is called before the memory reach the heap

figure 10.16b we can see a system GC call triggered when the application reaches the upper bound of the memory heap allocated for the application process. The OS waits

to trigger a GC call until the process fills and hit the memory heap. Lollipop GC policy is then completely different from the Jelly Bean (Samsung N7000). It allows application to fill the heap and only after that it triggers a GC and estimates the new heap size. For a test purpose We choose to trigger a GC call before the process can fill the heap. You can see in figure 10.16b that the forced GC call cleared the heap (the first big step) and resized the heap according to the resource need. After that the heap keeps being filled with new allocations and also old allocations, but at 10m37s we triggered another forced GC call (second small step). We can now realize for real how many resources are really used during plotting phase and how many are just put in the GC queue. If we keep calling forced GC, probably we would achieve the same behaviour as the one observed on the Samsung N7000 (a series of memory waves).

A detailed view of memory allocation by RAM types is shown below: As the

```

Applications Memory Usage (kB):
Uptime: 10316126 Realtime: 13018467

** MEMINFO in pid 7388 [it.polimi.anto.chai.ecg_ir] **
              Pss   Private  Private  Swapped   Heap   Heap   Heap
              Total    Dirty    Clean    Dirty    Size   Alloc   Free
----- ----- ----- ----- ----- -----
Native Heap      6158     6120      0      0  16384    7761    8622
Dalvik Heap     6899     6512      0      0  33527   21963   11564
Dalvik Other     868      868      0      0      0      0      0
Stack           184      184      0      0      0      0      0
Gfx dev          4162     3752      0      0      0      0      0
Other dev          5       0       4      0      0      0      0
.so mmap         887      332     124      0      0      0      0
.apk mmap        150       0      28      0      0      0      0
.ttf mmap         55       0       4      0      0      0      0
.dex mmap        3692      0     3688      0      0      0      0
.oat mmap        1162      0     260      0      0      0      0
.art mmap        1582     836     488      0      0      0      0
Other mmap         24       4       0      0      0      0      0
EGL mtrack      39808    39808      0      0      0      0      0
Unknown          156      156      0      0      0      0      0
TOTAL          65792    58572    4596      0  49911   29724   20186

Objects
      Views:        221  ViewRootImpl:        3
  AppContexts:        7  Activities:        5
      Assets:        3  AssetManagers:        3
Local Binders:       26  Proxy Binders:       22
  Parcel memory:       4  Parcel count:       16
Death Recipients:       0  OpenSSL Sockets:       0

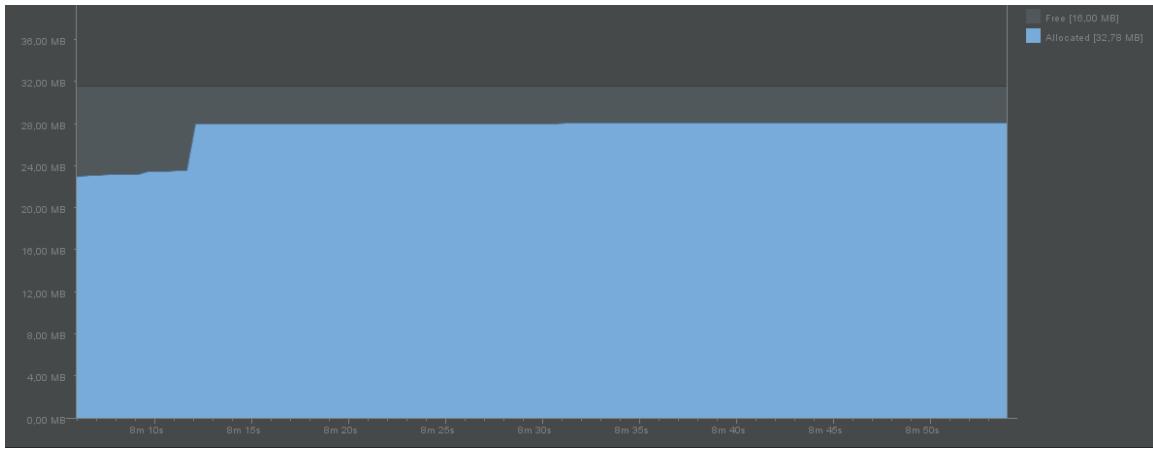
SQL
      MEMORY_USED:        0
PAGECACHE_OVERFLOW:        0  MALLOC_SIZE:        0

```

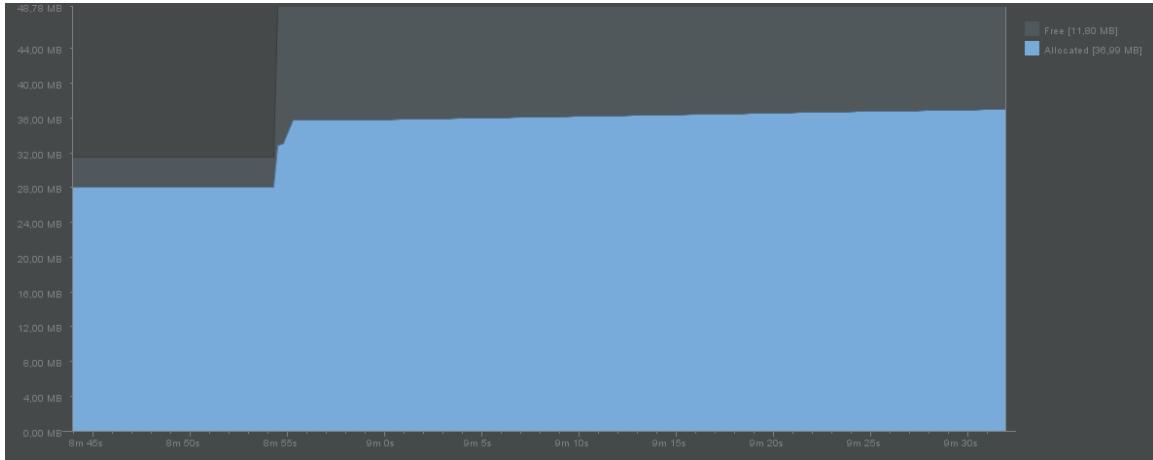
Figure 10.17: Memory usage details by RAM types.

plotting state requires more drawings and dynamic data usage (updates of data structure and views updates) there is an obvious increase of resources needs.

Analysis state Considering how the GC behaves the analysis state is kind of predictable. Instead of calling a GC after the buffer becomes useless as the analysis algorithm finishes, the OS instantiates more memory heap for the process as it needs and sets up a new and higher upper bound. The GC call will be triggered only when the new bound is reached. We already know that most of the allocated resources are to be garbage collected but since the limit of the heap size will not be hit no GC calls will be triggered. We can observe that all the resources used during the analysis are kept in memory.



(a) Memory usage when a record analysis is called. The step in the memory usage graph is due to the buffer allocation for the record.



(b) Memory usage when a record analysis is called. The step in the memory usage graph is due to the analysis algorithm running on the record.

Figure 10.18: Two memory usage views during the execution of the analysis algorithms, on the first figure 10.18a the step sign the execution satrting point, in figure 10.18b the step is due to the end of the execution

Due to the creation of the copy record buffer the Dalvik Private Dirty is doubled (from 6512 kB in plotting state to 14268 kB in analysis state) and since we pause all the views the EGL mTrack is reduced (less views need memory).

The overall memory usage keeps being the same also for the OnePlus X device, even if on this device the GC behaviour is completely different, the resource used are limited and deallocated when there is no more need of them. Having less GC calls improved the application performance since any GC call affects and keeps the CPU busy for while. Also the use of an additional graphical layer improves the application

```

Applications Memory Usage (kB):
Uptime: 10242259 Realtime: 12944600

** MEMINFO in pid 7388 [it.polimi.anto.chai.ecg_ir] **
              Pss  Private_Dirty  Private_Clean  Swapped  Heap_Size  Heap_Alloc  Heap_Free
-----  -----  -----  -----  -----  -----  -----  -----
Native Heap      6910        6872          0          0    16384     8708      7675
Dalvik Heap    14655       14268          0          0    32596    28939      3657
Dalvik Other      816         816          0          0
Stack            184         184          0          0
Gfx dev          3990        3580          0          0
Other dev           5          0          4          0
.so mmap         895         340         124          0
.apk mmap        150          0          28          0
.ttf mmap          55          0          4          0
.dex mmap       3692          0        3688          0
.oat mmap       1159          0        260          0
.art mmap       1582        836        488          0
Other mmap          24          4          0          0
EGL mtrack      30688       30688          0          0
Unknown          156         156          0          0
TOTAL          64961       57744        4596          0    48980    37647     11332

Objects
      Views:        232   ViewRootImpl:          4
  AppContexts:          6   Activities:          4
      Assets:          3 AssetManagers:          3
Local Binders:        23   Proxy Binders:        21
  Parcel memory:        4   Parcel count:        16
Death Recipients:        0 OpenSSL Sockets:          0

SQL
      MEMORY_USED:          0
PAGECACHE_OVERFLOW:        0   MALLOC_SIZE:          0

```

Figure 10.19: Memory usage when a record analysis is called. The memory usage is divided by RAM types.

responsiveness and performance, but it also comes with a more intensive memory usage that not all devices can afford. In case of the OnePlus X hardware memory limitation is not a great deal with its 3GB RAM memory, nor the performance limitation is really a limit, but since the application has to be run on other devices as well, these principles should always be considered.

CPU

The CPU performance tests were done by executing the same actions on the same views. We decide to provide CPU usage during an application run on different devices. As in the memory tests we decide to track CPU usage with the application on the three different states:

Idle state, Plotting state and Analysis state. The CPU usage is reported into percentages and is related to the amount of CPU power used by the Kernel and by the User (in this case the User is the application process).

Samsung Galaxy N7000 This device has a Dual-core 1.4 GHz Cortex-A9, performed fairly well even with its large screen 5.3 inches and its 800x1280 pixels (285 ppi pixel density).

Idle state During Idle state, when the application is simply kept on foreground on a static screen, we can observe there are no cpu usage leakage. The peaks in the

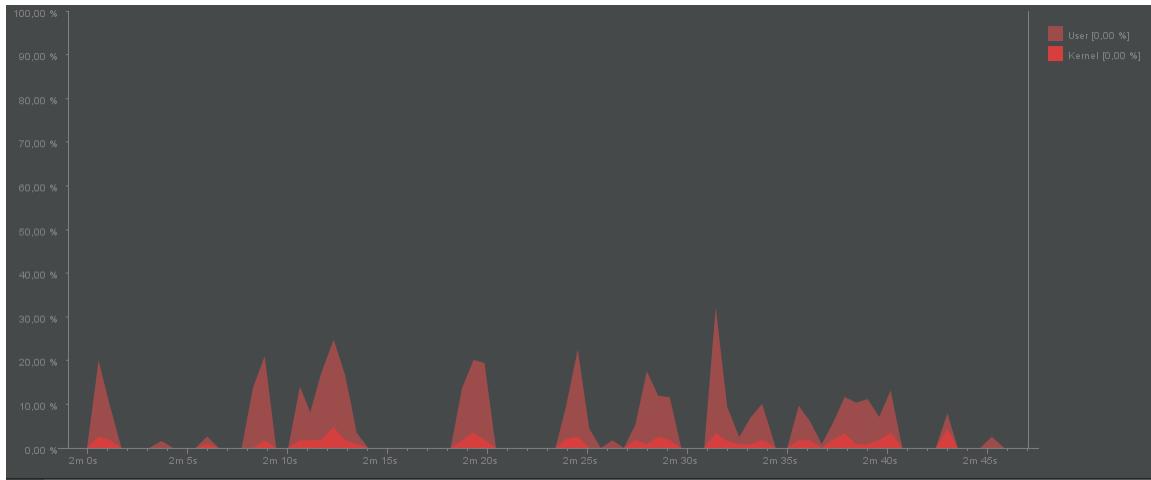


Figure 10.20: Cpu usage during idle state, the application is on foreground on a static page.

figure 10.28 are due to the user's interaction with the listView, otherwise in case of no interaction at all the graph is plain. At each peak the ListAdapter (in this case the RecycleViewAdapter) instantiates and so inflates the next item list represented by a record in the list so that the user can visualize it on the scrolling action.

Plotting state During plotting state the CPU usage results to be quite constant. In the figure 10.21, we can observe that it never exceed 60%. At minute 7m10s we selected the record 100.dat and opened it. The plotting screen keeps the cpu quite busy at a constant rate (between 40% and 58%) . This value doesn't change and it is independent from the user's interaction (not like the listView which requires cpu at each user's interaction as scrolling). Also during the automatic scrolling of the ECG strip, the cpu usage keeps being the same. This can be explained by observing that this screen and all the views were completely customized and implemented by us. As the views on screen are plotting canvases and to get the animation features we are forced to use a second thread to manage the drawings. By this way the screen content

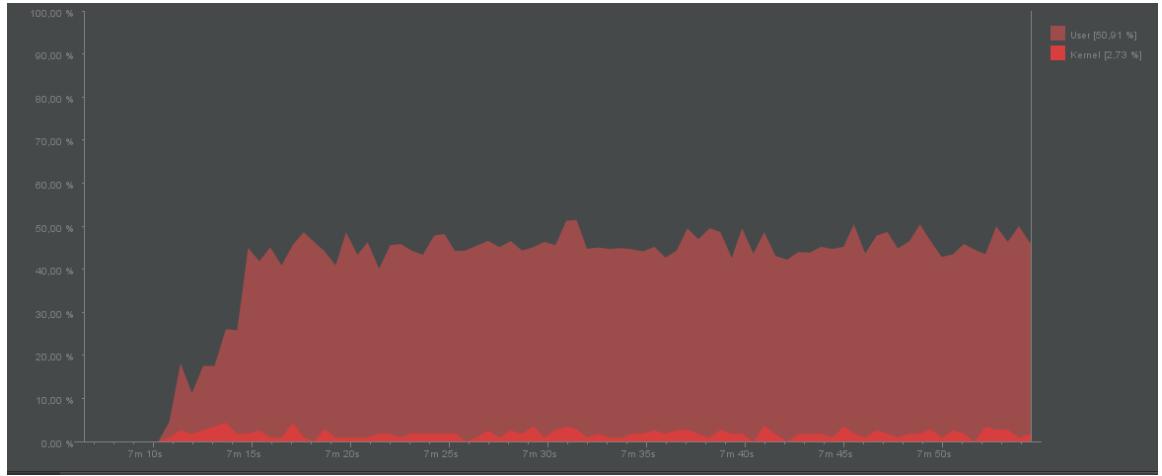


Figure 10.21: Cpu usage when a record is plot on screen.

is redrawn as often as it is possible according to the maximum frame rate set by the user but compatible the frame rate the thread can get.

As soon as the user closes the Plotting screen, the cpu usage dropped and resources are kept free. In figure 10.30 at minute 13m25s we closed the Activity with the ECG



Figure 10.22: Cpu usage when we close a plotting screen and reopen it.

strip plot. We can observe that the application released immediately the cpu resources as they are no more needed. Also all the drawing threads are killed and streams closed. But as at minute 13m35s we reopened another record, the resources are reallocated and the new views show ups. New drawing threads are instantiated and a new data stream is opened to read the new record data.

Analysis state Interesting results comes from the Analysis state. As you can see in the figure 10.31, at minute 8m30s we launched an analysis command. The down peak is due to the drawing thread freezing as we stop all drawing for a while. If during the memory test we observed an increasing of memory allocation, here there is no significant changes into the cpu usage. We just pause the drawing thread and launch an AsyncTask to compute the analysis over the ECG record.

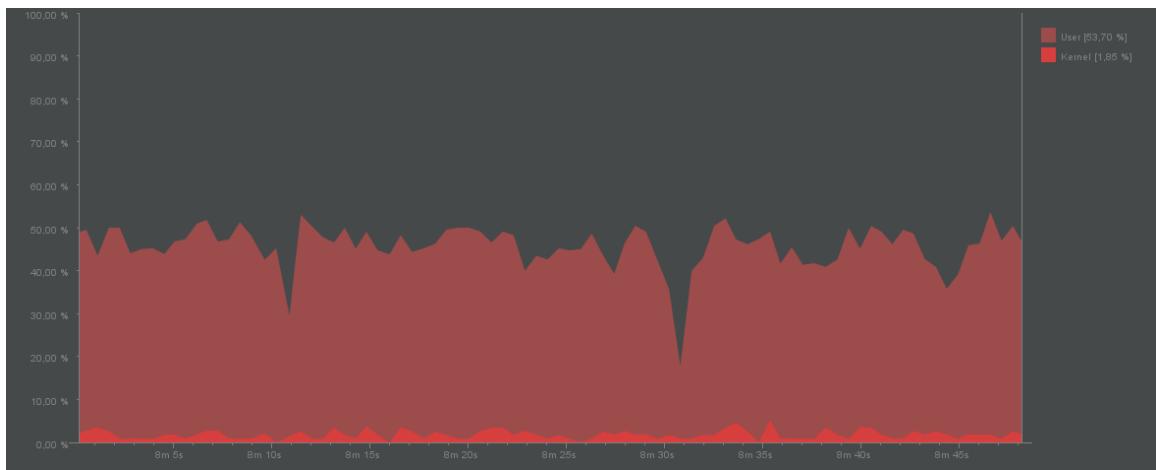


Figure 10.23: CPU usage during analysis state. An analysis command is launched over an ECG record strip.

OnePlus X CPU usage on this device didn't differ that much on the previous results on the Samsung N7000. The application achieved same performance and never exceeded 60% of the CPU capabilities. The graph suggests that in the Quad-core the average CPU usage is also smaller than the one found on the Samsung N7000's Dual-core.

Idle state On idle state the cpu usage is reduced at minimum. Less than 2% of CPU power usage is registered on static screen pages.

Plotting state During the plotting of a record the application CPU usage is constant in a range between 25% to 45%. The slope down is due to a change in activity view. The Record Activity is closed at minute 8m5s. The process (and all the threads) in charge to draw the ECG signals is killed so also resources are released as well. At minute 8m15s another record is opened and new resources and threads

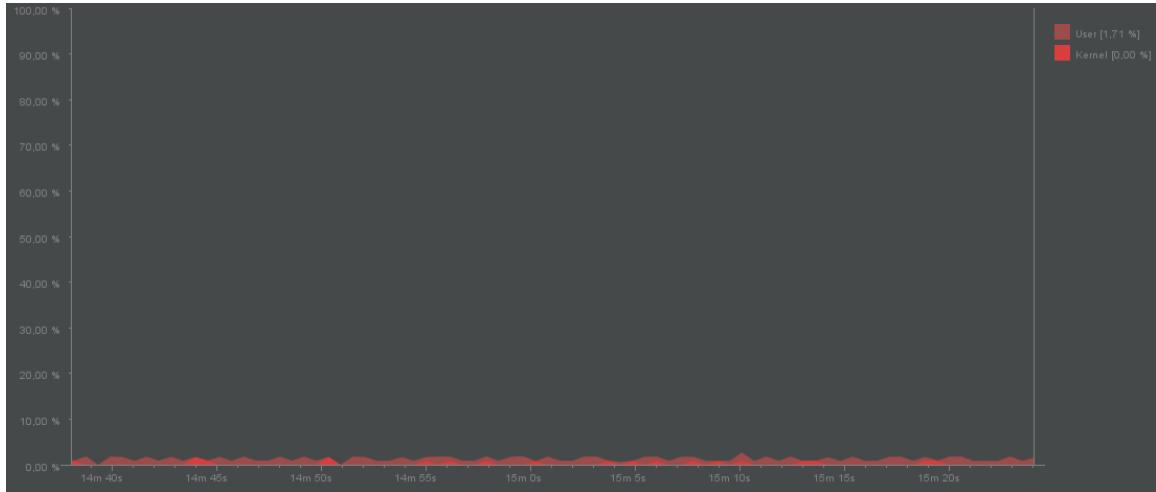


Figure 10.24: CPU usage on idle pages(app open on static screen page)

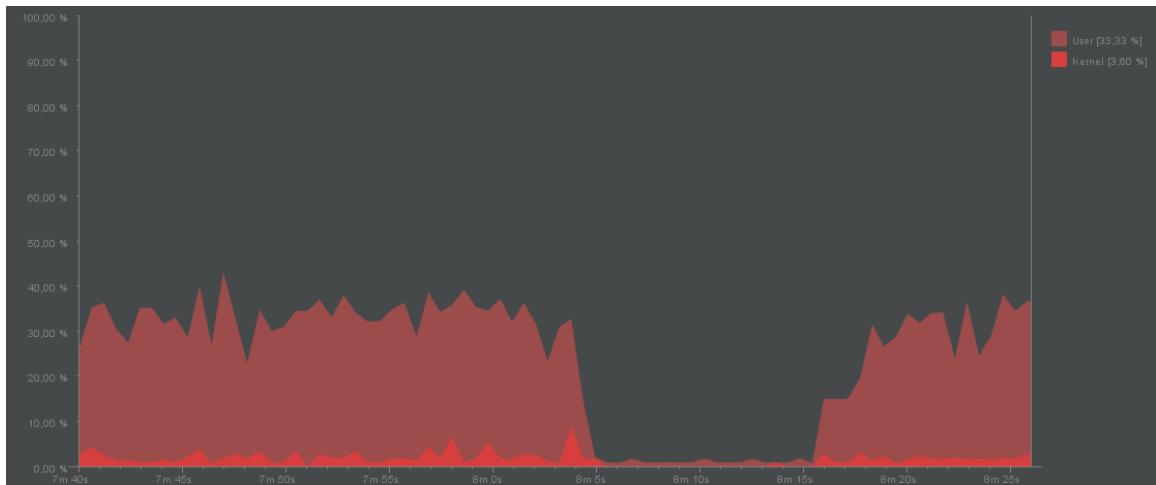


Figure 10.25: CPU usage during a record plotting)

are allocated.

We can observe that the overall cpu usage on the OnePlus X device is on average less than the average cpu usage on the Samsung N7000, even if the application overall performance results to have the same behaviour.

Analysis state During the analysis state the maximum cpu usage is not changed. It worth mentioning that even if there is no significant change in the % of cpu usage, on the OnePlus X the execution time of analysis is much shorter. In the figure 10.26 during a plotting state the analysis is launched at minute 3m30s (first white bar). The process terminates at minute 4m20s as shown by the second white bar.

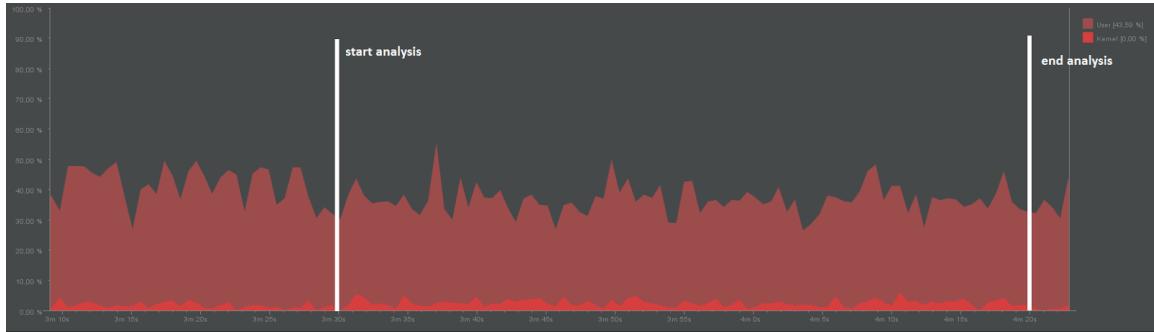


Figure 10.26: usage when the analysis algorithm is launched over a record)

10.1.3 Response Time

We performed a test about the response time and the time spent by the processor to load the plotting page and the time spent on executing the analysis algorithm.

To have accurate measurements we put inside the application additional lines of code as checkpoints.

Android OS offers an utility class to measure execution time of methods that is the `TimingLogger` class.

As documented into the Android references guide the class is used as follows:

```
TimingLogger timings = new TimingLogger(TAG, "methodA");
//do some work A
timings.addSplit("work A");
// do some work
timings.addSplit("work B");
//do some work C
timings.addSplit("work C");
timings.dumpToLog();
```

After a class declaration we split the time between the methods we want to check the execution time. We used such a class to measure the longest operations within ECG-ira, that is when we load the record into memory and when we execute the analysis on the record.

```
/*A timingLogger implementation to measure the execution time of the
two longest processes.*/
@Override
protected Boolean doInBackground(Void params) {
```

```

TimingLogger timeLogger = new
    TimingLogger("it.polimi.anto.chai.ecg_ira.Activity.AsyncTask",
    "RECORD ANALYSIS");
short[] result = parser.getSamplesForAnalysis(index);
timeLogger.addSplit("sample extraction from file to memory");
Analysis.execute(result, rate);
timeLogger.addSplit("analysis execution");
timeLogger.dumpToLog();
}

```

The same code was executed on different devices and as expected the OnePlus X performed much better in term of execution time.

Below is reported the execution time in milliseconds for the sample extraction from the .dat file to the memory and the time spent to execute the analysis on the record.

The memory load on the OnePlus X is extremely fast taking on average about 0.1

#of executions	Memory load (ms)	Analysis execution (ms)
run1	115	55038
run2	101	46429
run3	98	47775
run4	87	47910
run5	92	47087
average	98.6	48847.8

Table 10.2: Execution time for a record load in memory and its analysis on the OnePlus X device

second (98.6 milliseconds) on a record which lasts 30 minutes and dimension 1.9 MB. The analysis method is the longest in term of execution time. On average on the OnePlus X it takes about 48 seconds to execute (48.847 milliseconds). The results are not precise as the execution time takes in consideration many other independent external factors. The overall time can be less than the average or even much higher due to the number of processes running on background. Five runs are not exhaustive at all for a detailed performance analysis but it gives a general view of the performance we can get during a normal usage of the phone. The execution time on the Samsung N7000 are three time slower on average with a memory load time of 0.3 seconds and an Analysis execution time of 2.30 minutes.

#of executions	Memory load (ms)	Analysis execution (ms)
run1	275	121483
run2	320	128352
run3	358	136027
run4	242	144923
run5	334	156327
average	305.8	137422.4

Table 10.3: Execution time for a record load in memory and its analysis on the Samsung N7000 device

We observed that during other tests the application overall behaviour with respect to memory usage and amount of CPU usage is regular and the application doesn't show any kind of leakages nor crashes. The huge difference in the execution time is mostly due to the difference in the hardware and software strategy. The OnePlus X quad-core is obviously performing much faster than the Dual-core of the Samsung N7000. Also the availability of 3 GB of RAM can allow the OS not to call the GC as often as the OS on the Samsung device which has limited 1GB RAM to be shared across all the applications.

Chapter 11

Conclusions

ECG-ira is a reliable application for mobile devices for the acquisition, visualization and analysis of electrocardiographic recordings. The application is designed to be flexible, customizable, and easy to extend to additional features. With the increasing trend of the healthcare business, ECG-ira is a good help both for non-professionals and for professionals in healthcare, enabling them to monitor and read ECG recordings at any time and everywhere. In its current state, ECG-ira is compliant with the ZEcg acquisition device, only: however, ECG-ira can be easily enhanced to cope with more devices by exploiting its highly customizable settings. ECG-ira can be connected to other open electrocardiographic recordings and file formats.

The application was designed mobile first, considering all the trade-off and limitations of mobile devices. We exploited the best performance and provided the best user experience according to the last UI design and development patterns, independently from which (compatible) device the application is executed on.

In the near future, mobile applications will become more and more widespread, and will also replace desktop applications, as people moved their habit from desktop PCs to smartphone devices. The increasing phenomenon of the IoT (Internet of Things) will bring smart objects connected all to each other and smart devices and mobile applications for healthcare are becoming of common use for every consumer. The use of devices to measure an athlete's performance has been quite common for some years in the professional environment: nowadays, these technologies are more and more affordable also to a wider arena of non professional athletes and sport practitioners in general.

As people became more and more familiar with mobile applications, ECG-ira is

a complete solution for electrocardiographic recording acquisition, visualization and analysis. Due to its ease of use, taking an electrocardiographic recording does no longer require a professional healthcare expert (nurse or physician), and most people are already familiar with the smartphone devices.

11.1 Future works

ECG-ira represents a starting point to build up a complete software tool for the acquisition, visualization and analysis of electrocardiographic recordings. More ECG formats can be added, beyond those already implemented: ZEcg format (format “0” and format “1”) and the MIT/BIH format (format 212). Among the possible formats to enhance the performances of the application, we mention here the format “16” as defined by the American Heart Association (and also used by the Poli/Cá Granda ECG/VCG database) and the SCP standard [28].

More analysis techniques can also be added, such as ischemia detection techniques. Those techniques are mainly devoted to detect insufficient flow of blood to the heart, which may lead to acute myocardial infarction (IMA).

More enhancements include the capability of remote transmission of the recordings to an external web server: this could be the starting point for a revolutionary way of home-monitoring for patients suffering from chronic heart diseases.

Acknowledgement

*To my parents, mother Stella and father Gregorio who are always the fixed stars in
my sky.*

To my brother Dung who is in my blood.

To Dora who is in my deeper heart.

*To all my best friends who may be physically far away but I know they will answer to
any my calls. To all the ones who never believed on me.*

Thanks to all of them I am the man I wished to be.

*The future is a step ahead and thank to my beloved supporters I am not scared to go
further...*

Chai

To my family and all my friends.

To the destiny. To God...

Antonello

Bibliography

- [1] Blausen Medical Communications. Heart illustrations.
- [2] AllianceTek. Comparison chart for mobile app development methods. Norwegian University of Science and Technology, 2011.
- [3] Simone Battaglia Diego Ulisse Pizzagalli. Algoritmi per il riconoscimento automatico di eventi aritmici ed ischemici in segnali ecg. Master's thesis, Politecnico di Milano.
- [4] Android Developers. Android design principles. <http://developer.android.com/design/get-started/principles.html>.
- [5] Mediatek helio x20 - the world's first mobile processor with tri-cluster cpu architecture and ten processing cores (deca-core) and 30consumption: <http://mediatek-helio.com/x20/>.
- [6] Nicos et al Maglaveras. An adaptive backpropagation neural network for real-time ischemia episodes detection: development and performance analysis using the european st-t database. *Biomedical Engineering, IEEE Transactions*, 45(7):805–813, 1998.
- [7] GSMA Intelligence. Mobile platform wars, Feb. 2014.
- [8] Gartner. Global market share held by smartphone operating systems, Dec. 2015.
- [9] Lionbridge. Mobile web apps vs. mobile native apps, 2012.
- [10] Redda Yonathan Aklilu. Cross platform mobile applications development. Norwegian University of Science and Technology, 2012.

- [11] Palmieri Manuel, Inderjeet Singh, and Antonio Cicchetti. Comparison of cross-platform mobile development tools. In *Intelligence in Next Generation Networks (ICIN)*, pages 179–186, Oct. 2012.
- [12] Emily Blem, Jaikrishnan Menon, and Karthikeyan Sankaralingam. A detailed analysis of contemporary arm and x86 architectures. *UW-Madison Technical Report*, 2013.
- [13] Blanco-Velasco, Binwei Weng Manuel, and Kenneth E Barner. Ecg signal denoising and baseline wander correction based on the empirical mode decomposition. *Computers in biology and medicine*, 38(1):1–13, Jan. 2008.
- [14] Burattini L, W Zareba, and R Burattini. The effect of baseline wandering in automatic t-wave alternans detection from holter recordings. *Computers in Cardiology*, pages 257–260, Sep. 2006.
- [15] Leski, Jacek M, and Norbert Henzel. Ecg baseline wander and powerline interference reduction using nonlinear filter bank. *Signal processing*, 85(4):781–793, 2005.
- [16] Xu and Lisheng et al. Baseline wander correction in pulse waveforms using wavelet-based cascaded adaptive filter. *Computers in Biology and Medicine*, 37(5):716–731, 2007.
- [17] Hu, Xiao, Ying Gao, and Wai-Xi Liu. Pattern recognition of surface electromyography signal based on wavelet coefficient entropy. *Health*, 1(02):121, 2009.
- [18] Hu, Xiao, Zhong Xiao, and Ni Zhang. Removal of baseline wander from ecg signal based on a statistical weighted moving average filter. *Journal of Zhejiang University SCIENCE C*, 12(5):397–403, 2011.
- [19] Manjoo and Farhad. A murky road ahead for android, despite market dominance. *The New York Times*, May 2015.
- [20] Goransson Anders. *Efficient Android Threading: Asynchronous Processing Techniques for Android Applications*. O'Reilly Media, 1 edition, May 2014.
- [21] Lee Sangchul and Jae Wook Jeon. Evaluating performance of android platform using native c for embedded systems. In *Control Automation and Systems (ICCAS)*, pages 1160–1163, Oct. 2012.

- [22] Tahir Mir Nauman. *Learning Android Canvas*. Packt Publishing Ltd, 2013.
- [23] Crespi Alessandro and Ulisse Pizzagalli. Zecg: Acquisizione, trasmissione e memorizzazione del segnale ecg. Master's thesis, Politecnico di Milano.
- [24] Muschko and Benjamin. *Gradle in Action*. Manning, 2014.
- [25] Wikipedia. Dots per inch. https://en.wikipedia.org/wiki/Dots_per_inch, 2011.
- [26] Android Developers. Canvas and drawables. <http://developer.android.com/guide/topics/graphics/2d-graphics.html>, 2009.
- [27] Android Design Patterns. How to leak a context: Handlers & inner classes. <http://www.androiddesignpatterns.com/2013/01/inner-class-handler-memory-leak.html>, 2013.
- [28] V. A. Plotnikov, D. A. Prilutskii, and S. V. Selishchev. The scp-ecg standard in electrocardiographic software systems. *Biomedical Engineering*, 33(3):128–135, 1999.